

DSA 8020 R Session 0: A Quick Introduction to R

Whitney Huang

December 27, 2021

Contents

Basic	1
What is in the workspace?	1
Where is my R working directory?	2
What is in my working directory?	2
“c” combines sets of numbers (or datasets)	2
Now recheck workspace	2
Reassign “x” to another name	2
Remove “x” and create another object “x3”	2
Question: how would you combine “x2” and “x3” to make a new data set ?	3
Arithmetic	3
Add numbers in R	3
Some other operations	3
Generating a sequence	3
Generating sin wave	3
Use R to generate random values	5
Subsetting	6
Load a data set	6
Print the first 10 values	6
An indicator for all values over 70	7
Working with these data as a matrix	7
Apply Functions in R	8
<i>apply</i> functions	8
Load the Boulder temperature data set into R	8
The “ <i>apply</i> ” function	9
Writing Functions in R	10
Finding the inter quartile range (IQR)	10
Building your own function	10
Building your own IQR function	10
Modify this function to work with NAs	11
Adding warning message	11

This R session is modified from a R tutorial from Doug Nychka [\[link\]](#).

Basic

What is in the workspace?

```
ls()
```

```
## character(0)
```

Where is my R working directory?

```
getwd()
```

```
## [1] "/Users/whitneyhuang/Desktop/Desktop - mass-mini19-huang/Teaching/DSA/DSA8020/22 Spring"
```

What is in my working directory?

```
dir()
```

```
## [1] "BoulderTemperature.RData" "BT.RData"
## [3] "DSA8020_22sp_RCode0.pdf"  "DSA8020_22sp_RCode0.Rmd"
## [5] "DSA8020_22sp_Syllabus.pdf" "DSA8020_RCode1.Rmd"
## [7] "DSA8020_Slides0.pdf"      "notes0.pdf"
```

“c” combines sets of numbers (or datasets)

```
x <- c(2, 3, 20)
```

Note: R is case sensitive. Type “X” in R console and then click Enter to see what happen

Now recheck workspace

```
ls()
```

```
## [1] "x"
```

```
# print out x
```

```
x
```

```
## [1]  2  3 20
```

Reassign “x” to another name

```
x2 <- x
```

```
ls()
```

```
## [1] "x"  "x2"
```

Remove “x” and create another object “x3”

```
rm(x)
```

```
x3 <- c(3, 4, 5)
```

Question: how would you combine “x2” and “x3” to make a new data set ?

```
#give a try here
```

Arithmetic

Add numbers in R

```
A <- 2
B <- 10
Y <- A + B

A <- c(2, 3, 4)
B <- c(10, 100, 1000)
Y <- A + B
Y # note that the numbers have been added row by row like a spread sheet.
```

```
## [1] 12 103 1004
```

Some other operations

```
2^4
```

```
## [1] 16
```

```
2 * (1 + 4)
```

```
## [1] 10
```

```
sqrt(81)
```

```
## [1] 9
```

```
exp(2)
```

```
## [1] 7.389056
```

Generating a sequence

```
1:10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
-5:5
```


```
## [1] -5 -4 -3 -2 -1 0 1 2 3 4 5
```

Question: How would you generate the values in order 5 to 1?


Generating sin wave

```
x <- 0:100
# hundred values between 0 and 2*pi
```

```
x <- 2 * pi * (x / 100)
y <- sin(x)
# plot the sin wave
plot(x, y)
```



```
# change the plot to connect points with a line instead of points
plot(x, y, type = "l")
```



```
# another way of creating the x
x <- seq(0, 2 * pi, length.out = 101)
```

Use R to generate random values

```
# generating 10 random numbers between 0 and 1
```

```
U <- runif(10)
```

```
U
```

```
## [1] 0.614253061 0.687874304 0.009322416 0.736307974 0.878694482 0.611269986
```


```
## [7] 0.216010526 0.920248689 0.457006954 0.534641163
```

```
# generating 30 random numbers from a standard normal distribution
```

```
V <- rnorm(n = 30, mean = 0, sd = 1)
```


```
hist(V)
```

Histogram of V



```
qqnorm(V); qqline(V)
```

Normal Q-Q Plot



Subsetting

Load a data set

```
load("BT.RData")  
# copy for easy typing  
BT <- BoulderJuneTemperature$Temp  
BA11 <- BoulderJuneTemperature  
head(BT); head(BA11, 10)  
  
## [1] 65.51667 68.58333 69.21667 68.58333 70.91667 64.25000  
  
##   Year   Temp  
## 1  1984 65.51667  
## 2  1985 68.58333  
## 3  1986 69.21667  
## 4  1987 68.58333  
## 5  1988 70.91667  
## 6  1989 64.25000  
## 7  1990 69.95000  
## 8  1991 66.56667  
## 9  1992 62.90000  
## 10 1993 64.66667
```

Print the first 10 values

```
BT[1:10]
```

```
## [1] 65.51667 68.58333 69.21667 68.58333 70.91667 64.25000 69.95000 66.56667
## [9] 62.90000 64.66667
```

An indicator for all values over 70

```
ind70 <- BT > 70
ind70
```

```
## [1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE
## [25] FALSE FALSE FALSE FALSE TRUE FALSE
```

```
# temperatures with values over 70
BT[ind70]
```

```
## [1] 70.91667 70.05000 70.36667 71.56667 74.13333
```

```
# the years with values over 70
BA11$Year[ind70]
```

```
## [1] 1988 1994 2002 2006 2012
```

Question: How many years exceed 70 degrees?

Working with these data as a matrix

```
dim(BA11)
```

```
## [1] 30 2
```

```
# This is the first row and first column
BA11[1, 1]
```

```
## [1] 1984
```

```
# first row
BA11[1,]
```

```
## Year Temp
## 1 1984 65.51667
```

```
# first column
BA11[, 1]
```

```
## [1] 1984 1985 1986 1987 1988 1989 1990 1991 1992 1993 1994 1995 1996 1997 1998
## [16] 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012 2013
```

```
# column with Year (this is also column 1)
BA11[, "Year"]
```

```
## [1] 1984 1985 1986 1987 1988 1989 1990 1991 1992 1993 1994 1995 1996 1997 1998
## [16] 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012 2013
```

```
# second column (could also use "Temp" to refer to this)
BA11[, 2]
```

```
## [1] 65.51667 68.58333 69.21667 68.58333 70.91667 64.25000 69.95000 66.56667
## [9] 62.90000 64.66667 70.05000 62.33333 66.93333 66.40000 62.68333 64.85000
## [17] 67.40000 68.85000 70.36667 62.83333 62.70000 65.41667 71.56667 67.65000
```

```
## [25] 66.01667 63.16667 66.35000 67.56667 74.13333 69.82000
```

```
# rows 10 through 20  
BA11[10:20,]
```

```
##      Year      Temp  
## 10 1993 64.66667  
## 11 1994 70.05000  
## 12 1995 62.33333  
## 13 1996 66.93333  
## 14 1997 66.40000  
## 15 1998 62.68333  
## 16 1999 64.85000  
## 17 2000 67.40000  
## 18 2001 68.85000  
## 19 2002 70.36667  
## 20 2003 62.83333
```

Exercise: Plot the temperatures by year

Apply Functions in R

apply functions

1. a family of functions in R which allow you to repetitively perform an action on multiple chunks of data
2. run faster than loops and often require less code.

Let's take a look at some examples

Load the Boulder temperature data set into R

```
load("BoulderTemperature.RData") # monthly mean temperatures  
dim(BoulderTemperature)
```

```
## [1] 118 12
```

```
# check out first row  
BoulderTemperature[1,]
```

```
##      jan feb mar apr      may  jun      jul      aug      sep      oct  
## 1897 NaN NaN NaN NaN 60.25806 64.65 70.56452 69.06452 66.81667 52.41935  
##      nov      dec  
## 1897 41.86667 30.40323
```

```
#extract 1991 - 2010  
yr <- rownames(BoulderTemperature)  
index <- which(yr %in% 1991:2010)  
tempData <- BoulderTemperature[index,]  
# check this out  
tempData
```

```
##      jan      feb      mar      apr      may      jun      jul      aug  
## 1991 29.61290 40.96429 42.69355 47.73333 58.22581 66.56667 70.46774 69.11290  
## 1992 35.28571 40.56897 43.25806 54.21667 59.06452 62.90000 68.14516 66.29032  
## 1993 28.33871 30.58929 42.30645 47.56667 57.91935 64.66667 69.46774 67.37097
```



```
## 1994 35.50000 32.10714 43.82258 47.61667 60.80645 70.05000 71.14516 71.01613
## 1995 34.67742 38.28571 42.11290 44.51667 50.85484 62.33333 70.48387 73.96774
## 1996 29.70968 37.67241 37.62903 50.41667 58.87097 66.93333 71.45161 69.48387
## 1997 31.37097 32.96429 45.53226 42.81667 57.17742 66.40000 71.40323 68.85484
## 1998 36.50000 36.39286 38.54839 46.50000 58.61290 62.68333 72.75806 70.37097
## 1999 36.20968 42.10714 45.98387 44.55000 55.58065 64.85000 73.33871 69.30645
## 2000 36.41935 41.06897 42.85484 51.23333 60.98387 67.40000 74.66129 73.03226
## 2001 32.93548 32.32143 40.75806 50.63333 58.40323 68.85000 75.14516 71.85484
## 2002 33.11290 35.98214 37.27419 53.01667 56.16129 70.36667 76.90323 71.30645
## 2003 40.20968 32.08929 43.66129 50.60000 57.35484 62.83333 75.66129 72.79032
## 2004 35.40323 33.67241 48.17742 49.18333 59.96774 62.70000 69.16129 66.40323
## 2005 35.43548 37.87500 41.96774 48.40000 57.72581 65.41667 75.04839 69.70968
## 2006 40.66129 33.67857 39.38710 53.88333 60.95161 71.56667 74.37097 71.56452
## 2007 27.22581 34.58929 47.56452 47.81667 58.00000 67.65000 74.75806 73.56452
## 2008 31.62903 36.10345 40.75806 47.80000 57.03226 66.01667 75.01613 69.62903
## 2009 38.19355 39.33929 44.20968 47.30000 59.30645 63.16667 69.53226 69.48387
## 2010 33.01613 30.05357 42.37097 48.75000 53.90323 66.35000 72.45161 72.41935
##      sep      oct      nov      dec
## 1991 61.56667 52.14516 37.05000 35.48387
## 1992 64.41667 53.87097 34.00000 29.77419
## 1993 59.01667 48.61290 35.61667 35.41935
## 1994 64.83333 50.69355 36.53333 36.08065
## 1995 60.38333 51.33871 44.95000 36.24194
## 1996 60.76667 53.01613 40.58333 36.46774
## 1997 64.01667 52.66129 37.86667 33.83871
## 1998 67.20000 50.32258 44.01667 32.16129
## 1999 58.48333 51.90323 47.98333 36.91935
## 2000 63.10000 49.59677 31.31667 31.20968
## 2001 65.00000 53.83871 43.85000 34.98387
## 2002 64.06667 45.75806 40.26667 36.58065
## 2003 60.50000 57.38710 38.91667 36.35484
## 2004 62.85000 51.85484 39.66667 36.45161
## 2005 66.35000 53.09677 44.93333 33.30645
## 2006 58.40000 50.98387 43.36667 35.29032
## 2007 64.43333 55.17742 44.86667 30.06452
## 2008 60.90000 51.80645 46.20000 31.09677
## 2009 63.10000 44.46774 43.76667 26.66129
## 2010 66.55000 54.77419 39.76667 37.19355
```

The “*apply*” function

```
# means by rows of this table
byYear <- apply(tempData, 1, FUN = mean) # by rows, 1 = first index
byYear
```

```
##      1991      1992      1993      1994      1995      1996      1997      1998
## 50.96857 50.98260 48.90762 51.68375 50.84554 51.08345 50.40858 51.33892
##      1999      2000      2001      2002      2003      2004      2005      2006
## 52.26798 51.90642 52.38118 51.73297 52.36322 51.29098 52.43878 52.84208
##      2007      2008      2009      2010
## 52.14257 51.16565 50.71062 51.46661
```

```
# means by columns
byMonth <- apply(tempData, 2, FUN = mean) # by cols, 2 = second index
byMonth

##      jan      feb      mar      apr      may      jun      jul      aug
## 34.07235 35.92127 42.54355 48.72750 57.84516 65.98500 72.56855 70.37661
##      sep      oct      nov      dec
## 62.79667 51.66532 40.77583 34.07903
```

Writing Functions in R

Finding the inter quartile range (IQR)

```
# 75% quantile
BT75 <- quantile(BT, .75)
#Question: Find the interquartile range 75% - 25% quantiles
# and check this against the built in function
IQR(BT)

## [1] 4.4125
```

Building your own function

Here is a function that adds the squares of two numbers. It has three parts, the *calling arguments*, the *body* where you do the work and then *returning any results*.

```
myFun <- function(a, b){
  result <- a^2 + b^2
  return(result)
}

test1 <- myFun(2, 3)
test1
```

```
## [1] 13

test2 <- myFun(1:5, 11:15)
test2
```

```
## [1] 122 148 178 212 250
```

Note that the “a”, “b” and result are only used inside the function and do not appear in your workspace. Also since the body is normal R code, this works for vectors automatically.

Building your own IQR function

```
myIQR <- function(y){
  IQR <- quantile(y, .75, names = FALSE) - quantile(y, .25, names = FALSE)
  return(IQR)
}
myIQR(BT)

## [1] 4.4125
```

Modify this function to work with NAs

```
test <- c(BT, NA)
myIQR <- function(y, na.rm = FALSE){
  IQR <- quantile(y, .75, names = FALSE, na.rm = na.rm) -
    quantile(y, .25, names = FALSE, na.rm = na.rm)
  return(IQR)
}
myIQR(test, na.rm = T)

## [1] 4.4125
```

Adding warning message

```
myIQR <- function(y, na.rm = FALSE){
  # an example of adding a warning
  if(na.rm){
    warning("NAs removed from the data")
  }
  IQR <- quantile(y, .75, names = FALSE, na.rm = na.rm) -
    quantile(y, .25, names = FALSE, na.rm = na.rm)
  return(IQR)
}
myIQR(test, na.rm = T)

## Warning in myIQR(test, na.rm = T): NAs removed from the data
## [1] 4.4125
```