

DSA 8020 R Session 10: Random and Mixed Effects Models and Computer Experiments

Whitney

March 20, 2023

Contents

Random Effects Example	1
Read the data into R	1
Fitting a fixed effects model	2
Fitting a random effects model	3
RCBD: Fixed vs. Random Block	3
Load R libraries	3
Read the data	4
Fixed block	4
Random block	4
Computer Experiments	5
Design: Latin hypercube	5
Analysis: Gaussian Process	7

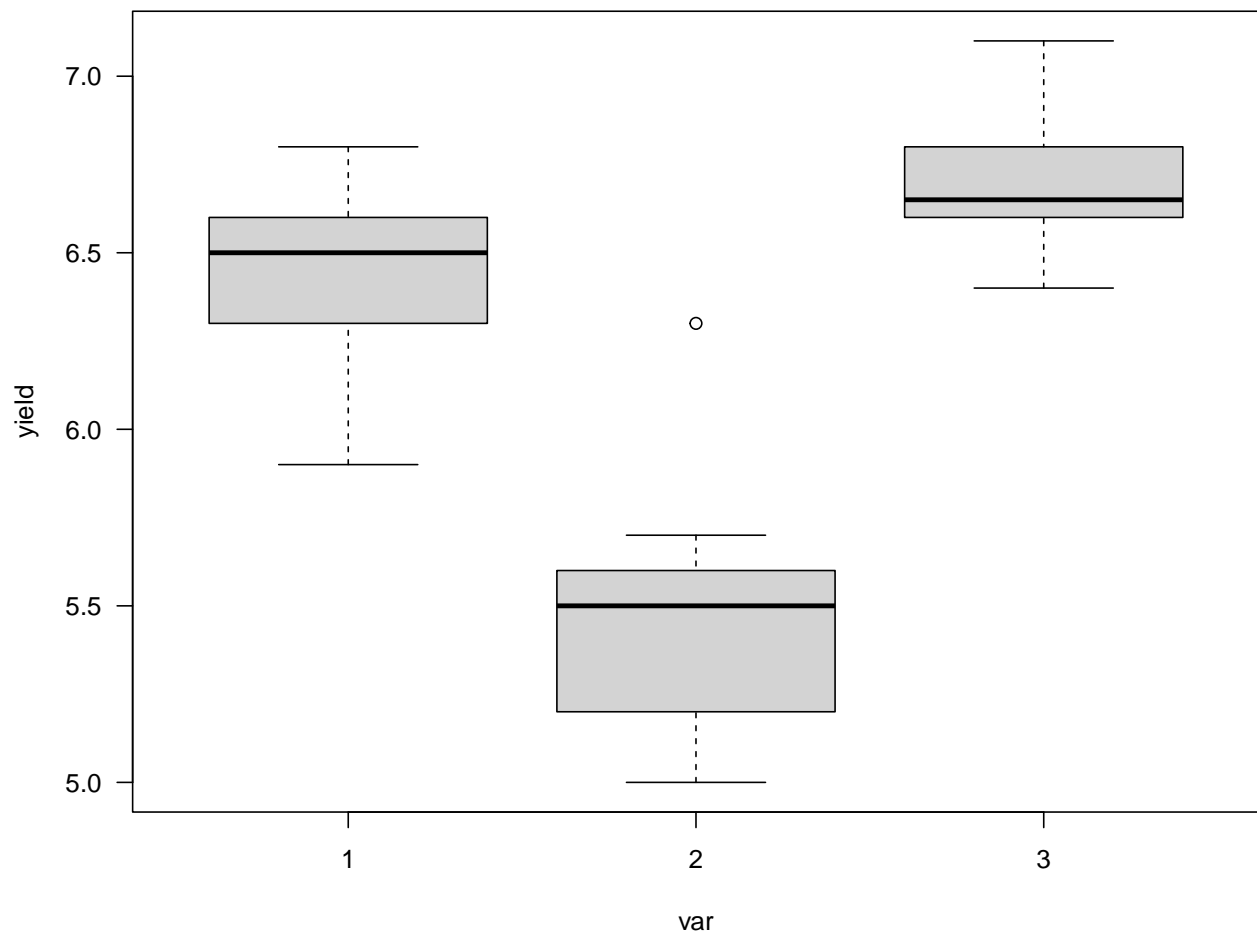
Random Effects Example

Suppose that an agronomist is studying a large number of varieties of soybeans for yield. The agronomist randomly selects three varieties, and then randomly assigns each of those varieties to 10 of 30 available plots.

Model: $y_{ij} = \mu + \alpha_i + \epsilon_{ij}$, $\alpha_i s \stackrel{i.i.d.}{\sim} N(0, \sigma_\alpha^2)$, $\epsilon_{ij} s \stackrel{i.i.d.}{\sim} N(0, \sigma^2)$. $\alpha_i s$ and $\epsilon_{ij} s$ are independent to each other

Read the data into R

```
v1 <- c(6.6, 6.4, 5.9, 6.6, 6.2, 6.7, 6.3, 6.5, 6.5, 6.8)
v2 <- c(5.6, 5.2, 5.3, 5.1, 5.7, 5.6, 5.6, 6.3, 5.0, 5.4)
v3 <- c(6.9, 7.1, 6.4, 6.7, 6.5, 6.6, 6.6, 6.6, 6.8, 6.8)
yield <- c(v1, v2, v3)
var <- factor(c(rep(1, 10), rep(2, 10), rep(3, 10)))
plot(yield ~ var, las = 1)
```



Fitting a fixed effects model

```
fixef <- lm(yield ~ var)
anova(fixef)
```

```
## Analysis of Variance Table
##
## Response: yield
##          Df Sum Sq Mean Sq F value    Pr(>F)
## var         2  8.306   4.1530  49.593 9.114e-10 ***
## Residuals  27  2.261   0.0837
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coefficients(fixef)
```

```
## (Intercept)      var2      var3
##          6.45      -0.97       0.25
```

Fitting a random effects model

```
library(lme4)
```

```
## Loading required package: Matrix
```

```
randef <- lmer(yield ~ 1 + (1|var), REML = TRUE)
summary(randef)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: yield ~ 1 + (1 | var)
##
## REML criterion at convergence: 21.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.8839 -0.6181  0.1118  0.4962  2.7828
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##   var      (Intercept) 0.40693  0.6379
##   Residual                0.08374  0.2894
## Number of obs: 30, groups: var, 3
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)   6.2100    0.3721   16.69
```

Let's construct CIs for σ_α^2 , σ^2 , and μ

```
## Compute the confidence intervals (CIs) using profile likelihood
CIs <- confint(randef, oldNames = FALSE)
```

```
## Computing profile confidence intervals ...
```

```
CIs
```

```
##              2.5 %    97.5 %
## sd_(Intercept)|var 0.2637525 1.5512218
## sigma              0.2265053 0.3877781
## (Intercept)       5.3618584 7.0581407
```

RCBD: Fixed vs. Random Block

Load R libraries

```
library(lsmeans)
library(lmerTest)
```

Read the data

```
### Create the data set
x <- c(52, 47, 44, 51, 42, 60, 55, 49, 52, 43, 56, 48, 45, 44, 38)
trt <- rep(c("A", "B", "C"), each = 5)
blk <- rep(1:5, 3)
dat <- data.frame(x = x, trt = trt, blk = as.factor(blk))
```

Fixed block

```
fixef <- lm(x ~ trt + blk, data = dat)
anova(fixef)
```

```
## Analysis of Variance Table
##
## Response: x
##           Df Sum Sq Mean Sq F value    Pr(>F)
## trt         2   89.2    44.60   7.6239 0.0140226 *
## blk         4  363.6    90.90  15.5385 0.0007684 ***
## Residuals   8   46.8     5.85
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
lsmeans(fixef, list(pairwise ~ trt), adjust = "none")
```

```
## $'lsmeans of trt'
##   trt lsmean   SE df lower.CL upper.CL
## A     47.2 1.08  8     44.7     49.7
## B     51.8 1.08  8     49.3     54.3
## C     46.2 1.08  8     43.7     48.7
##
## Results are averaged over the levels of: blk
## Confidence level used: 0.95
##
## $'pairwise differences of trt'
##    1      estimate    SE df t.ratio p.value
## A - B     -4.6 1.53  8    -3.007  0.0169
## A - C      1.0 1.53  8     0.654  0.5316
## B - C      5.6 1.53  8     3.661  0.0064
##
## Results are averaged over the levels of: blk
```

Random block

```
randef <- lmer(x ~ trt + (1|blk), REML = TRUE, data = dat)
summary(randef)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: x ~ trt + (1 | blk)
## Data: dat
##
## REML criterion at convergence: 71.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.1417 -0.6147 -0.1494  0.5772  1.3390
##
## Random effects:
## Groups   Name                Variance Std.Dev.
## blk      (Intercept) 28.35     5.324
## Residual                    5.85     2.419
## Number of obs: 15, groups: blk, 5
##
## Fixed effects:
##              Estimate Std. Error    df t value Pr(>|t|)
## (Intercept)   47.200      2.615   5.054  18.047 8.76e-06 ***
## trtB           4.600      1.530   8.000   3.007  0.0169 *
## trtC          -1.000      1.530   8.000  -0.654  0.5316
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##      (Intr) trtB
## trtB -0.292
## trtC -0.292  0.500
```

Computer Experiments

Design: Latin hypercube

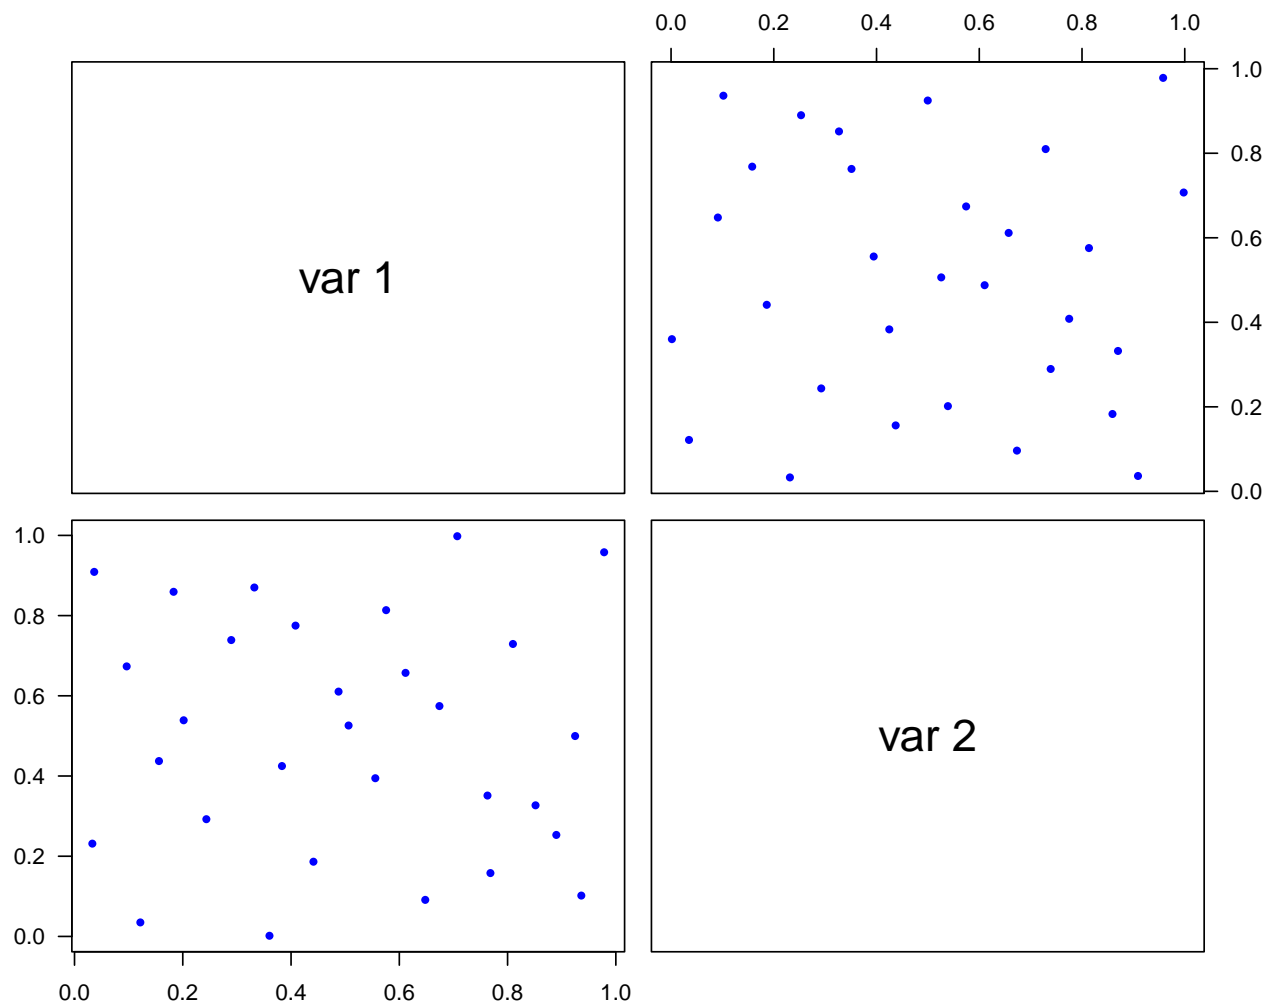
```
library(lhs)
LHD = maximinLHS(n = 30, k = 2, dup = 5)
# "dup" is an integer tuning parameter that determines the number of
# candidate points considered. Larger values should improve results
# but require more computational resources.

# Display the LHD
LHD
```

```
##              [,1]      [,2]
## [1,] 0.85163980 0.326915406
## [2,] 0.61148855 0.657195815
## [3,] 0.50628303 0.525948417
## [4,] 0.15601821 0.437273694
## [5,] 0.28953798 0.739029594
## [6,] 0.18301146 0.859320828
## [7,] 0.89002961 0.253087406
## [8,] 0.40827579 0.775041648
```

```
## [9,] 0.76830765 0.157968310
## [10,] 0.38320445 0.424840000
## [11,] 0.44138172 0.186290866
## [12,] 0.67412738 0.574413634
## [13,] 0.09636128 0.673347404
## [14,] 0.12161340 0.034926174
## [15,] 0.76287225 0.351277499
## [16,] 0.70708137 0.997937587
## [17,] 0.48775888 0.610443650
## [18,] 0.80988089 0.729239290
## [19,] 0.24355766 0.292339308
## [20,] 0.33220007 0.870074028
## [21,] 0.64791544 0.091138286
## [22,] 0.92476507 0.499698644
## [23,] 0.20163000 0.538966427
## [24,] 0.36014331 0.001642486
## [25,] 0.57562602 0.813562347
## [26,] 0.93621045 0.101869528
## [27,] 0.03293315 0.231395050
## [28,] 0.97835142 0.957936952
## [29,] 0.03637950 0.909008997
## [30,] 0.55566473 0.394488435
```

```
pairs(LHD, col = "blue", cex = 0.8, pch = 16, las = 1)
```



Analysis: Gaussian Process

```
# Load the data
neuron <- read.table("http://deanvosdraguljic.ietsandbox.net/DeanVossDraguljic/R-data/neuron.txt",
                     header = T)
head(neuron, 10)
```

```
##      gNaFsc    gKdrsc fr
## 1 0.38593729 0.2120652 33
## 2 0.04666927 0.4594742  0
## 3 1.00000000 0.4473344 46
## 4 0.95467637 0.3351407 44
## 5 0.53334929 0.7981310 41
## 6 0.59166751 0.6042714 41
## 7 0.18570301 0.3799469 31
## 8 0.49927784 0.2444170 36
## 9 0.74609113 0.3949591 42
## 10 0.07269414 1.0000000  0
```

```

# Fit a GP
library(mlegp)
GPFit <- mlegp(neuron[, 1:2], neuron[, 3])

## no reps detected - nugget will not be estimated
##
## ===== FITTING GP # 1 =====
## running simplex # 1...
## ...done
## ...simplex #1 complete, loglike = -104.446501 (convergence)
## running simplex # 2...
## ...done
## ...simplex #2 complete, loglike = -104.446501 (convergence)
## running simplex # 3...
## ...done
## ...simplex #3 complete, loglike = -104.446502 (convergence)
## running simplex # 4...
## ...done
## ...simplex #4 complete, loglike = -104.446501 (convergence)
## running simplex # 5...
## ...done
## ...simplex #5 complete, loglike = -104.446501 (convergence)
##
## using L-BFGS method from simplex #1...
## iteration: 1,loglike = -104.446501
## ...L-BFGS method complete
##
## Maximum likelihood estimates found, log like = -104.446501
## creating gp object.....done

```

```

summary(GPFit)

##
## Total observations = 30
## Dimensions = 2
##
## mu = 27.61157
## sig2: 251.8751
## nugget: 0
##
## Correlation parameters:
##
##      beta a
## 1  5.027878 2
## 2 50.228477 2
##
## Log likelihood = -104.4465
##
## CV RMSE: 7.312618
## CV RMaxSE: 1020.777

```



```

# Make prediction
predictedX = expand.grid(g_NaF = seq(0, 1, 0.02), g_KDR = seq(0, 1, 0.02))
yhats = predict(GPfit, predictedX, se.fit = T)
# Visualize predictions and their uncertainty
library(fields)
par(mfrow = c(1, 2))
image.plot(seq(0, 1, 0.02), seq(0, 1, 0.02), matrix(yhats$fit, 51, 51),
           xlab = "g NaF (mS/cm^2)", ylab = "g KDR (mS/cm^2)", las = 1,
           main = "Predictions")
points(neuron[, 1:2], pch = 16, cex = 0.75)
image.plot(seq(0, 1, 0.02), seq(0, 1, 0.02), matrix(yhats$se.fit, 51, 51),
           xlab = "g NaF (mS/cm^2)", ylab = "g KDR (mS/cm^2)", las = 1,
           main = "Predictions Uncertainty")
points(neuron[, 1:2], pch = 16, cex = 0.75)

```

