

DSA 8020 R Session 6: Non-parametric Regression and Shrinkage Methods

Whitney

February 07, 2021

Contents

Non-parametric Regression: Motorcycle Accident Simulation Data	1
Load and plot the data	1
Linear and polynomial regression fits	2
Kernel regression	3
Local Polynomial Regression Fitting (<i>loess</i>)	4
Regression Splines	5
Generalized additive models	6
Smoothing splines	8
Comparing Regression spline/GAM/smoothing spline fits	10
Shrinkage Methods	11
Ridge Regression	11
The Lasso	17

Non-parametric Regression: Motorcycle Accident Simulation Data


A data frame giving a series of measurements of head acceleration in a simulated motorcycle accident, used to test crash helmets.

- **times**: time in milliseconds after impact
- **accel**: head acceleration in g

Data Source: Silverman, B. W. (1985) Some aspects of the spline smoothing approach to non-parametric curve fitting. *Journal of the Royal Statistical Society series B* 47, 1–52.

Load and plot the data


```
library(MASS)
data(mcycle)
attach(mcycle)
plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
```



Linear and polynomial regression fits

```
rg <- range(times)
xg = seq(rg[1], rg[2], 0.1) # prediction grids

plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lmFit <- lm(accel ~ times, data = mcycle)
abline(lmFit, col = "red")
Cub.polyFit <- lm(accel ~ poly(times, 3), data = mcycle)
Cub.polyPred <- predict(Cub.polyFit, data.frame(times = xg))
lines(xg, Cub.polyPred, col = "blue")
```




Kernel regression

$\hat{f}(x) = \hat{\mathbb{E}}(Y|X = x) = \frac{\sum_{i=1}^n K_h(x-x_i)y_i}{\sum_{i=1}^n K_h(x-x_i)}$, where K_h is a kernel with a bandwidth h .

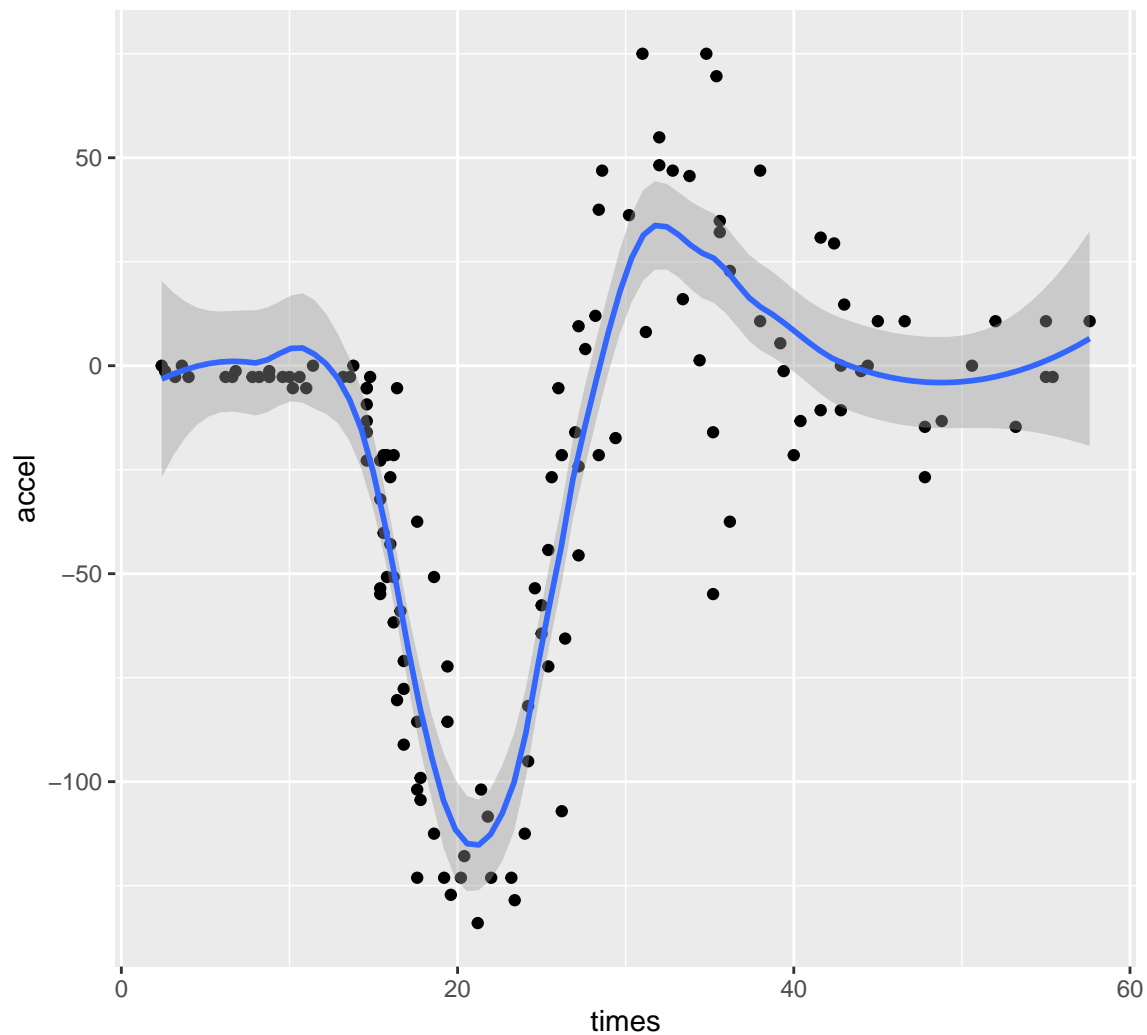
```
KernFit <- with(mcycle, ksmooth(times, accel, kernel = "normal", bandwidth = 0.5))
KernFit2 <- with(mcycle, ksmooth(times, accel, kernel = "normal", bandwidth = 5))

plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lines(KernFit$x, KernFit$y, col = "darkgreen")
lines(KernFit2$x, KernFit2$y, col = "blue")
```



Local Polynomial Regression Fitting (*loess*)

```
library(ggplot2)
plot <- ggplot(aes(x = times, y = accel), data = mcycle)
plot <- plot + geom_point()
(plot <- plot + geom_smooth(method = "loess", degree = 2, span = 0.4, se = TRUE))
```



Regression Splines

```
library(splines)
RegSplineFit <- lm(accel ~ bs(times, df = 10), data = mcycle)
summary(RegSplineFit)
```

```
##
## Call:
## lm(formula = accel ~ bs(times, df = 10), data = mcycle)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-76.673	-12.362	-0.557	13.139	51.740


```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.9312	14.4492	0.064	0.94872
bs(times, df = 10)1	-12.2008	37.5144	-0.325	0.74556
bs(times, df = 10)2	6.2223	23.6415	0.263	0.79284

```
## bs(times, df = 10)3      -7.3726      18.2652    -0.404    0.68718
## bs(times, df = 10)4    -118.7497      17.9975    -6.598    1.13e-09 ***
## bs(times, df = 10)5    -152.4486      20.0955    -7.586    7.25e-12 ***
## bs(times, df = 10)6      50.0827      18.7966      2.664    0.00875 **
## bs(times, df = 10)7      19.4271      19.3827      1.002    0.31819
## bs(times, df = 10)8      -8.1814      23.9354     -0.342    0.73308
## bs(times, df = 10)9     -11.1443      29.2202     -0.381    0.70358
## bs(times, df = 10)10      8.6378      23.6119      0.366    0.71513
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22.68 on 122 degrees of freedom
## Multiple R-squared:  0.7964, Adjusted R-squared:  0.7797
## F-statistic: 47.72 on 10 and 122 DF,  p-value: < 2.2e-16
```

```
RegSplinePred <- predict(RegSplineFit, data.frame(times = xg))

plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lines(xg, RegSplinePred, col = "darkgreen")
```




Generalized additive models

```
library(mgcv)
GAMFit <- gam(accel ~ s(times), data = mcycle)
summary(GAMFit)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## accel ~ s(times)
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -25.546      1.951   -13.1   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F p-value
## s(times)  8.693  8.972 53.52  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.783   Deviance explained = 79.8%
## GCV = 545.78   Scale est. = 506         n = 133
```

```
GAMPred <- predict(GAMFit, data.frame(times = xg))

plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lines(xg, GAMPred, col = "red")
```



Smoothing splines

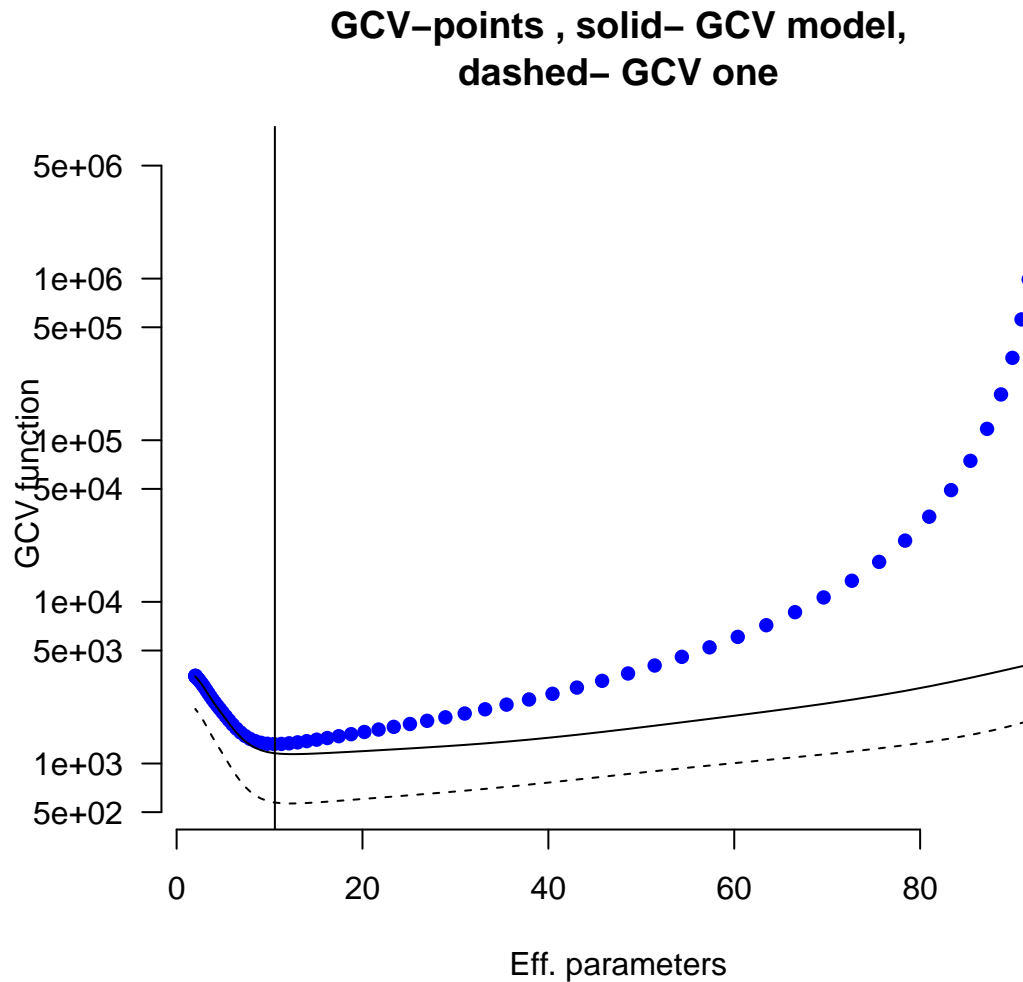
```
library(fields)
SpFit <- sreg(times, accel)
summary(SpFit)
```

```
## CALL:
## sreg(x = times, y = accel)
##
## Number of Observations:      133
## Number of unique points:     133
## Eff. degrees of freedom for spline: 10.6
## Residual degrees of freedom:  122.4
## GCV est. sigma                22.97
## Pure error sigma              24.49
## lambda                       0.3826
##
## RESIDUAL SUMMARY:
##      min    1st Q   median    3rd Q     max
## -78.1500 -13.8800  -0.7238  13.6300  49.6300
##
## DETAILS ON SMOOTHING PARAMETER:
## Method used:      Cost:
##   lambda      trA      GCV   GCV.one GCV.model      shat
```




```
##      0.3826   10.5726 1318.0646  573.4152 1156.4850   22.9746
##
## Summary of estimates for lambda
##      lambda   trA    GCV  shat converge
## GCV        0.3826 10.573 1318.1 22.97      13
## GCV.model   0.1835 12.467 1142.5 22.64      12
## GCV.one     0.1981 12.253  565.5 22.66      12
## pure error  1.1041  8.375 1380.7 24.49      NA
```

```
plot(SpFit, which = 3, col = "blue", pch = 16, las = 1)
```



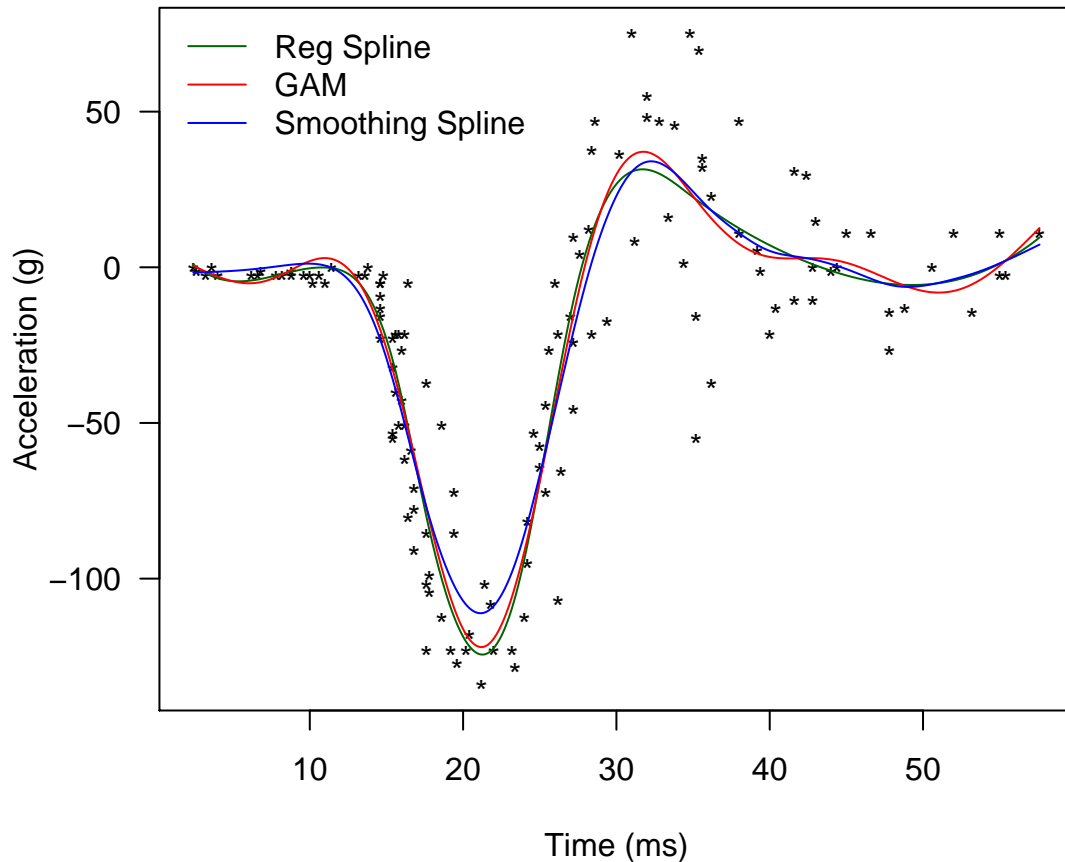
```
SpPred <- predict(SpFit, xg)

plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lines(xg, SpPred, col = "blue")
```



Comparing Regression spline/GAM/smoothing spline fits

```
plot(times, accel, pch = "*", cex = 1, las = 1,
      xlab = "Time (ms)", ylab = "Acceleration (g)")
lines(xg, RegSplinePred, col = "darkgreen")
lines(xg, GAMpred, col = "red")
lines(xg, SpPred, col = "blue")
legend("topleft", legend = c("Reg Spline", "GAM", "Smoothing Spline"),
      col = c("darkgreen", "red", "blue"), lty = 1, bty = "n")
```



Shrinkage Methods

The rest of this R session is largely based on the R lab: Ridge Regression and the Lasso of the book “Introduction to Statistical Learning with Applications in R” by *Gareth James, Daniela Witten, Trevor Hastie* and *Robert Tibshirani*. We will use the `glmnet` package to perform ridge regression and the lasso to predict `Salary` on the `Hitters` data.

Ridge Regression

1. Data Setup

```
library(ISLR)
data(Hitters)
Hitters = na.omit(Hitters)
head(Hitters)
```

```
##           AtBat Hits HmRun Runs RBI Walks Years CAtBat CHits CHmRun
## -Alan Ashby    315   81     7   24  38   39    14   3449   835    69
## -Alvin Davis   479  130    18   66  72   76     3   1624   457    63
## -Andre Dawson  496  141    20   65  78   37    11   5628  1575   225
## -Andres Galarra 321   87    10   39  42   30     2    396   101    12
## -Alfredo Griffin 594  169     4   74  51   35    11  4408  1133    19
## -Al Newman    185   37     1   23   8   21     2   214    42     1
##           CRuns CRBI CWalks League Division PutOuts Assists Errors
```

## -Alan Ashby	321	414	375	N	W	632	43	10
## -Alvin Davis	224	266	263	A	W	880	82	14
## -Andre Dawson	828	838	354	N	E	200	11	3
## -Andres Galarrraga	48	46	33	N	E	805	40	4
## -Alfredo Griffin	501	336	194	A	W	282	421	25
## -Al Newman	30	9	24	N	E	76	127	7
##	Salary NewLeague							
## -Alan Ashby	475.0		N					
## -Alvin Davis	480.0		A					
## -Andre Dawson	500.0		N					
## -Andres Galarrraga	91.5		N					
## -Alfredo Griffin	750.0		A					
## -Al Newman	70.0		A					

summary(Hitters)

##	AtBat	Hits	HmRun	Runs
##	Min. : 19.0	Min. : 1.0	Min. : 0.00	Min. : 0.00
##	1st Qu.: 282.5	1st Qu.: 71.5	1st Qu.: 5.00	1st Qu.: 33.50
##	Median : 413.0	Median : 103.0	Median : 9.00	Median : 52.00
##	Mean : 403.6	Mean : 107.8	Mean : 11.62	Mean : 54.75
##	3rd Qu.: 526.0	3rd Qu.: 141.5	3rd Qu.: 18.00	3rd Qu.: 73.00
##	Max. : 687.0	Max. : 238.0	Max. : 40.00	Max. : 130.00
##	RBI	Walks	Years	CAtBat
##	Min. : 0.00	Min. : 0.00	Min. : 1.000	Min. : 19.0
##	1st Qu.: 30.00	1st Qu.: 23.00	1st Qu.: 4.000	1st Qu.: 842.5
##	Median : 47.00	Median : 37.00	Median : 6.000	Median : 1931.0
##	Mean : 51.49	Mean : 41.11	Mean : 7.312	Mean : 2657.5
##	3rd Qu.: 71.00	3rd Qu.: 57.00	3rd Qu.: 10.000	3rd Qu.: 3890.5
##	Max. : 121.00	Max. : 105.00	Max. : 24.000	Max. : 14053.0
##	CHits	CHmRun	CRuns	CRBI
##	Min. : 4.0	Min. : 0.00	Min. : 2.0	Min. : 3.0
##	1st Qu.: 212.0	1st Qu.: 15.00	1st Qu.: 105.5	1st Qu.: 95.0
##	Median : 516.0	Median : 40.00	Median : 250.0	Median : 230.0
##	Mean : 722.2	Mean : 69.24	Mean : 361.2	Mean : 330.4
##	3rd Qu.: 1054.0	3rd Qu.: 92.50	3rd Qu.: 497.5	3rd Qu.: 424.5
##	Max. : 4256.0	Max. : 548.00	Max. : 2165.0	Max. : 1659.0
##	CWalks	League Division	PutOuts	Assists
##	Min. : 1.0	A:139 E:129	Min. : 0.0	Min. : 0.0
##	1st Qu.: 71.0	N:124 W:134	1st Qu.: 113.5	1st Qu.: 8.0
##	Median : 174.0		Median : 224.0	Median : 45.0
##	Mean : 260.3		Mean : 290.7	Mean : 118.8
##	3rd Qu.: 328.5		3rd Qu.: 322.5	3rd Qu.: 192.0
##	Max. : 1566.0		Max. : 1377.0	Max. : 492.0
##	Errors	Salary	NewLeague	
##	Min. : 0.000	Min. : 67.5	A:141	
##	1st Qu.: 3.000	1st Qu.: 190.0	N:122	
##	Median : 7.000	Median : 425.0		
##	Mean : 8.593	Mean : 535.9		
##	3rd Qu.: 13.000	3rd Qu.: 750.0		
##	Max. : 32.000	Max. : 2460.0		

```
library(glmnet)
X <- model.matrix(Salary ~ ., data = Hitters)[, -1]
y <- Hitters$Salary
```

The `glmnet()` function has an `alpha` argument that determines what type of model is fit. If `alpha = 0` then a ridge regression model is fit, and if `alpha = 1` then a lasso model is fit. We first fit a ridge regression model, which minimizes

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p \beta_j^2,$$

where $\lambda \geq 0$ is a *tuning parameter* to be determined.

2. Fit Ridge Regression over a grid of λ values

```
grid <- 10^seq(10, -2, length = 100)
ridge.mod <- glmnet(X, y, alpha = 0, lambda = grid)
```

3. Ridge Regression Coefficients

```
dim(coef(ridge.mod))
```

```
## [1] 20 100
```

We expect the coefficient estimates to be much smaller, in terms of ℓ_2 norm, when a large value of λ is used.

```
ridge.mod$lambda[50] #Display 50th lambda value
```

```
## [1] 11497.57
```

```
coef(ridge.mod)[, 50] # Display coefficients associated with 50th lambda value
```

##	(Intercept)	AtBat	Hits	HmRun	Runs
##	407.356050200	0.036957182	0.138180344	0.524629976	0.230701523
##	RBI	Walks	Years	CAtBat	CHits
##	0.239841459	0.289618741	1.107702929	0.003131815	0.011653637
##	CHmRun	CRuns	CRBI	CWalks	LeagueN
##	0.087545670	0.023379882	0.024138320	0.025015421	0.085028114
##	DivisionW	PutOuts	Assists	Errors	NewLeagueN
##	-6.215440973	0.016482577	0.002612988	-0.020502690	0.301433531

```
sqrt(sum(coef(ridge.mod)[-1, 50]^2)) # Calculate l2 norm
```

```
## [1] 6.360612
```

In contrast, here are the coefficients when $\lambda = 705$, along with their ℓ_2 norm. Note the much larger ℓ_2 norm of the coefficients associated with this smaller value of λ .

```
ridge.mod$lambda[60] #Display 60th lambda value
```

```
## [1] 705.4802
```

```
coef(ridge.mod)[, 60] # Display coefficients associated with 60th lambda value
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 54.32519950  0.11211115  0.65622409  1.17980910  0.93769713  0.84718546
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  1.31987948  2.59640425  0.01083413  0.04674557  0.33777318  0.09355528
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  0.09780402  0.07189612 13.68370191 -54.65877750  0.11852289  0.01606037
##      Errors      NewLeagueN
## -0.70358655  8.61181213
```

```
sqrt(sum(coef(ridge.mod)[-1, 60]^2)) # Calculate l2 norm
```

```
## [1] 57.11001
```

We can use the `predict()` function for a number of purposes. For instance, we can obtain the ridge regression coefficients for a new value of λ , say 50:

```
predict(ridge.mod, s = 50, type = "coefficients")[1:20, ]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs
## 4.876610e+01 -3.580999e-01  1.969359e+00 -1.278248e+00  1.145892e+00
##      RBI      Walks      Years      CAtBat      CHits
## 8.038292e-01  2.716186e+00 -6.218319e+00  5.447837e-03  1.064895e-01
##      CHmRun      CRuns      CRBI      CWalks      LeagueN
## 6.244860e-01  2.214985e-01  2.186914e-01 -1.500245e-01  4.592589e+01
##      DivisionW      PutOuts      Assists      Errors      NewLeagueN
## -1.182011e+02  2.502322e-01  1.215665e-01 -3.278600e+00 -9.496680e+00
```

4. Training/Testing

We now split the samples into a training set and a test set in order to estimate the test error of ridge regression and later on the lasso.

```
set.seed(1)
train <- sample(1:nrow(X), nrow(X) / 2)
test <- (-train)
y.test <- y[test]

# Fit Ridge regression to the training data
ridge.mod <- glmnet(X[train,], y[train], alpha = 0, lambda = grid, thresh = 1e-12)
# Predict the salary to the testing data with lambda = 4
ridge.pred <- predict(ridge.mod, s = 4, newx = X[test,])
# Calculate the Root Mean Square Error (RMSE)
sqrt(mean((ridge.pred - y.test)^2))
```

```
## [1] 377.093
```

```
# Compute the RMSE for the intercept-only model
sqrt(mean((mean(y[train]) - y.test)^2))
```

```
## [1] 473.9936
```

```
# Change to a much larger lambda
ridge.pred <- predict(ridge.mod, s = 1e10, newx = X[test,])
sqrt(mean((ridge.pred - y.test)^2))
```

```
## [1] 473.9935
```

```
# Change lambda to 0
ridge.pred <- predict(ridge.mod, s = 0, newx = X[test,])
sqrt(mean((ridge.pred - y.test)^2))
```

```
## [1] 409.6215
```

```
lm(y ~ X, subset = train)
```

```
##
## Call:
## lm(formula = y ~ X, subset = train)
##
## Coefficients:
## (Intercept)      XAtBat      XHits      XHmRun      XRuns      XRBI
## 274.0145      -0.3521     -1.6377      5.8145      1.5424      1.1243
##      XWalks      XYears      XCatBat      XCHits      XCHmRun      XCRuns
## 3.7287     -16.3773     -0.6412      3.1632      3.4008     -0.9739
##      XCRBI      XCWalks      XLeagueN      XDivisionW      XPutOuts      XAssists
## -0.6005      0.3379     119.1486     -144.0831      0.1976      0.6804
##      XErrors      XNewLeagueN
## -4.7128     -71.0951
```

```
predict(ridge.mod, s = 0, type = "coefficients")[1:20,]
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 274.2089049    -0.3699455    -1.5370022    5.9129307    1.4811980    1.0772844
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
## 3.7577989    -16.5600387    -0.6313336    3.1115575    3.3297885    -0.9496641
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
## -0.5694414      0.3300136    118.4000592   -144.2867510    0.1971770    0.6775088
##      Errors      NewLeagueN
## -4.6833775    -70.1616132
```

Instead of arbitrarily choosing $\lambda = 4$, it would be better to use cross-validation (CV) to choose the tuning parameter λ . We can do this using the built-in cross-validation function, `cv.glmnet()`. By default, the function performs 10-fold cross-validation, though this can be changed using the argument `folds`.

5. Cross-Validation (CV)

```

set.seed(1)
# Fit ridge regression model on training data
cv.out <- cv.glmnet(X[train,], y[train], alpha = 0)
# Select lambda that minimizes training MSE
(bestLambda = cv.out$lambda.min)

```

```
## [1] 326.0828
```


```

ridge.pred <- predict(ridge.mod, s = bestLambda, newx = X[test,])
sqrt(mean((ridge.pred - y.test)^2))

```

```
## [1] 373.9741
```

```
plot(cv.out) # Draw plot of training MSE as a function of lambda
```



Finally, we refit our ridge regression model on the full data set, using the value of λ chosen by cross-validation, and examine the coefficient estimates.

```

# Fit ridge regression model on full dataset
out <- glmnet(X, y, alpha = 0)
# Display coefficients using lambda chosen by CV
predict(out, type = "coefficients", s = bestLambda)[1:20,]

```



```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 15.44383135  0.07715547  0.85911581  0.60103107  1.06369007  0.87936105
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##  1.62444616  1.35254780  0.01134999  0.05746654  0.40680157  0.11456224
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##  0.12116504  0.05299202  22.09143189 -79.04032637  0.16619903  0.02941950
##      Errors      NewLeagueN
## -1.36092945  9.12487767
```


The Lasso

We saw that ridge regression with a wise choice of λ can outperform least squares as well as the null model on the Hitters data set. We now ask whether the lasso, which minimizes

$$\sum_{i=1}^n (y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij})^2 + \lambda \sum_{j=1}^p |\beta_j|$$


can yield either a more accurate or a more interpretable model than ridge regression. In order to fit a lasso model, we once again use the `glmnet()` function; however, this time we use the argument `alpha=1`.

```
# Fit lasso model on training data
lasso.mod <- glmnet(X[train,], y[train], alpha = 1, lambda = grid)
# Draw plot of coefficients
plot(lasso.mod, las = 1)
```



Notice that in the coefficient plot that depending on the choice of tuning parameter, some of the coefficients are exactly equal to zero. We now perform cross-validation and compute the associated test error:

```
set.seed(1)
# Fit lasso model on training data
cv.out <- cv.glmnet(X[train,], y[train], alpha = 1)
# Draw plot of training MSE as a function of lambda
plot(cv.out)
```



```
# Select lamda that minimizes training MSE
bestLambda <- cv.out$lambda.min
# Use best lambda to predict test data
lasso.pred <- predict(lasso.mod, s = bestLambda, newx = X[test,])
# Calculate test RMSE
sqrt(mean((lasso.pred - y[test])^2))
```

```
## [1] 379.043
```

This is substantially lower than the test set RMSE of the null model and of least squares, and very similar to the test RMSE of ridge regression with λ chosen by cross-validation.

However, the lasso has a substantial advantage over ridge regression in that the resulting coefficient estimates are sparse. Here we see that 8 of the 19 coefficient estimates are exactly zero:

```

# Fit lasso model on full dataset
out <- glmnet(X, y, alpha = 1, lambda = grid)
# Display coefficients using lambda chosen by CV
(lasso.coef <- predict(out, type = "coefficients", s = bestLambda)[1:20,])

```

```

##      (Intercept)      AtBat      Hits      HmRun      Runs
##      1.27479059   -0.05497143   2.18034583   0.00000000   0.00000000
##           RBI      Walks      Years      CAtBat      CHits
##      0.00000000   2.29192406  -0.33806109   0.00000000   0.00000000
##      CHmRun      CRuns      CRBI      CWalks      LeagueN
##      0.02825013   0.21628385   0.41712537   0.00000000   20.28615023
##      DivisionW      PutOuts      Assists      Errors      NewLeagueN
## -116.16755870   0.23752385   0.00000000  -0.85629148   0.00000000

```

```

lasso.coef[lasso.coef != 0] # Display only non-zero coefficients

```

```

##      (Intercept)      AtBat      Hits      Walks      Years
##      1.27479059   -0.05497143   2.18034583   2.29192406  -0.33806109
##      CHmRun      CRuns      CRBI      LeagueN      DivisionW
##      0.02825013   0.21628385   0.41712537   20.28615023  -116.16755870
##      PutOuts      Errors
##      0.23752385  -0.85629148

```