

DSA 8020 R Session 10: Random and Mixed Effects Models and Computer Experiments

Whitney

March 18, 2021

Contents

Random Effects Example	1
Read the data into R	1
Fitting a fixed effects model	2
Fitting a random effects model	3
RCBD: Fixed vs. Random Block	3
Load R libraries	3
Read the data	4
Fixed block	4
Random block	4
Computer Experiments	6
Design: Latin hypercube	6
Analysis: Gaussian Process	7


Random Effects Example

Suppose that an agronomist is studying a large number of varieties of soybeans for yield. The agronomist randomly selects three varieties, and then randomly assigns each of those varieties to 10 of 30 available plots.

Model: $y_{ij} = \mu + \alpha_i + \epsilon_{ij}$, $\alpha_i s \stackrel{i.i.d.}{\sim} N(0, \sigma_\alpha^2)$, $\epsilon_{ij} s \stackrel{i.i.d.}{\sim} N(0, \sigma^2)$. $\alpha_i s$ and $\epsilon_{ij} s$ are independent to each other

Read the data into R

```
v1 <- c(6.6, 6.4, 5.9, 6.6, 6.2, 6.7, 6.3, 6.5, 6.5, 6.8)
v2 <- c(5.6, 5.2, 5.3, 5.1, 5.7, 5.6, 5.6, 6.3, 5.0, 5.4)
v3 <- c(6.9, 7.1, 6.4, 6.7, 6.5, 6.6, 6.6, 6.6, 6.8, 6.8)
yield <- c(v1, v2, v3)
var <- factor(c(rep(1, 10), rep(2, 10), rep(3, 10)))
plot(yield ~ var, las = 1)
```



Fitting a fixed effects model

```
fixef <- lm(yield ~ var)
anova(fixef)
```

```
## Analysis of Variance Table
##
## Response: yield
##          Df Sum Sq Mean Sq F value    Pr(>F)
## var         2  8.306   4.1530  49.593 9.114e-10 ***
## Residuals  27  2.261   0.0837
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coefficients(fixef)
```

```
## (Intercept)      var2      var3
##          6.45      -0.97       0.25
```

Fitting a random effects model

```
library(lme4)
```

```
## Loading required package: Matrix
```

```
randef <- lmer(yield ~ 1 + (1|var), REML = TRUE)
summary(randef)
```

```
## Linear mixed model fit by REML ['lmerMod']
## Formula: yield ~ 1 + (1 | var)
##
## REML criterion at convergence: 21.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.8839 -0.6181  0.1118  0.4962  2.7828
##
## Random effects:
##  Groups   Name      Variance Std.Dev.
##  var      (Intercept) 0.40693  0.6379
##  Residual              0.08374  0.2894
## Number of obs: 30, groups:  var, 3
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)   6.2100    0.3721   16.69
```

Let's construct CIs for σ_α^2 , σ^2 , and μ

```
## Compute the confidence intervals (CIs) using profile likelihood
CIs <- confint(randef, oldNames = FALSE)
```

```
## Computing profile confidence intervals ...
```

```
CIs
```

```
##              2.5 %    97.5 %
## sd_(Intercept)|var 0.2637525 1.5512218
## sigma              0.2265053 0.3877781
## (Intercept)       5.3618584 7.0581407
```

RCBD: Fixed vs. Random Block

Load R libraries

```
library(lsmeans)
```

```
## Loading required package: emmeans

## The 'lsmeans' package is now basically a front end for 'emmeans'.
## Users are encouraged to switch the rest of the way.
## See help('transition') for more information, including how to
## convert old 'lsmeans' objects and scripts to work with 'emmeans'.
```

```
library(lmerTest)
```

```
##
## Attaching package: 'lmerTest'

## The following object is masked from 'package:lme4':
##
##      lmer

## The following object is masked from 'package:stats':
##
##      step
```

Read the data

```
### Create the data set
x <- c(52, 47, 44, 51, 42, 60, 55, 49, 52, 43, 56, 48, 45, 44, 38)
trt <- rep(c("A", "B", "C"), each = 5)
blk <- rep(1:5, 3)
dat <- data.frame(x = x, trt = trt, blk = as.factor(blk))
```

Fixed block

```
fixef <- lm(x ~ trt + blk, data = dat)
anova(fixef)
```

```
## Analysis of Variance Table
##
## Response: x
##           Df Sum Sq Mean Sq F value    Pr(>F)
## trt         2   89.2    44.60   7.6239 0.0140226 *
## blk         4  363.6    90.90  15.5385 0.0007684 ***
## Residuals   8   46.8     5.85
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Random block

```
randef <- lmer(x ~ trt + (1|blk), REML = TRUE, data = dat)
summary(randef)
```

```
## Linear mixed model fit by REML. t-tests use Satterthwaite's method [
## lmerModLmerTest]
## Formula: x ~ trt + (1 | blk)
## Data: dat
##
## REML criterion at convergence: 71.1
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.1417 -0.6147 -0.1494  0.5772  1.3390
##
## Random effects:
## Groups Name Variance Std.Dev.
## blk (Intercept) 28.35 5.324
## Residual 5.85 2.419
## Number of obs: 15, groups: blk, 5
##
## Fixed effects:
## Estimate Std. Error df t value Pr(>|t|)
## (Intercept) 47.200 2.615 5.054 18.047 8.76e-06 ***
## trtB 4.600 1.530 8.000 3.007 0.0169 *
## trtC -1.000 1.530 8.000 -0.654 0.5316
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
## (Intr) trtB
## trtB -0.292
## trtC -0.292 0.500
```

```
lsmeans(randef, list(pairwise ~ trt), adjust = "none")
```

```
## $lsmeans of trt'
## trt lsmean SE df lower.CL upper.CL
## A 47.2 2.62 5.05 40.5 53.9
## B 51.8 2.62 5.05 45.1 58.5
## C 46.2 2.62 5.05 39.5 52.9
##
## Degrees-of-freedom method: kenward-roger
## Confidence level used: 0.95
##
## $pairwise differences of trt'
## 1 estimate SE df t.ratio p.value
## A - B -4.6 1.53 8 -3.007 0.0169
## A - C 1.0 1.53 8 0.654 0.5316
## B - C 5.6 1.53 8 3.661 0.0064
##
## Degrees-of-freedom method: kenward-roger
```

Computer Experiments

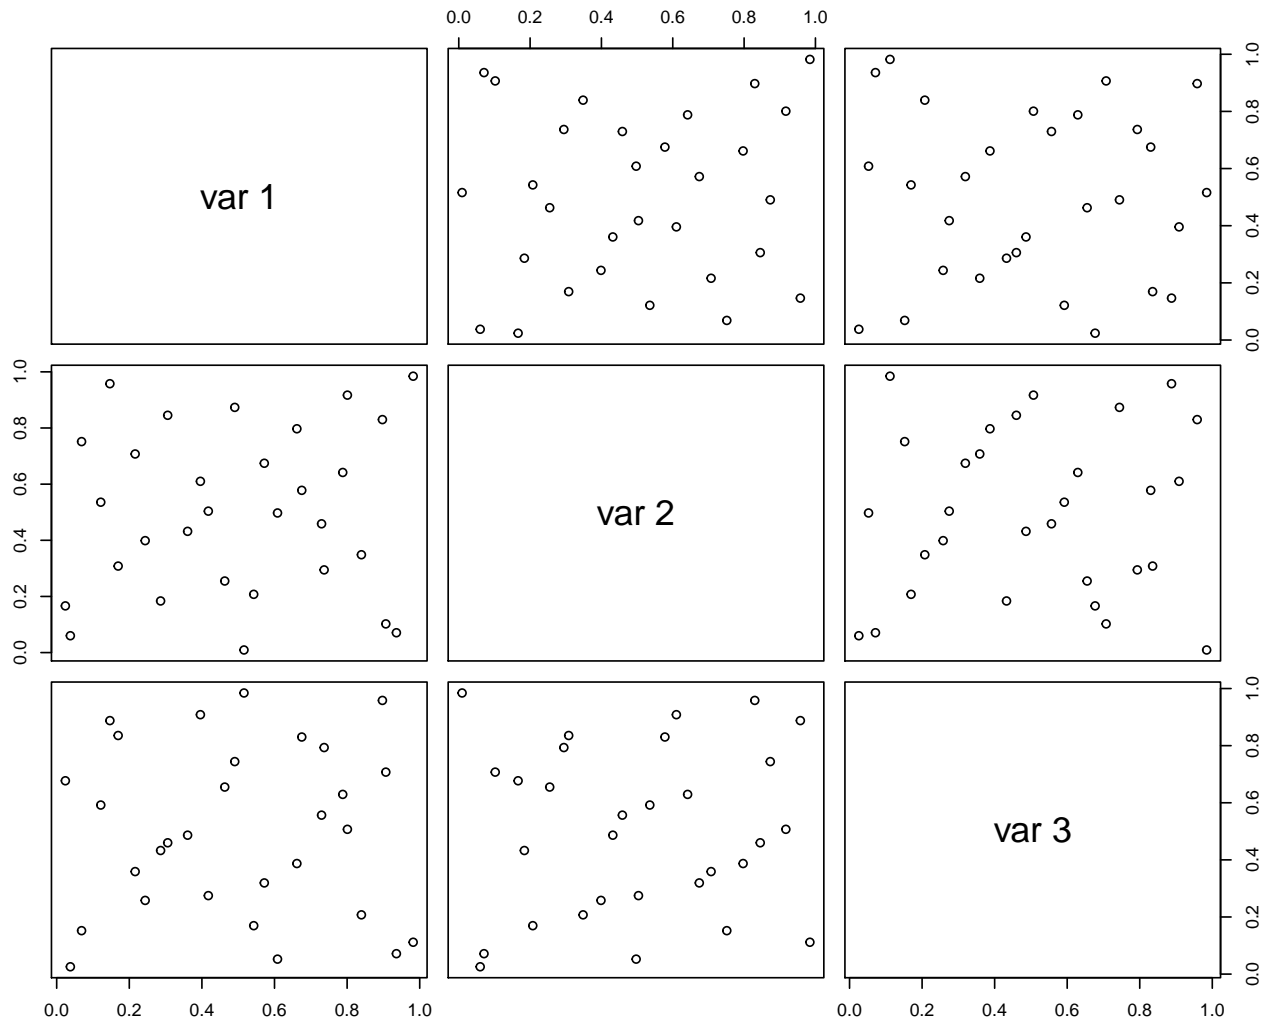
Design: Latin hypercube

```
# install.packages("lhs") # Latin Hypercube Sample Package
library(lhs)
# Generate a good n x k LHD
LHD = maximinLHS(n = 30, k = 3, dup = 5)
# "dup" is an integer tuning parameter that determines the number of
# candidate points considered. Larger values should improve results
# but require more computational resources.

# Display the LHD
LHD
```

```
##           [,1]      [,2]      [,3]
## [1,] 0.66163013 0.797215717 0.38716040
## [2,] 0.21616791 0.707447850 0.35897696
## [3,] 0.41772287 0.503840763 0.27480117
## [4,] 0.73670929 0.294607013 0.79334025
## [5,] 0.36065087 0.431912882 0.48648056
## [6,] 0.78795996 0.641612941 0.62929182
## [7,] 0.28625740 0.183888209 0.43277182
## [8,] 0.16932053 0.308183630 0.83551799
## [9,] 0.67511452 0.577965552 0.83027151
## [10,] 0.49050355 0.873412126 0.74413599
## [11,] 0.54275082 0.207537900 0.16957182
## [12,] 0.83939817 0.348544478 0.20747059
## [13,] 0.12142653 0.535666697 0.59201455
## [14,] 0.57189301 0.674414029 0.31927994
## [15,] 0.46290057 0.254986353 0.65496816
## [16,] 0.90663529 0.102199413 0.70732637
## [17,] 0.06853557 0.751505786 0.15205673
## [18,] 0.80082366 0.917088582 0.50702441
## [19,] 0.14656878 0.957735816 0.88770341
## [20,] 0.24372752 0.398750090 0.25805307
## [21,] 0.60829811 0.497352016 0.05242416
## [22,] 0.39588275 0.610093321 0.90843868
## [23,] 0.72964467 0.458725665 0.55666923
## [24,] 0.02391046 0.166382465 0.67701698
## [25,] 0.98197503 0.984524633 0.11155958
## [26,] 0.89726012 0.829905764 0.95841707
## [27,] 0.03764711 0.060056816 0.02564730
## [28,] 0.51594097 0.009333771 0.98445368
## [29,] 0.93576288 0.070723224 0.07146833
## [30,] 0.30579710 0.845264521 0.45985709
```

```
pairs(LHD)
```



Analysis: Gaussian Process

```
# Load the data
neuron <- read.table("http://deanvosdraguljic.ietsandbox.net/DeanVossDraguljic/R-data/neuron.txt", head=10)
head(neuron, 10)
```

```
##      gNaFsc    gKdrsc fr
## 1  0.38593729 0.2120652 33
## 2  0.04666927 0.4594742  0
## 3  1.00000000 0.4473344 46
## 4  0.95467637 0.3351407 44
## 5  0.53334929 0.7981310 41
## 6  0.59166751 0.6042714 41
## 7  0.18570301 0.3799469 31
## 8  0.49927784 0.2444170 36
## 9  0.74609113 0.3949591 42
## 10 0.07269414 1.0000000  0
```

```

# Fit a GP
library(mlegp)
GPFit <- mlegp(neuron[, 1:2], neuron[, 3])

## no reps detected - nugget will not be estimated
##
## ===== FITTING GP # 1 =====
## running simplex # 1...
## ...done
## ...simplex #1 complete, loglike = -104.446501 (convergence)
## running simplex # 2...
## ...done
## ...simplex #2 complete, loglike = -104.446501 (convergence)
## running simplex # 3...
## ...done
## ...simplex #3 complete, loglike = -104.446502 (convergence)
## running simplex # 4...
## ...done
## ...simplex #4 complete, loglike = -104.446501 (convergence)
## running simplex # 5...
## ...done
## ...simplex #5 complete, loglike = -104.446501 (convergence)
##
## using L-BFGS method from simplex #1...
## iteration: 1,loglike = -104.446501
## ...L-BFGS method complete
##
## Maximum likelihood estimates found, log like = -104.446501
## creating gp object.....done

```

```

summary(GPFit)

##
## Total observations = 30
## Dimensions = 2
##
## mu = 27.61157
## sig2: 251.8751
## nugget: 0
##
## Correlation parameters:
##
##      beta a
## 1  5.027878 2
## 2 50.228477 2
##
## Log likelihood = -104.4465
##
## CV RMSE: 7.312618
## CV RMaxSE: 1020.777

```



```

# Make prediction
predictedX = expand.grid(g_NaF = seq(0, 1, 0.02), g_KDR = seq(0, 1, 0.02))
yhats = predict(GPfit, predictedX, se.fit = T)
# Visualize predictions and their uncertainty
library(fields)
par(mfrow = c(1, 2))
image.plot(seq(0, 1, 0.02), seq(0, 1, 0.02), matrix(yhats$fit, 51, 51),
           xlab = "g NaF (mS/cm^2)", ylab = "g KDR (mS/cm^2)", las = 1,
           main = "Predictions")
points(neuron[, 1:2], pch = 16, cex = 0.75)
image.plot(seq(0, 1, 0.02), seq(0, 1, 0.02), matrix(yhats$se.fit, 51, 51),
           xlab = "g NaF (mS/cm^2)", ylab = "g KDR (mS/cm^2)", las = 1,
           main = "Predictions Uncertainty")
points(neuron[, 1:2], pch = 16, cex = 0.75)

```

