Lecture 5

Multivariate Linear Regression

Readings: Johnson & Wichern 2007, Chapter 7; DSA 8020 Lectures 1-4 [Link]; Zelterman, 2015, Chapter 9

DSA 8070 Multivariate Analysis

Whitney Huang
Clemson University

Notes

Agenda

1. Motivation
2. Model and Assumptions
3. Parameter Estimation
4. Inference and Prediction
5. Motor Trend Car Road Tests Data Analysis

Notes

Example: Motor Trend Car Road Tests

```
> head(mtcars)
                   mpg cyl disp  hp drat    wt  qsec vs am gear carb
Mazda RX4         21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
Mazda RX4 Wag     21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
Datsun 710        22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
Hornet 4 Drive    21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
Valiant           18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

Suppose we would like to study the (linear) relationship between `mpg`, `disp`, `hp`, `wt` (responses) and `cyl`, `am`, `carb` (predictors)

Multiple regression predicts one outcome; multivariate regression predicts several simultaneously

Notes

Why Multivariate Regression Instead of Separate Regressions?

- **Estimation:**
  - Coefficient estimates are the same as running separate regressions
- **Inference:** The real gain comes from joint modeling
  - Test hypotheses across multiple outcomes simultaneously
  - Accounts for correlations among responses $\Rightarrow$ more powerful and accurate tests
- **Examples:**
  - Does a predictor affect all outcomes jointly?
  - Multivariate analog of ANOVA (MANOVA)

Notes

Review: Linear Regression Model

The multiple linear regression model has the form:

$$y_i = \beta_0 + \sum_{j=1}^{p} \beta_j x_{ij} + \varepsilon_i, \quad i = 1, \cdots, n,$$

where

- y_i is the response for the i-th observation
- x_{ij} is the j-th predictor for the i-th observation
- β_0 and β_j's are the regression intercept and slopes for the response, respectively
- ε_i is the error term for the response of the i-th observation

Notes

The Multivariate Linear Regression Model: Scalar Form

The multivariate (multiple) linear regression model has the form:

$$y_{ik} = \beta_{0k} + \sum_{j=1}^{p} \beta_{jk} x_{ij} + \varepsilon_{ik}, \quad i = 1, \cdots, n, \quad k = 1, \cdots, d,$$

where

- y_{ik} is the k-th response for the i-th observation
- x_{ij} is the j-th predictor for the i-th observation
- β_{0k} and β_{jk}'s are the regression intercept and slopes for k-th response, respectively
- ε_{ik} is the error term for the k-th response of the i-th observation

Notes

The Multivariate Linear Regression Model: Assumptions

The assumptions of the model are:

- Relationship between $\{x_j\}_{j=1}^p$ and Y_k is linear for each $k \in \{1, \cdots, d\}$
- $(\varepsilon_{i1}, \cdots, \varepsilon_{id})^T \overset{i.i.d.}{\sim} \mathrm{N}(\mathbf{0}, \Sigma)$ is an unobserved random vector
- $[Y_{ik}|x_{i1}, \cdots, x_{ip}] \sim \mathrm{N}(\beta_{0k} + \sum_{j=1}^p \beta_{jk} x_{ij}, \sigma_{kk})$ for each $k \in \{1, \cdots, d\}$

Notes

The Multivariate Linear Regression Model: Matrix Form

The multivariate multiple linear regression model has the form

$$\boldsymbol{Y} = \boldsymbol{X}\boldsymbol{B} + \boldsymbol{E},$$

where

- $\boldsymbol{Y} = [\boldsymbol{y}_1, \cdots, \boldsymbol{y}_d]$ is the $n \times d$ response matrix, where $\boldsymbol{y}_k = (y_{1k}, \cdots, y_{nk})^T$ is the k-th response vector
- $\boldsymbol{X} = [\mathbf{1}, \boldsymbol{x}_1, \cdots, \boldsymbol{x}_p]$ is the $n \times (p+1)$ design matrix
- $\boldsymbol{B} = [\boldsymbol{\beta}_1, \cdots, \boldsymbol{\beta}_d]$ is the $(p+1) \times d$ matrix of regression coefficients
- $\boldsymbol{E} = [\boldsymbol{\varepsilon}_1, \cdots, \boldsymbol{\varepsilon}_d]$ is the $n \times d$ error matrix

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

Notes

Another Look of the Matrix Form

Matrix form writes the multivariate linear regression model for all $n \times d$ points simultaneously as

$$\boldsymbol{Y} = \boldsymbol{X}\boldsymbol{B} + \boldsymbol{E}$$

$$\begin{bmatrix} y_{11} & \cdots & y_{1d} \\ y_{21} & \cdots & y_{2d} \\ \vdots & \ddots & \vdots \\ y_{n1} & \cdots & y_{nd} \end{bmatrix} = \begin{bmatrix} 1 & \cdots & x_{1p} \\ 1 & \cdots & x_{2p} \\ \vdots & \ddots & \vdots \\ 1 & \cdots & x_{np} \end{bmatrix} \begin{bmatrix} \beta_{01} & \cdots & \beta_{0d} \\ \beta_{11} & \cdots & \beta_{1d} \\ \vdots & \ddots & \vdots \\ \beta_{p1} & \cdots & \beta_{pd} \end{bmatrix} + \begin{bmatrix} \varepsilon_{11} & \cdots & \varepsilon_{1d} \\ \varepsilon_{21} & \cdots & \varepsilon_{2d} \\ \vdots & \ddots & \vdots \\ \varepsilon_{n1} & \cdots & \varepsilon_{nd} \end{bmatrix}$$

Assuming that n subjects are independent, we have

- $\boldsymbol{\varepsilon}_k \sim \mathrm{N}(0, \sigma_{kk}), \quad k \in \{1, \cdots, d\}$
- $\boldsymbol{\varepsilon}_i \overset{i.i.d.}{\sim} \mathrm{N}(\mathbf{0}, \Sigma), \quad i = 1, \cdots, n$

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

Notes

Ordinary Least Squares

The ordinary least squares OLS estimate is

$$\underset{\boldsymbol{B}\in\mathbb{R}^{(p+1)\times d}}{\operatorname{argmin}} ||\boldsymbol{Y}-\boldsymbol{X}\boldsymbol{B}||^2 = \underset{\boldsymbol{B}\in\mathbb{R}^{(p+1)\times d}}{\operatorname{argmin}} \sum_{i=1}^{n}\sum_{k=1}^{d}\left(y_{ik}-\beta_{0k}-\sum_{j=1}^{p}\beta_{jk}x_{ij}\right)^2,$$

where $||\cdot||$ denotes the Frobenius norm.

- $\mathrm{OLS}(\boldsymbol{B}) = ||\boldsymbol{Y}-\boldsymbol{X}\boldsymbol{B}||^2 = \mathrm{tr}(\boldsymbol{Y}^T Y) - 2\mathrm{tr}(\boldsymbol{Y}^T\boldsymbol{X}\boldsymbol{B}) + \mathrm{tr}(\boldsymbol{B}^T\boldsymbol{X}^T\boldsymbol{X}\boldsymbol{B})$
- $\frac{\partial \mathrm{OLS}(\boldsymbol{B})}{\partial \boldsymbol{B}} = -2\boldsymbol{X}^T\boldsymbol{Y} + 2\boldsymbol{X}^T\boldsymbol{X}\boldsymbol{B}$

The OLS estimate has the form

$$\hat{\boldsymbol{B}} = (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{Y} \Rightarrow \hat{\boldsymbol{\beta}}_k = (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{y}_k, \quad k\in\{1,\cdots,d\}$$

Notes

Expected Value of Least Squares Coefficients

The expected value of the estimated coefficients is given by

$$\begin{aligned}\mathbb{E}(\hat{\boldsymbol{B}}) &= \mathbb{E}\left[(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{Y}\right]\\ &= (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\mathbb{E}(\boldsymbol{Y})\\ &= (\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{X}\boldsymbol{B}\\ &= \boldsymbol{B}\end{aligned}$$

$\Rightarrow \hat{\boldsymbol{B}}$ is an unbiased estimator of $\boldsymbol{B}$

Notes

Fitted Values and Residuals

- Fitted values are given by

$$\hat{\boldsymbol{Y}} = \boldsymbol{X}\hat{\boldsymbol{B}},$$

  i.e., $\hat{y}_{ik} = \hat{\beta}_{0k} + \sum_{j=1}^{p}\hat{\beta}_{jk}x_{ij}, \quad i=1,\cdots,n, \quad k=1,\cdots,d$

- Residuals are given by

$$\hat{\boldsymbol{E}} = \boldsymbol{Y} - \hat{\boldsymbol{Y}},$$

  i.e., $\hat{\varepsilon}_{ik} = y_{ik} - \hat{y}_{ik}, \quad i=1,\cdots,n, \quad k=1,\cdots,d$

Notes

Hat Matrix

Just like in univariate linear regression we can write the fitted values as

$$\begin{aligned}\hat{\boldsymbol{Y}} &= \boldsymbol{X}\hat{\boldsymbol{B}}\\ &= \boldsymbol{X}(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T\boldsymbol{Y}\\ &= \boldsymbol{H}\boldsymbol{Y},\end{aligned}$$

where $\boldsymbol{H} = \boldsymbol{X}(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{X}^T$ is the hat matrix

$\Rightarrow$ $\boldsymbol{H}$ projects $\boldsymbol{y}_k$ onto the column space of $\boldsymbol{X}$ for $k \in \{1, \cdots, d\}$

Notes

Partitioning the Total Variation

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

We can partition the total covariation in $\{\boldsymbol{y}_i\}_{i=1}^n$ ($\mathrm{SSCP}_{\mathrm{Tot}}$)as

$$\begin{aligned}\mathrm{SSCP}_{\mathrm{tot}} &= \sum_{i=1}^n (\boldsymbol{y}_i - \bar{\boldsymbol{y}})^T(\boldsymbol{y}_i - \bar{\boldsymbol{y}})\\ &= \sum_{i=1}^n (\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i + \hat{\boldsymbol{y}}_i - \bar{\boldsymbol{y}})(\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i + \hat{\boldsymbol{y}}_i - \bar{\boldsymbol{y}})^T\\ &= \underbrace{\sum_{i=1}^n (\hat{\boldsymbol{y}}_i - \bar{\boldsymbol{y}})(\hat{\boldsymbol{y}}_i - \bar{\boldsymbol{y}})^T}_{\mathrm{SSCP}_{\mathrm{Reg}}} + \underbrace{\sum_{i=1}^n (\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i)(\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i)^T}_{\mathrm{SSCP}_{\mathrm{Err}}}\\ &\quad + \underbrace{2\sum_{i=1}^n (\hat{\boldsymbol{y}}_i - \bar{\boldsymbol{y}})(\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i)}_{=0}\\ &= \mathrm{SSCP}_{\mathrm{Reg}} + \mathrm{SSCP}_{\mathrm{Err}}\end{aligned}$$

The corresponding degrees of freedom are $d(n-1)$ for $\mathrm{SSCP}_{\mathrm{Tot}}$; dp for $\mathrm{SSCP}_{\mathrm{Reg}}$; and $d(n-p-1)$ for $\mathrm{SSCP}_{\mathrm{Err}}$

Notes

Estimated Error Covariance

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

The estimated error covariance matrix is

$$\begin{aligned}\hat{\boldsymbol{\Sigma}} &= \frac{\sum_{i=1}^n (\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i)(\boldsymbol{y}_i - \hat{\boldsymbol{y}}_i)^T}{n-p-1}\\ &= \frac{\mathrm{SSCP}_{Err}}{n-p-1}\end{aligned}$$

- $\hat{\boldsymbol{\Sigma}}$ is an unbiased estimate of $\boldsymbol{\Sigma}$
- The estimate $\hat{\boldsymbol{\Sigma}}$ is the mean SSCP_{Err}

Notes

Sampling Distributions of $\hat{B}$, $\hat{Y}$, and $\hat{E}$

We would need to figure out the sampling distributions of estimator and predictor in order to drawn inference

Given the model assumptions, we have

$$\mathrm{vec}(\hat{\boldsymbol{B}}) \sim \mathrm{N}(\mathrm{vec}(\boldsymbol{B}), \boldsymbol{\Sigma} \otimes (\boldsymbol{X}^T\boldsymbol{X})^{-1})$$
$$\mathrm{vec}(\hat{\boldsymbol{Y}}) \sim \mathrm{N}(\mathrm{vec}(\boldsymbol{X}\boldsymbol{B}), \boldsymbol{\Sigma} \otimes \boldsymbol{H})$$
$$\mathrm{vec}(\hat{\boldsymbol{E}}) \sim \mathrm{N}(\boldsymbol{0}, \boldsymbol{\Sigma} \otimes (\boldsymbol{I} - \boldsymbol{H})),$$

where $\mathrm{vec}(\cdot)$ is the vectorization operator and $\otimes$ is the Kronecker product

Notes

Inference about Multiple $\hat{\beta}_{jk}$

Assume that $q < p$ and want to test if a reduced model is sufficient:

$$H_0 : \boldsymbol{B}_2 = \boldsymbol{0}_{p-q} \times d, \quad \text{versus} \quad H_a : \boldsymbol{B}_2 \neq \boldsymbol{0}_{p-q} \times d,$$

where

$$\boldsymbol{B} = \begin{bmatrix} \boldsymbol{B}_1 \\ \boldsymbol{B}_2 \end{bmatrix}$$

is the partitioned of the coefficient vector

We can compare the SSCP_{Err} for the full model:

$$y_{ik} = \beta_{0k} + \sum_{j=1}^{p} \beta_{jk} x_{ij} + \varepsilon_{ik}, \quad k - 1, \cdots, d$$

and the reduced model:

$$y_{ik} = \beta_{0k} + \sum_{j=1}^{q} \beta_{jk} x_{ij} + \varepsilon_{ik}, \quad k - 1, \cdots, d$$

Notes

Some Test Statistics

Let $\tilde{\boldsymbol{E}} = n\tilde{\boldsymbol{\Sigma}}$ denote the SSCP_{Err} matrix from the full model, and let $\tilde{\boldsymbol{H}} = n\left(\tilde{\boldsymbol{\Sigma}}_1 - \tilde{\boldsymbol{\Sigma}}\right)$ denote the hypothesis SSCP_{Err} matrix

Some test statistics for

$$H_0 : \boldsymbol{B}_2 = \boldsymbol{0}_{p-q} \times d, \quad \text{versus} \quad H_a : \boldsymbol{B}_2 \neq \boldsymbol{0}_{p-q} \times d :$$

- Wilks Lambda
$$\Lambda^* = \frac{|\tilde{\boldsymbol{E}}|}{|\tilde{\boldsymbol{H}} + \tilde{\boldsymbol{E}}|}$$
Reject H_0 if Λ^* is "small"
- Hotelling-Lawley Trace
$$T_0^2 = \mathrm{tr}(\tilde{\boldsymbol{H}}\tilde{\boldsymbol{E}}^{-1})$$
Reject H_0 if T_0^2 is "large"
- Pillai Trace
$$V = \mathrm{tr}(\tilde{\boldsymbol{H}}(\tilde{\boldsymbol{H}} + \tilde{\boldsymbol{E}})^{-1})$$
Reject H_0 if V is "large"

Notes

Interval Estimation

We would like to estimate the expected value of the response for a given predictor $\boldsymbol{x}_h = (1, x_{h1}, \cdots, x_{hp})$.

Note that we have

$$\hat{\boldsymbol{y}}_h \sim \mathrm{N}(\boldsymbol{B}^T\boldsymbol{x}_h, \boldsymbol{x}_h^T(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{x}_h\boldsymbol{\Sigma})$$

We can exploit the duality between interval estimation and hypothesis testing. That is, we can test

$$H_0 : \mathbb{E}(\boldsymbol{y}_h) = \boldsymbol{y}_h^* \text{ versus } H_a : \mathbb{E}(\boldsymbol{y}_h) \neq \boldsymbol{y}_h^*$$

> The $100(1-\alpha)$% confidence region is the collection of $\boldsymbol{y}_h^*$ values that fail to reject H_0 at α level

Notes

Interval Estimation (Cont'd)

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

Test statistics:

$$T^2 = \left(\frac{\hat{\boldsymbol{B}}^T\boldsymbol{x}_h - \boldsymbol{B}^T\boldsymbol{x}_h}{\sqrt{\boldsymbol{x}_h^T(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{x}_h}}\right)^T \hat{\boldsymbol{\Sigma}}^{-1} \left(\frac{\hat{\boldsymbol{B}}^T\boldsymbol{x}_h - \boldsymbol{B}^T\boldsymbol{x}_h}{\sqrt{\boldsymbol{x}_h^T(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{x}_h}}\right)$$
$$\overset{H_0}{\sim} \frac{d(n-p-1)}{n-p-d} F_{d,n-p-d}$$

Therefore, the $100(1-\alpha)$% simultaneous confidence interval for y_{hk} is

$$\hat{y}_{hk} \pm \sqrt{\frac{d(n-p-1)}{n-p-d} F_{d,n-p-d}} \sqrt{\boldsymbol{x}_h^T(\boldsymbol{X}^T\boldsymbol{X})^{-1}\boldsymbol{x}_h\hat{\sigma}_{kk}},$$

$k \in \{1, \cdots, d\}$

Notes

Predicting New Observations

Multivariate Linear Regression
CLEMSON UNIVERSITY
Motivation
Model and Assumptions
Parameter Estimation
Inference and Prediction
Motor Trend Car Road Tests Data Analysis

Here we want to predict the observed value of response for a given predictor

- **Note**: interested in actual $\hat{\boldsymbol{y}}_h$ instead of $\mathbb{E}(\hat{\boldsymbol{y}}_h)$
- Given $\boldsymbol{x}_h = (1, x_{h1}, \cdots, x_{hp})$, the fitted value is still $\hat{\boldsymbol{y}}_h = \hat{\boldsymbol{B}}^T\boldsymbol{x}_h$

We can exploit the duality between interval estimation and hypothesis testing. That is, we can test

$$H_0 : \boldsymbol{y}_h = \boldsymbol{y}_h^* \text{ versus } H_a : \boldsymbol{y}_h \neq \boldsymbol{y}_h^*$$

> The $100(1-\alpha)$% prediction interval is the collection of $\boldsymbol{y}_h^*$ values that fail to reject H_0 at α level

Notes

Predicting New Observations (Cont'd)

Test statistics:

$$T^2 = \left(\frac{\hat{\boldsymbol{B}}^T \boldsymbol{x}_h - \boldsymbol{B}^T \boldsymbol{x}_h}{\sqrt{1 + \boldsymbol{x}_h^T (\boldsymbol{X}^T \boldsymbol{X})^{-1} \boldsymbol{x}_h}} \right)^T \hat{\boldsymbol{\Sigma}}^{-1} \left(\frac{\hat{\boldsymbol{B}}^T \boldsymbol{x}_h - \boldsymbol{B}^T \boldsymbol{x}_h}{\sqrt{1 + \boldsymbol{x}_h^T (\boldsymbol{X}^T \boldsymbol{X})^{-1} \boldsymbol{x}_h}} \right)$$

$$\overset{H_0}{\sim} \frac{d(n-p-1)}{n-p-d} F_{d,n-p-d}$$

Therefore, the $100(1-\alpha)\%$ simultaneous prediction interval for y_{hk} is

$$\hat{y}_{hk} \pm \sqrt{\frac{d(n-p-1)}{n-p-d} F_{d,n-p-d}} \sqrt{\left(1 + \boldsymbol{x}_h^T (\boldsymbol{X}^T \boldsymbol{X})^{-1} \boldsymbol{x}_h\right) \hat{\sigma}_{kk}},$$

$$k \in \{1, \cdots, d\}$$

Notes

Motor Trend Car Road Tests Data Analysis

Study the linear relationship between `mpg, disp, hp, wt` (responses) and `cyl, am, carb` (predictors) in the `mtcars` dataset

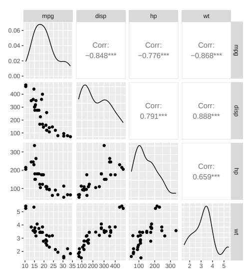

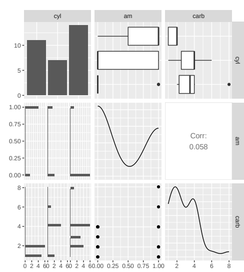

Notes

Multivariate Regression Fit

- Model: `lm(Y ~ cyl + am + carb, data = mtcars)`
- Key findings:
  - `mpg`: decreases with more cylinders, increases with manual transmission
  - `disp`: strongly influenced by `cyl`
  - `hp`: influenced by `cyl8` and `carb`
  - `wt`: influenced by `cyl8`, `am`, `carb`
- **Note:** Multivariate regression produces the *same point estimates* as running separate regressions for each response

Notes

SSCP & Error Covariance

SSCP decomposition:

$$SSCP_{Tot} = SSCP_{Reg} + SSCP_{Err}$$

- Estimated error covariance matrix ($\hat{\Sigma}$)

```
             mpg        disp          hp          wt
mpg    7.8680094  -53.27166 -19.7015979 -0.6575443
disp -53.2716607 2504.87095 425.1328988 18.1065416
hp   -19.7015979  425.13290 577.2703337  0.4662491
wt    -0.6575443   18.10654   0.4662491  0.2573503
            mpg        disp          hp          wt
mpg   1.0000000 -0.3794645 -0.29233405 -0.46209388
disp -0.3794645  1.0000000  0.35354314  0.71314929
hp   -0.2923340  0.3535431  1.00000000  0.03825304
wt   -0.4620939  0.7131493  0.03825304  1.00000000
```

$\Rightarrow$ Captures residual dependencies among responses

Notes

Do We Need "cyl"?

```
> mvlm0 <- lm(Y ~ am + carb, data = mtcars)
> anova(mvlm, mvlm0, test = "Wilks")
Analysis of Variance Table

Model 1: Y ~ cyl + am + carb
Model 2: Y ~ am + carb
  Res.Df Df Gen.var.   Wilks approx F num Df
1     27       29.862
2     29  2    43.692 0.16395   8.8181      8
  den Df    Pr(>F)
1
2     48 2.525e-07 ***
```

Interpretation: cylinder count explains variation across responses

Notes

Confidence and Prediction Intervals

Default `predict()` lacks multivariate CI/PI, so we use an `R` function from Prof. Helwig at the University of Minnesota

```
> newdata <- data.frame(cyl = factor(6, levels = c(4, 6, 8)),
+                        am = 1, carb = 4)
> # confidence interval
> pred.mlm(mvlm, newdata)
         mpg     disp        hp       wt
fit 21.51824 159.2707 136.98500 2.631108
lwr 16.65593  72.5141  95.33649 1.751736
upr 26.38055 246.0273 178.63351 3.510479
> # prediction interval
> pred.mlm(mvlm, newdata, interval = "prediction")
          mpg      disp        hp        wt
fit 21.518240 159.27070 136.98500 2.6311076
lwr  9.680053 -51.95435  35.58397 0.4901152
upr 33.356426 370.49576 238.38603 4.7720999
```


Notes

Summary

In this lecture, we learned about Multivariate Linear Regression

- Model and Assumptions
- Parameter Estimation
- Inference and Prediction

Notes

Notes

Notes