

DSA 8070 R Session 6: Reduced-Rank Regression and Repeated Measures Analysis

Whitney Huang, Clemson University

Contents

Reduced-Rank Regression	1
Multivariate Regression Refresher	1
Reduced-Rank Regression (RRR): SVD-Based Estimator	2
Rank Selection via Cross-Validation	3
Compare OLS vs. RRR (Predictive Performance)	5
Inspect Low-Rank Structure	6
Repeated Measures Analysis	7
Dog Experiment	7
Split-plot ANOVA	7
MANOVA	8
Mixed Model with AR(1) temporal correlation structure	8

Reduced-Rank Regression

Multivariate Regression Refresher

We model a multivariate response matrix $Y \in \mathbb{R}^{n \times q}$ with predictors $X \in \mathbb{R}^{n \times p}$:

$$Y = XB + E,$$

where $B \in \mathbb{R}^{p \times q}$ is the coefficient matrix.

A naive approach fits each response column separately (OLS for each). But this **ignores cross-response correlation**.

```
set.seed(123)
n <- 200; p <- 6; q <- 4

# Design matrix with intercept
X <- cbind(1, matrix(rnorm(n * (p - 1)), n, p - 1))
colnames(X) <- c("Intercept", paste0("X", 1:(p - 1)))

# True low-rank coefficient: rank = 2
r_true <- 2
```

```

A <- matrix(rnorm(p * r_true), p, r_true)
C <- matrix(rnorm(r_true * q), r_true, q)
B_true <- A %*% C # p x q

# Correlated errors across responses
Sigma_eps <- matrix(c(1.0, 0.6, 0.3, 0.2,
                      0.6, 1.0, 0.5, 0.4,
                      0.3, 0.5, 1.0, 0.7,
                      0.2, 0.4, 0.7, 1.0), q, q)

library(MASS)
E <- MASS::mvrnorm(n, mu = rep(0, q), Sigma = 2 * Sigma_eps)

Y <- X %*% B_true + E
colnames(Y) <- paste0("Y", 1:q)

# Fit separate OLS for each response (baseline)
ols_fits <- vector("list", length = q)
for(j in 1:q){
  ols_fits[[j]] <- lm(Y[, j] ~ X - 1) # lm adds intercept
}
names(ols_fits) <- colnames(Y)
#Compare the estimate and the true parameter
B_est <- array(unlist(lapply(ols_fits, coef)), dim = c(p, q))

B_est

```

```

##           [,1]      [,2]      [,3]      [,4]
## [1,] 1.29633691 0.3063739 -0.5678854 0.1460612
## [2,] 2.32378304 0.6533733 -3.0533527 4.0340122
## [3,] 0.36436374 0.1528000 -0.8591614 1.1869485
## [4,] 0.07241715 0.1164822 0.6869218 -0.7811949
## [5,] 4.43669300 1.1061847 -4.1532537 4.1316973
## [6,] 0.03901217 0.4363386 -3.4027784 5.2680089

```

B_true

```

##           [,1]      [,2]      [,3]      [,4]
## [1,] 1.33072102 0.25707090 -0.5417158 0.1524423
## [2,] 2.40068397 0.67164166 -3.2856474 3.9295441
## [3,] 0.34329260 0.13196539 -0.8687432 1.1934462
## [4,] -0.04855480 -0.05568266 0.5339427 -0.8195902
## [5,] 4.41898000 1.06179603 -4.1101136 4.1652634
## [6,] 0.06024535 0.30995319 -3.3372124 5.2514408

```

Reduced-Rank Regression (RRR): SVD-Based Estimator

Idea: Fit full multivariate OLS, then project fitted values to rank r via SVD.

Algorithm:

1. $\hat{B}^{OLS} = (X^\top X)^{-1} X^\top Y$
2. $\hat{Y} = X \hat{B}^{OLS}$

3. SVD: $\hat{Y} = U\Sigma V^\top$

4. Keep top r right singular vectors V_r , set $\hat{B}^{RRR} = \hat{B}^{OLS}V_rV_r^\top$

```
RRR_fit <- function(X, Y, r){
  # X: n x p, Y: n x q
  XtX <- crossprod(X)
  ridge <- 1e-8 * diag(ncol(X))
  B_ols <- solve(XtX + ridge, crossprod(X, Y)) # p x q
  Yhat <- X %*% B_ols # n x q
  svdY <- svd(Yhat)
  Vr <- svdY$v[, 1:r, drop = FALSE] # q x r
  B_rrr <- B_ols %*% Vr %*% t(Vr) # p x q
  out <- list(B_ols = B_ols, B_rrr = B_rrr, Vr = Vr, svd = svdY)
  return(out)
}

predict_multivar <- function(Xnew, B){
  Xnew %*% B
}
```

Try a few ranks and compare in-sample fit:

```
r_grid <- 0:min(ncol(X), ncol(Y)) # from 0 up to min(p, q)
RSS_by_r <- matrix(NA_real_, nrow = length(r_grid), ncol = 2)
colnames(RSS_by_r) <- c("r", "RSS")

for(i in seq_along(r_grid)){
  r <- r_grid[i]
  if(r == 0){
    B <- matrix(0, nrow = ncol(X), ncol = ncol(Y))
  } else {
    B <- RRR_fit(X, Y, r)$B_rrr
  }
  Yhat <- X %*% B
  RSS_by_r[i, ] <- c(r, sum((Y - Yhat)^2))
}
RSS_by_r
```

```
##      r      RSS
## [1,] 0 26482.308
## [2,] 1  4007.351
## [3,] 2  1495.498
## [4,] 3  1479.386
## [5,] 4  1473.253
```

Rank Selection via Cross-Validation

We evaluate predictive performance for candidate ranks using K-fold CV on rows (observations).

```
cv_rrr <- function(X, Y, ranks, K = 5, seed = 1){
  set.seed(seed)
```

```

n <- nrow(X)
folds <- sample(rep(1:K, length.out = n))
# store (r, fold, mse)
res <- matrix(NA_real_, nrow = length(ranks) * K, ncol = 3)
colnames(res) <- c("r", "fold", "mse")
row_idx <- 1
for(rr in ranks){
  for(k in 1:K){
    idx_te <- which(folds == k)
    idx_tr <- setdiff(seq_len(n), idx_te)
    Xtr <- X[idx_tr, , drop = FALSE]
    Ytr <- Y[idx_tr, , drop = FALSE]
    Xte <- X[idx_te, , drop = FALSE]
    Yte <- Y[idx_te, , drop = FALSE]
    if(rr == 0){
      B <- matrix(0, nrow = ncol(X), ncol = ncol(Y))
    } else {
      B <- RRR_fit(Xtr, Ytr, rr)$B_rrr
    }
    Ypred <- Xte %*% B
    mse <- mean((Yte - Ypred)^2)
    res[row_idx, ] <- c(rr, k, mse)
    row_idx <- row_idx + 1
  }
}
return(res)
}

ranks <- 0:min(ncol(X), ncol(Y))
cv_res <- cv_rrr(X, Y, ranks = ranks, K = 5, seed = 123)

# summarize by rank
cv_summary <- matrix(NA_real_, nrow = length(ranks), ncol = 2)
colnames(cv_summary) <- c("r", "cv_mse")
for(i in seq_along(ranks)){
  r <- ranks[i]
  cv_summary[i, 1] <- r
  cv_summary[i, 2] <- mean(cv_res[cv_res[,1]==r, 3])
}
cv_summary

```

```

##      r    cv_mse
## [1,] 0 33.102885
## [2,] 1  5.069399
## [3,] 2  1.958754
## [4,] 3  1.996958
## [5,] 4  1.994387

```

Plot CV curve and select \hat{r} :

```

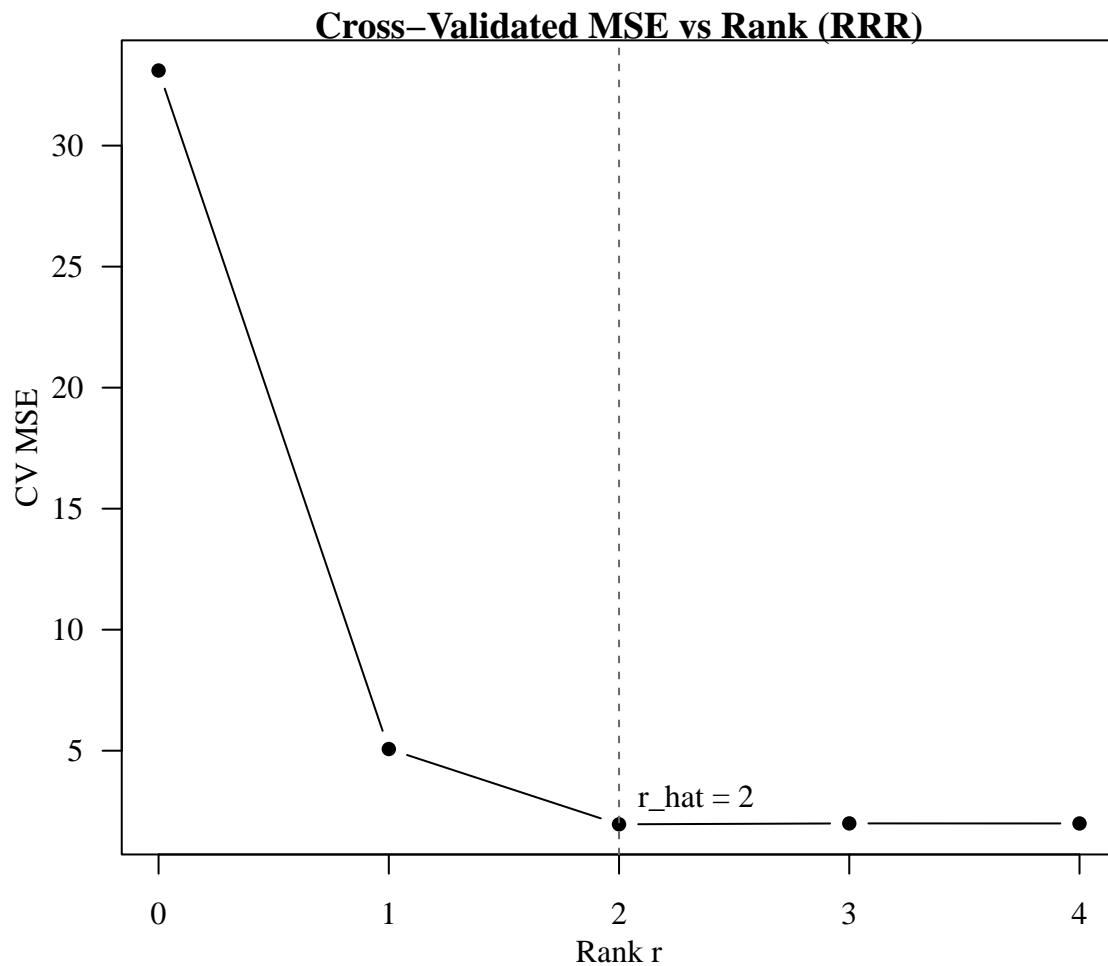
par(mar = c(3, 3.5, 0.8, 0.6), mgp = c(2, 1, 0), las = 1,
    family = "serif")
plot(cv_summary[,1], cv_summary[,2], type = "b", xlab = "Rank r",

```

```

ylab = "CV MSE", main = "Cross-Validated MSE vs Rank (RRR)",
pch = 16)
r_hat <- cv_summary[which.min(cv_summary[,2]), 1]
abline(v = r_hat, lty = 2, col = "gray40")
text(r_hat, min(cv_summary[,2]) + 1, labels = paste("r_hat =", r_hat), pos = 4)

```



```
r_hat
```

```
## r
## 2
```

Compare OLS vs. RRR (Predictive Performance)

Hold out a test set to compare OLS and RRR predictions.

```

set.seed(1234)
test_idx <- sample(1:n, size = round(0.3 * n))
train_idx <- setdiff(1:n, test_idx)

Xtr <- X[train_idx, , drop = FALSE]
Ytr <- Y[train_idx, , drop = FALSE]

```

```

Xte <- X[test_idx, , drop = FALSE]
Yte <- Y[test_idx, , drop = FALSE]

# OLS (multivariate): fit B_ols on train
B_ols_tr <- solve(crossprod(Xtr) + 1e-8 * diag(ncol(Xtr)), crossprod(Xtr, Ytr))
Yhat_ols <- Xte %*% B_ols_tr
mse_ols <- mean((Yte - Yhat_ols)^2)

# RRR (rank = r_hat) on train
r_use <- max(1, r_hat) # ensure at least 1
B_rrr_tr <- RRR_fit(Xtr, Ytr, r = r_use)$B_rrr
Yhat_rrr <- Xte %*% B_rrr_tr
mse_rrr <- mean((Yte - Yhat_rrr)^2)

res_compare <- rbind(
  cbind(Method = "OLS (separate)", Test_MSE = mse_ols),
  cbind(Method = paste0("RRR (r=", r_use, ")"), Test_MSE = mse_rrr)
)
res_compare

##      Method      Test_MSE
## [1,] "OLS (separate)" "1.90507096441179"
## [2,] "RRR (r=2)"      "1.89022281117133"

```

Inspect Low-Rank Structure

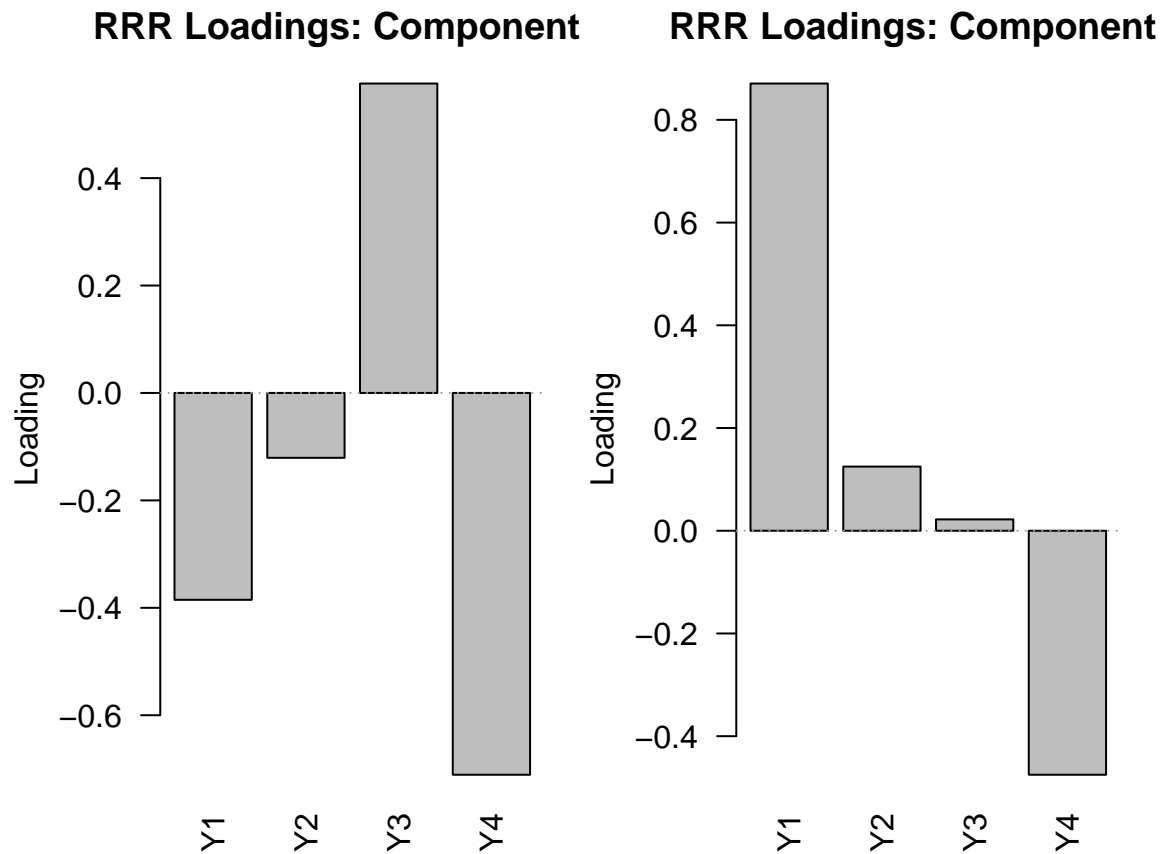
Visualize the right singular vectors (response-space directions) that RRR uses.

```

fit_full <- RRR_fit(X, Y, r = r_use)
Vr <- fit_full$Vr # q x r_use

# Barplot of loadings by response for each component
op <- par(mfrow = c(1, max(1, ncol(Vr))), mar = c(4, 4, 3, 1))
for(j in 1:ncol(Vr)){
  barplot(Vr[, j],
    names.arg = colnames(Y),
    las = 2,
    main = paste("RRR Loadings: Component", j),
    ylab = "Loading")
  abline(h = 0, col = "gray60", lty = 3)
}

```



```
par(op)
```

Repeated Measures Analysis

Dog Experiment

```
dat <- read.table("dog1.txt")
temp <- array(dim = c(144, 4))
temp[, 1] <- rep(dat$V1, 4)
temp[, 2] <- rep(dat$V2, 4)
temp[, 3] <- rep(c(1, 5, 9, 13), each = 36)
temp[, 4] <- c(dat$V3, dat$V4, dat$V5, dat$V6)
dat2 <- data.frame(temp)
names(dat2) <- c("Treatment", "Dog_id", "Time", "Response")
dat2$Treatment <- as.factor(dat2$Treatment)
dat2$Dog_id <- as.factor(dat2$Dog_id)
dat2$Time <- as.factor(dat2$Time)
```

Split-plot ANOVA

```
# computing the cell means (by treatment and time combinations)
tapply(dat2$Response, list(dat2$Treatment, dat2$Time), mean)
# interaction plot
```

```

par(las = 1, mgp = c(2.2, 1, 0), mar = c(3.6, 3.6, 0.8, 0.6))
with(dat2, interaction.plot(x.factor = Time, trace.factor = Treatment,
                           response = Response, col = 1:4, lwd = 1.5))

# Split-plot anova
library(lmerTest)
fit <- lmer(Response ~ Treatment * Time + (1 | Dog_id), data = dat2)
fit
anova(fit)

```

MANOVA

```

out <- manova(cbind(V3, V4, V5, V6) ~ as.factor(V1), data = dat)
summary(out, test = "Wilks")

```

```

##              Df  Wilks approx F num Df den Df  Pr(>F)
## as.factor(V1)  3 0.48452    2.022    12 77.018 0.03316 *
## Residuals      32
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
summary(out)
```

```

##              Df Pillai approx F num Df den Df  Pr(>F)
## as.factor(V1)  3 0.5978    1.9286    12   93 0.04048 *
## Residuals      32
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Mixed Model with AR(1) temporal correlation structure

```

library(nlme)
fit1 = gls(Response ~ Treatment * Time,
            correlation = corCompSymm(form = ~ 1 | Dog_id), data = dat2)
fit1

```

```

## Generalized least squares fit by REML
## Model: Response ~ Treatment * Time
## Data: dat2
## Log-restricted-likelihood: -119.9032
##
## Coefficients:
##      (Intercept)      Treatment2      Treatment3      Treatment4
##      4.111111111      -0.511111111      -0.466666667      -0.571111111
##           Time5           Time9           Time13  Treatment2:Time5
##      0.28888889      0.955555556      0.611111111      0.311111111
## Treatment3:Time5 Treatment4:Time5 Treatment2:Time9 Treatment3:Time9
##      0.077777778      -0.208888889      -0.055555556      -0.622222222
## Treatment4:Time9 Treatment2:Time13 Treatment3:Time13 Treatment4:Time13
##      -0.835555556      0.013888889      -0.211111111      -0.691111111

```



```
##
## Correlation Structure: Compound symmetry
## Formula: ~1 | Dog_id
## Parameter estimate(s):
##      Rho
## 0.5538616
## Degrees of freedom: 144 total; 128 residual
## Residual standard error: 0.6446676
```

```
fit2 = gls(Response ~ Treatment * Time,
            correlation = corAR1(form = ~ 1 | Dog_id), data = dat2)
fit2
```

```
## Generalized least squares fit by REML
## Model: Response ~ Treatment * Time
## Data: dat2
## Log-restricted-likelihood: -120.7906
##
## Coefficients:
##      (Intercept)      Treatment2      Treatment3      Treatment4
##      4.11111111      -0.51111111      -0.46666667      -0.57111111
##      Time5          Time9          Time13      Treatment2:Time5
##      0.28888889      0.95555556      0.61111111      0.31111111
## Treatment3:Time5 Treatment4:Time5 Treatment2:Time9 Treatment3:Time9
##      0.07777778      -0.20888889      -0.05555556      -0.62222222
## Treatment4:Time9 Treatment2:Time13 Treatment3:Time13 Treatment4:Time13
##      -0.83555556      0.01388889      -0.21111111      -0.69111111
##
## Correlation Structure: AR(1)
## Formula: ~1 | Dog_id
## Parameter estimate(s):
##      Phi
## 0.5928708
## Degrees of freedom: 144 total; 128 residual
## Residual standard error: 0.6376364
```

```
anova(fit1, fit2)
```

```
##      Model df      AIC      BIC    logLik
## fit1      1 18 275.8063 327.1429 -119.9032
## fit2      2 18 277.5811 328.9177 -120.7906
```