

MATH 8090: ARMA Models: Prediction and Forecasting, Modeling Case Study

Whitney Huang, Clemson University

9/30-10/2/2025

Contents

NOAA wind data example	2
Load and plot the data	2
“Estimate” ϕ using sample ACF and center the data	2
One-step-ahead forecast	5
Fill in missing value example	6
Simulate an AR(-0.9)	6
Let’s remove some data to illustrate how to fill in missing values using forecasting algorithm	7
Fill in “missing” values	8
Prediction Errors from Best Linear Predictor	9
Ireland wind data case study	10
Load and plot the data	10
Deseasonalization: Harmonic Regression	11
ACF Plots: Original and Deseasonalized Series	12
Apply transformation to make wind speed more Gaussian like	13
Now take square roots of the original data and deseasonalizeagain!	14
Checking Normality ACF/PACF	15
Model identification, fitting, and selection	16
Let’s first fit an AR(1)	16
Fit an AR(2) model	19
Fit an ARMA(1,1) model	22
Use AIC to conduct model selection	28
Forecasting	29
Visualizing the Forecasts	30

NOAA wind data example

This example is taken from Don Percival's time series course (UW Stat 519).


The one-step-ahead forecast of an AR(1) process is:

$$P_n X_{n+1} = \hat{\mu} + \hat{\phi}(X_n - \hat{\mu}),$$

where $\hat{\phi}$ is our estimate of ϕ , and $\hat{\mu}$ is an estimate of μ .

Load and plot the data

```
ws <- scan("http://faculty.washington.edu/dbp/s519/Data/wind-speed.txt")
n <- length(ws)
par(las = 1, mgp = c(2, 1, 0), mar = c(3.6, 3.6, 1.4, 0.6), family = "serif")
plot(ws, col = "blue", xlab = "Time", typ = "b",
     ylab = expression(x[t]), main = "Wind Speed Time Series", cex = 0.5)
grid()
xbar_ws <- mean(ws)
abline(h = xbar_ws, lty = 2)
```



“Estimate” ϕ using sample ACF and center the data

```

# --- Quick AR(1) coefficient from lag-1 ACF -----
acf.ws <- acf(ws, lag.max = 40, plot = FALSE)$acf # includes lag 0
phi.ws <- acf.ws[2] # lag-1 sample ACF as a quick AR(1) phi estimate

# --- Bartlett variance factor for AR(1): w_hh(phi) -----
# Var( rho_hat[h] ) w_hh(phi) / n
gen.whh.ar <- function(h, phi){
  p2 <- phi^2
  p2h <- p2^h
  -2 * h * p2h + (1 - p2h) * (1 + p2) / (1 - p2)
}

# --- Plot Sample ACF, AR(1) Model ACF, and 95% CIs (Bartlett & IID) -----
plot.ACFbartlettAR <- function(ts, n.lags = 40){
  stopifnot(is.numeric(ts), length(ts) > 5)
  n.ts <- length(ts)
  lags <- 1:n.lags

  # Sample ACF for lags 1..n.lags (remove lag 0)
  acf.full <- acf(ts, lag.max = n.lags, plot = FALSE)$acf
  acf.est <- acf.full[-1]

  # AR(1) model ACF using phi_hat = sample acf at lag 1
  phi.hat <- acf.est[1]
  acf.model <- phi.hat^lags

  # Plot the sample ACF (stems + points)
  plot(lags, acf.est, type = "h",
       xlab = "h (lag)", ylab = "ACF",
       ylim = c(-1, 1),
       main = "Model & Sample ACFs with 95% Confidence Bounds",
       las = 1)
  points(lags, acf.est, pch = 8) # star

  # Add AR(1) model ACF
  points(lags, acf.model, col = "red")

  # Bartlett 95% CI under AR(1)
  CI.AR <- 1.96 * sqrt(sapply(lags, function(h) gen.whh.ar(h, phi.hat))) / sqrt(n.ts)
  lines(lags, acf.est + CI.AR, col = "red", lty = 2)
  lines(lags, acf.est - CI.AR, col = "red", lty = 2)

  # IID 95% CI for comparison: ±1.96/sqrt(n)
  CI.IID <- rep(1.96 / sqrt(n.ts), n.lags)
  lines(lags, -CI.IID, col = "gray", lty = 2)
  lines(lags, CI.IID, col = "gray", lty = 2)

  abline(h = 0, lty = "dashed")
  legend("bottomleft",
       legend = c("IID 95% CI", "AR(1) 95% CI", "AR(1) model"),
       lty = c(2, 2, NA),
       pch = c(NA, NA, 1),
       col = c("gray", "red", "red"),


```

```

    bty      = "n")
}

# --- Plot styling-----
par(mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.4, 0.6), family = "serif")
plot.ACFbartlettAR(ws)

```



```

# --- Alternatively: estimate phi via MLE -----
phi_hat_mle <- arima(ws, order = c(1, 0, 0)) # AR(1) fit by MLE
phi_hat_mle

```

```

##
## Call:
## arima(x = ws, order = c(1, 0, 0))
##
## Coefficients:
##      ar1  intercept
##      0.906   -1.1136
## s.e.  0.037    0.6035
##
## sigma^2 estimated as 0.4615:  log likelihood = -132.99,  aic = 271.99

```


```

# --- Center the series if you need a mean-zero version -----
ws.centered <- ws - mean(ws)

```

One-step-ahead forecast

```
# -- One-step-ahead predictions and errors -----
# Predictions available for t = 2..n using x_{t-1}
ws.hat <- phi.ws * ws.centered[1:(n - 1)] + xbar_ws      # \hat{x}_t
zt.ws   <- ws.hat - ws[2:n]                             # e_t = \hat{x}_t - x_t
## plot it
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.2, 0.6))
plot(ws, col = "blue", xlab = "Time", type = "b", ylab = expression(x[t]),
     main = "One-Step-Ahead Prediction", cex = 0.75)
points(2:n, ws.hat, pch = 8, col = "red", cex = 0.375)
lines(2:n, zt.ws, col = "gray")
legend("bottomright", legend = expression(x[t], hat(x)[t], hat(x)[t] - x[t]),
      col = c("blue", "red", "gray"), pch = c(1, 8, NA),
      lty = c(NA, NA, "solid"), pt.cex = c(0.75, 0.375, 1), inset = 0.01,
      bty = "n")
```



```
# -- Variance comparison (smaller error variance suggests better predictive fit)
cat("\nQuick AR(1) estimate (lag-1 ACF): phi_hat =", round(phi.ws, 4), "\n")
```

```
##
## Quick AR(1) estimate (lag-1 ACF): phi_hat = 0.8911
```

```
cat("Var(prediction errors) =", round(var(zt.ws), 4), "\n")
```

```
## Var(prediction errors) = 0.4629
```


```
cat("Var(series)          =", round(var(ws), 4), "\n")
```

```
## Var(series)          = 2.5025
```

Fill in missing value example

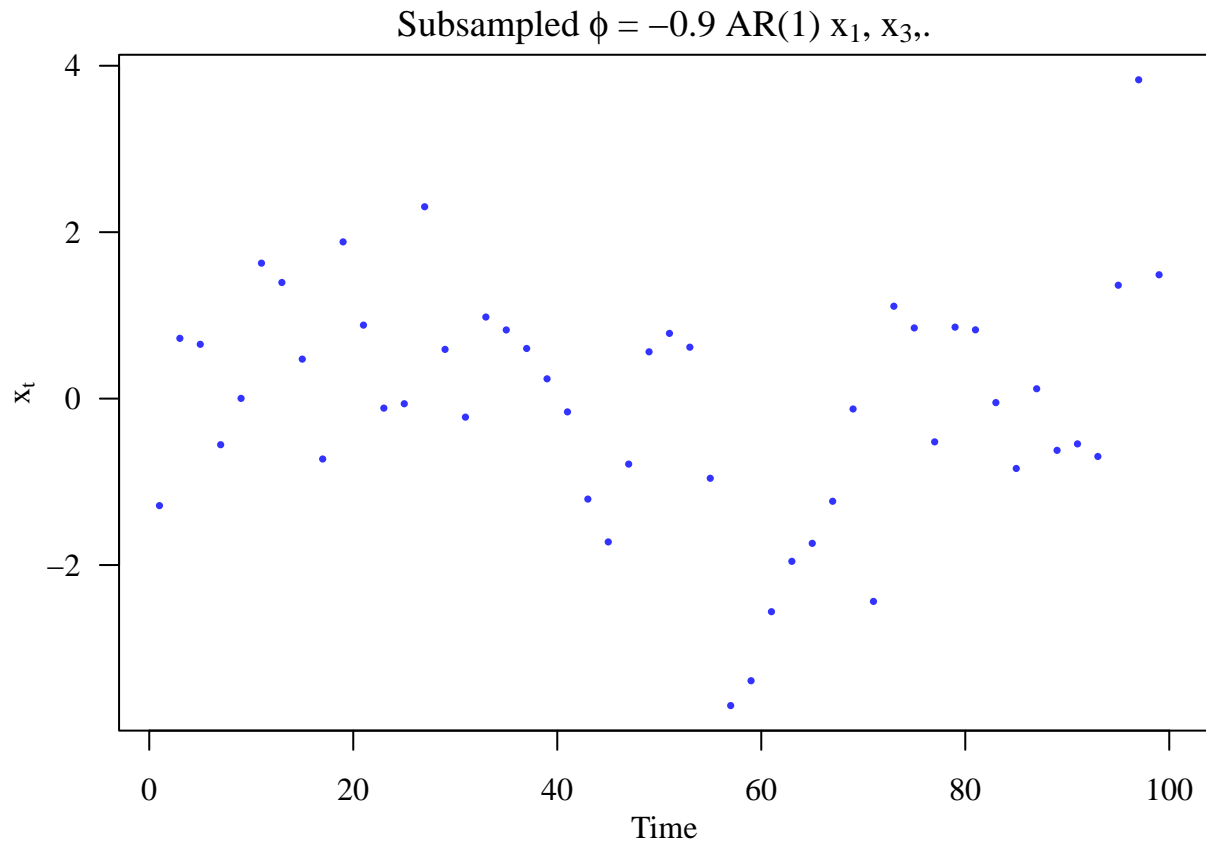
Simulate an AR(-0.9)

```
generate.AR1.ts <- function(phi = 0.0){  
  ts <- rep(0, 100)  
  ts[1] <- rnorm(1) / sqrt(1 - phi^2)  
  for(i in 2:100) ts[i] <- phi * ts[i - 1] + rnorm(1)  
  ts  
}  
set.seed(123)  
ar1.ts <- generate.AR1.ts(-0.9)  
  
library(scales)  
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.4, 0.6), family = "serif")  
plot(ar1.ts, col = alpha("blue", 0.4), xlab = "Time", type = "l",  
     ylab = expression(x[t]), cex = 0.5)  
points(ar1.ts, pch = 16, cex = 0.75, col = "blue")  
grid()
```



Let's remove some data to illustrate how to fill in missing values using forecasting algorithm

```
ar1.ts.subsampled <- ar1.ts
ar1.ts.subsampled[seq(2, 100, 2)] <- NA
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.4, 0.6), family = "serif")
plot(ar1.ts.subsampled, xlab = "Time", type = "b", ylab = expression(x[t]),
     main = expression(paste("Subsampled ", phi, " = -0.9 AR(1) ", x[1], ", ", "x[3], ", ".)),
     cex = 0.5, col = alpha("blue", 0.8), pch = 16)
```



Fill in “missing” values

$$\hat{X}_2 = \phi(X_1 + X_3)/(1 + \phi^2)$$

$$\text{MSPE} = \frac{\sigma^2}{1 + \phi^2}$$

```
ar1.ts.predicted <- ar1.ts
# Fill in even indices (2, 4, ..., 98) with prediction based on neighbors.
# Formula comes from conditional expectation of AR(1) with phi = -0.9
ar1.ts.predicted[seq(2, 98, 2)] <-
  -0.9 * (ar1.ts[seq(1, 97, 2)] + ar1.ts[seq(3, 99, 2)]) / 1.81

# Last element (t=100) set to NA since it has no neighbor on the right
ar1.ts.predicted[100] <- NA


# Plot styling
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.4, 0.6), family = "serif")

# Base plot: predicted series in light blue
plot(ar1.ts.predicted, col = alpha("blue", 0.4), xlab = "Time", type = "l",
     ylab = expression(x[t]), cex = 0.5)

# Mark predicted even-indexed points (red)
xs <- seq(2, 98, 2)
points(xs, ar1.ts.predicted[xs], pch = 19, col = "red", cex = 0.5)
```




```
# Mark observed odd-indexed points (blue)
xo <- seq(1, 99, 2)
points(xo, ar1.ts.subsampled[xo], col = "blue", cex = 0.75, pch = 16)
```



Prediction Errors from Best Linear Predictor

```
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1.4, 0.6), family = "serif")
plot(xs, (ar1.ts.predicted - ar1.ts)[xs], col = "blue", xlab = "Time", type = "p",
     ylab = expression(x[t]), main = "Prediction Errors from Best Linear Predictor",
     cex=0.5)
abline(h = 0, lty = 2)
```




Ireland wind data case study

Load and plot the data

In this case study, we use the data at the Rosslare station from 1965 to 1969.

```
library(gstat)
data(wind)
id <- which(wind$year %in% 65:69)
rosslare <- wind$ROS[id]
## set up the year variable
year <- seq(from = 1965, by = 1 / 365.25, length = length(rosslare))
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), mgp = c(2, 1, 0), las = 1)
plot(year, rosslare, type = "l", ylim = c(0, 35), lwd = 0.6,
      xlab = "year", ylab = "Wind speed (knots)")
grid()
```



Deseasonalization: Harmonic Regression


We use harmonic regression with 4 harmonics per year to model the seasonal components.

```
## create harmonic terms
Harmonic <- function(year, K){
  t <- outer(2 * pi * year, 1:K)
  return(cbind(apply(t, 2, cos), apply(t, 2, sin)))
}
harmonics <- Harmonic(year, 4)
## fit a harmonic regression
harm.model <- lm(rosslare ~ harmonics)
summary(harm.model)
```

```
##
## Call:
## lm(formula = rosslare ~ harmonics)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.8538  -3.3813  -0.4892   2.8395  20.8290
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 11.584141   0.112377 103.083 < 2e-16 ***
## harmonics1   1.687468   0.158936  10.617 < 2e-16 ***
```

```
## harmonics2 -0.435273 0.158936 -2.739 0.00623 **
## harmonics3 -0.060047 0.158936 -0.378 0.70562
## harmonics4 -0.251396 0.158936 -1.582 0.11388
## harmonics5 0.412363 0.158915 2.595 0.00954 **
## harmonics6 0.003874 0.158915 0.024 0.98055
## harmonics7 0.107245 0.158915 0.675 0.49985
## harmonics8 0.217870 0.158915 1.371 0.17055
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.802 on 1817 degrees of freedom
## Multiple R-squared: 0.06771, Adjusted R-squared: 0.06361
## F-statistic: 16.5 on 8 and 1817 DF, p-value: < 2.2e-16
```

```
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), mgp = c(2, 1, 0), las = 1)
plot(year, rosslare, type = "l",
      xlab = "year", ylab = "Wind speed (knots)", col = "grey")
lines(year, fitted(harm.model), lwd = 2)
```



ACF Plots: Original and Deseasonalized Series


Let's plot the ACF and PACF plots to investigate the possible order for the ARMA model.

```
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
## method          from
## as.zoo.data.frame zoo
```





```
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), mgp = c(2, 1, 0), las = 1,
    mfrow = c(2, 1))
Acf(rosslare, lag.max = 365, xlab = "", ylab = "Sample ACF", main = "")
Acf(resid(harm.model), lag.max = 365, xlab = "Lag in days",
    ylab = "Sample ACF", main = "")
```



Apply transformation to make wind speed more Gaussian like

Now look at a histogram of the values, along with the normal quantile-quantile plot.

```
par(mfrow = c(2, 2), bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1,
    mgp = c(2.2, 1, 0))
hist(rosslare, 40, main = "", xlab = "Wind speed", col = "lightblue", border = "gray")
qqnorm(rosslare, main = "")
qqline(rosslare)
## Histogram/Q-Q plot of 1/2 root transformation
hist(sqrt(rosslare), 25, main = "", xlab = "sqrt(Wind speed)", col = "lightblue", border = "gray")
qqnorm(sqrt(rosslare), main = ""); qqline(sqrt(rosslare))
```




Now take square roots of the original data and deseasonalizeagain!

```
## now we start again from the beginning with a sqrt transformation
sqrt.rosslare <- sqrt(rosslare)
## refit the periodicity, without the intercept term
harm.model <- lm(sqrt.rosslare ~ harmonics[, 1:4] - 1)
summary(harm.model)
```

```
##
## Call:
## lm(formula = sqrt.rosslare ~ harmonics[, 1:4] - 1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
##  1.146   2.848   3.316   3.799   5.656
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## harmonics[, 1:4]1  0.2391111  0.1126203   2.123  0.0339 *
## harmonics[, 1:4]2 -0.0606520  0.1126203  -0.539  0.5903
## harmonics[, 1:4]3 -0.0001588  0.1126203  -0.001  0.9989
## harmonics[, 1:4]4 -0.0363877  0.1126202  -0.323  0.7467
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

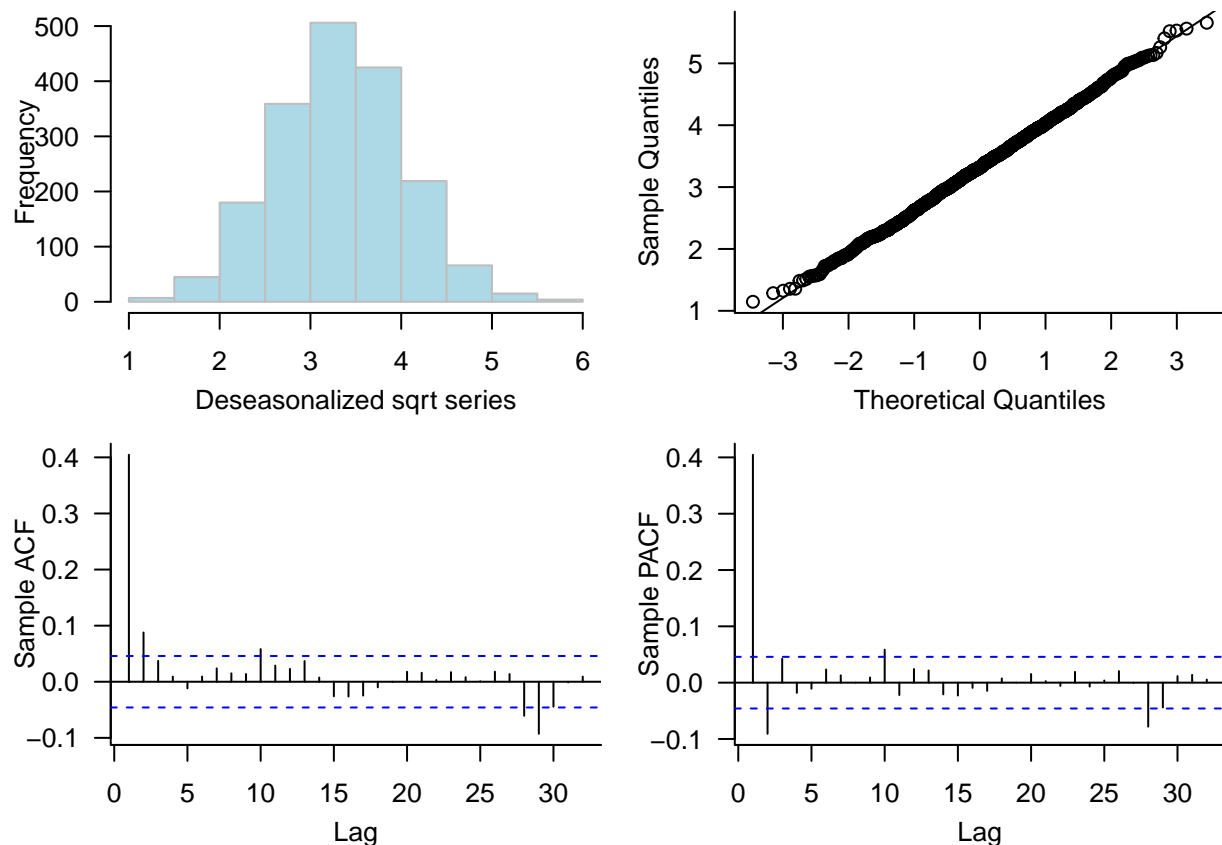
```
##
## Residual standard error: 3.403 on 1822 degrees of freedom
## Multiple R-squared:  0.002684,    Adjusted R-squared:  0.0004944
## F-statistic: 1.226 on 4 and 1822 DF,  p-value: 0.2978

## calculate the estimate of the deseasonalized series
sqrt.rosslare.ds <- resid(harm.model)
## Produce time series plots of the sqrt data
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0))
plot(year, sqrt.rosslare.ds, type = "l", col = "gray",
      xlab = "year", ylab = "Deaseas. sqrt series")
lines(year, fitted(harm.model) + mean(sqrt.rosslare.ds), lwd = 2)
```



Checking Normality ACF/PACF

```
## And check the distribution
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0),
      mfrow = c(2, 2))
hist(sqrt.rosslare.ds, main = "", xlab = "Deseasonalized sqrt series",
      col = "lightblue", border = "gray")
qqnorm(sqrt.rosslare.ds, main="")
qqline(sqrt.rosslare.ds)
## Now let's examine the sample ACF and PACF
Acf(sqrt.rosslare.ds, main = "", ylab = "Sample ACF")
Acf(sqrt.rosslare.ds, main = "", type = "partial", ylab = "Sample PACF")
```



Model identification, fitting, and selection

Let's first fit an **AR(1)** Fit an AR(1) model

```
ar1.model <- arima(sqrt.rosslare.ds, order = c(1, 0, 0))
```

Summarize the fitted model

```
ar1.model
```


```
##
## Call:
## arima(x = sqrt.rosslare.ds, order = c(1, 0, 0))
##
## Coefficients:
##          ar1  intercept
##      0.4060    3.3257
## s.e.  0.0214    0.0254
##
## sigma^2 estimated as 0.4148:  log likelihood = -1787.72,  aic = 3581.43
```

Sample and fitted ACF


```

par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0))
Acf(sqrt.rosslare.ds, main = "")
acf_true <- ARMAacf(ar = c(ar1.model$coef[1]), lag.max = 32)[-1]
points(1:32, acf_true, col = "red")

```




Extract residuals

```


ar1.resids <- resid(ar1.model)
## time series plot of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0),
    mfrow = c(1, 2))
plot(year, ar1.resids, type = "l", xlab = "year", ylab = "AR1 residuals",
     lwd = 0.6, col = "gray")
abline(h = 0, lty = 2)
## Normal Q-Q plot for the residuals
qqnorm(ar1.resids, main = "", cex = 0.75); qqline(ar1.resids)

```



Sample ACF and PACF of the residuals

```
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.4, 1, 0), mfrow = c(1, 2))
Acf(ar1.resids, ylab = "Sample ACF", xlab = "lag (day)", main = "")
Acf(ar1.resids, ylab = "Sample PACF", type = "partial", xlab = "lag (day)")
```



```
## Carry out the Ljung-Box test
Box.test(ar1.resids, lag = 32, type = "Ljung-Box", fitdf = 1)
```

```
##
## Box-Ljung test
##
## data: ar1.resids
## X-squared = 53.142, df = 31, p-value = 0.00794
```

```
(ar2.model <- arima(sqrt.rosslare.ds, order = c(2, 0, 0)))
```


Fit an AR(2) model

```
##
## Call:
## arima(x = sqrt.rosslare.ds, order = c(2, 0, 0))
##
## Coefficients:
##          ar1          ar2  intercept
##          0.4425  -0.0905      3.3254
## s.e.  0.0233   0.0233    0.0232
##
## sigma^2 estimated as 0.4114:  log likelihood = -1780.23,  aic = 3568.46
```

```

par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0))
acf(sqrt.rosslare.ds, main = "")
acf_true <- ARMAacf(ar = c(ar2.model$coef[1:2]), lag.max = 32)[-1]
points(1:32, acf_true, col = "red")

```




```

## extract the residuals
ar2.resids <- resid(ar2.model)


## time series plot of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0), mfrow = c(1, 2))
plot(year, ar2.resids, type = "l", xlab = "year",
     ylab = "AR(2) residuals", lwd = 0.6, col = "gray")
abline(h = 0, lty = 2)

## Normal Q-Q plot for the residuals
qqnorm(ar2.resids, main = "", cex = 0.75); qqline(ar2.resids)

```



```
## Sample ACF and PACF of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.4, 1, 0), mfrow = c(1, 2))
Acf(ar2.resids, ylab = "Sample ACF", xlab = "lag (day)", main = "")
pacf(ar2.resids, ylab = "Sample PACF", xlab = "lag (day)")
```



```
## Carry out the Ljung-Box test
Box.test(ar2.resids, lag = 32, type = "Ljung-Box", fitdf = 2)
```

```
##
## Box-Ljung test
##
## data: ar2.resids
## X-squared = 36.548, df = 30, p-value = 0.1907
```

```
(arma11.model <- arima(sqrt.rosslare.ds, order = c(1, 0, 1)))
```


Fit an ARMA(1,1) model

```
##
## Call:
## arima(x = sqrt.rosslare.ds, order = c(1, 0, 1))
##
## Coefficients:
##          ar1      ma1  intercept
##      0.1978  0.2502      3.3254
## s.e.  0.0556  0.0553      0.0234
##
## sigma^2 estimated as 0.4108: log likelihood = -1778.82, aic = 3565.64
```

```

par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0))
acf(sqrt.rosslare.ds, main = "")
acf_true <- ARMAacf(ar = arma11.model$coef[1], ma = arma11.model$coef[2], lag.max = 32)[-1]
points(1:32, acf_true, col = "red")

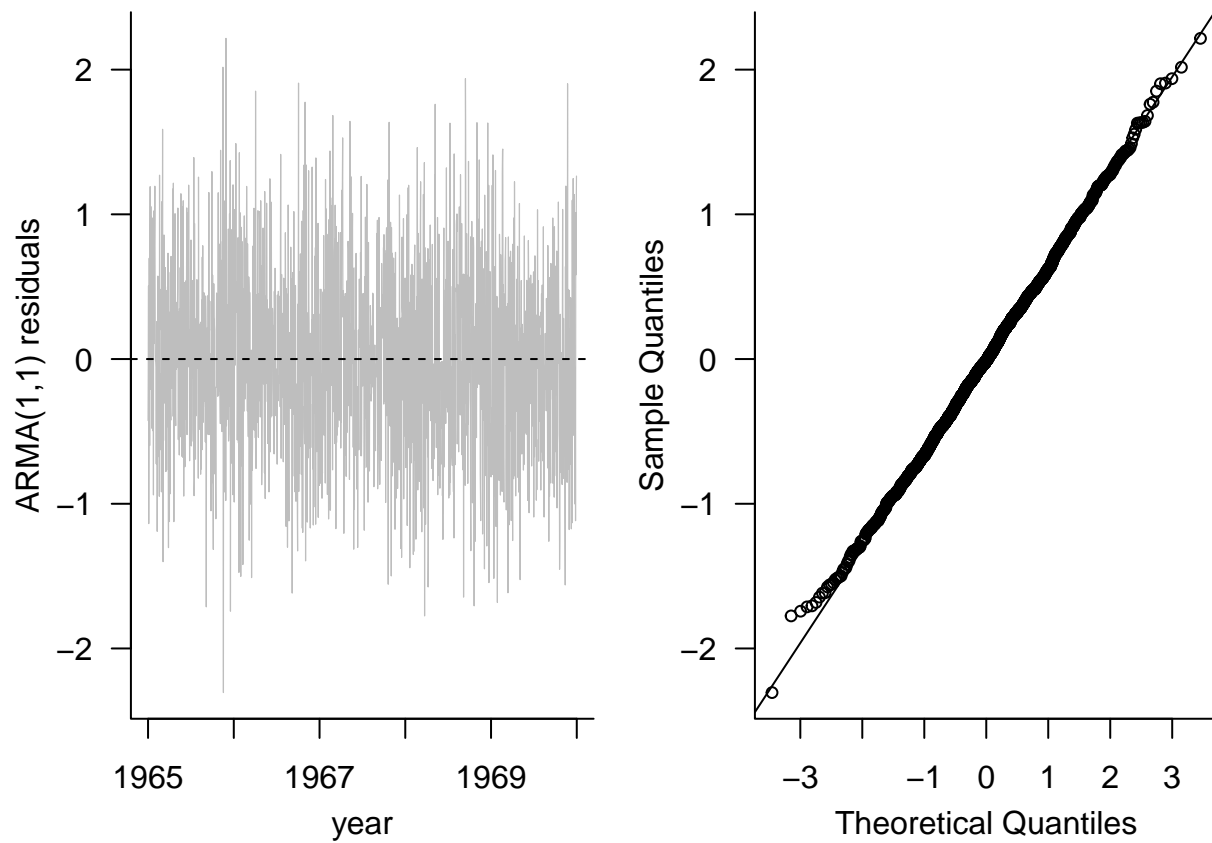
```




```

## extract the residuals
arma11.resids <- resid(arma11.model)
## time series plot of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0), mfrow = c(1, 2))
plot(year, arma11.resids, type = "l", xlab = "year",
     ylab = "ARMA(1,1) residuals", lwd = 0.6, col = "gray")
abline(h = 0, lty = 2)
## Normal Q-Q plot for the residuals
qqnorm(arma11.resids, main = "", cex = 0.75); qqline(arma11.resids)

```



```
## Sample ACF and PACF of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.4, 1, 0),
    mfrow = c(1, 2))
Acf(arma11.resids, ylab = "Sample ACF", xlab = "lag (day)", main = "")
pacf(arma11.resids, ylab = "Sample PACF", xlab = "lag (day)")
```

```
## Carry out the Ljung-Box test
Box.test(arma11.resids, lag = 32, type = "Ljung-Box", fitdf = 2)
```

```
##
## Box-Ljung test
##
## data: arma11.resids
## X-squared = 32.757, df = 30, p-value = 0.3332
```

```
(arma21.model <- arima(sqrt.rosslare.ds, order = c(2, 0, 1)))
```


Fit an ARMA(2,1) model

```
##
## Call:
## arima(x = sqrt.rosslare.ds, order = c(2, 0, 1))
##
## Coefficients:
##          ar1          ar2          ma1  intercept
##          0.0703  0.0587  0.3768      3.3253
## s.e.  0.1691  0.0772  0.1663      0.0237
##
## sigma^2 estimated as 0.4107:  log likelihood = -1778.56,  aic = 3567.11
```

```

par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0))
Acf(sqrt.rosslare.ds, main = "")
acf_true <- ARMAacf(ar = arma21.model$coef[1:2], ma = arma11.model$coef[3], lag.max = 32)[-1]
points(1:32, acf_true, col = "red")


```




```

## extract the residuals
arma21.resids <- resid(arma21.model)
## time series plot of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.2, 1, 0), mfrow = c(1, 2))
plot(year, arma21.resids, type = "l", xlab = "year",
     ylab = "ARMA(2,1) residuals", lwd = 0.6, col = "gray")
abline(h = 0, lty = 2)
## Normal Q-Q plot for the residuals
qqnorm(arma21.resids, main = "", cex = 0.75); qqline(arma21.resids)

```



```
## Sample ACF and PACF of the residuals
par(bty = "L", mar = c(3.6, 3.6, 0.5, 0.6), las = 1, mgp = c(2.4, 1, 0), mfrow = c(1, 2))
Acf(arma21.resids, ylab = "Sample ACF", xlab = "lag (day)", main = "")
pacf(arma21.resids, ylab = "Sample PACF", xlab = "lag (day)")
```



```
## Carry out the Ljung-Box test
Box.test(arma21.resids, lag = 32, type = "Ljung-Box", fitdf = 3)
```

```
##
## Box-Ljung test
##
## data: arma21.resids
## X-squared = 32.171, df = 29, p-value = 0.3124
```

Use AIC to conduct model selection

```
AIC.to.AICC <- function(aic, n, npars) {
  aic - 2 * npars * (1 - n/(n-1-npars))
}
# calculate the length of the time series
n <- length(sqrt.rosslare.ds)

# Here are the AIC values
ar1.model$aic
```

```
## [1] 3581.432
```

```
ar2.model$aic
```

```
## [1] 3568.46
```

```
arma11.model$aic
```

```
## [1] 3565.642
```

```
arma21.model$aic
```

```
## [1] 3567.112
```

```
# convert the AIC values to AICC values.  
AIC.to.AICC(ar1.model$aic, n, 2)
```

```
## [1] 3581.438
```

```
AIC.to.AICC(ar2.model$aic, n, 3)
```

```
## [1] 3568.473
```

```
AIC.to.AICC(arma11.model$aic, n, 3)
```

```
## [1] 3565.655
```

```
AIC.to.AICC(arma21.model$aic, n, 4)
```

```
## [1] 3567.134
```

Based on the AIC (and AICC as well), we choose the ARMA(1,1) model.

Forecasting

```
## How many days will we predict into the future?  
h <- 10  
## Predict 'h' days into the future using the ARMA(1,1) model.  
sqrt.rosslare.forecast <- predict(arma11.model, h)  
sqrt.rosslare.forecast$pred; sqrt.rosslare.forecast$se
```

```
## Time Series:  
## Start = 1827  
## End = 1836  
## Frequency = 1  
## [1] 3.997161 3.458299 3.351724 3.330646 3.326477 3.325652 3.325489 3.325457  
## [9] 3.325451 3.325449
```

```
## Time Series:
## Start = 1827
## End = 1836
## Frequency = 1
## [1] 0.6409326 0.7022959 0.7045876 0.7046771 0.7046806 0.7046807 0.7046807
## [8] 0.7046807 0.7046807 0.7046807

## define the forecast variable
forecast <- sqrt.rosslare.forecast$pred
## The plus or minus value is the z critical value
## times the standard error for the forecast
me <- qnorm(0.975) * sqrt.rosslare.forecast$se
lower <- forecast - me
upper <- forecast + me
## Define the prediction time
fyear <- 1970 + (0:(h - 1)) / 365.25
```

Visualizing the Forecasts

```
par(bty = "l", mar = c(3.6, 3.6, 0.75, 0.6), las = 1, mgp = c(2.4, 1, 0),
    mfrow = c(3, 1))
## Show the data for 1969 onwards
plot(year[year > 1969], sqrt.rosslare.ds[year > 1969], type = "l",
     xlim = c(1969, max(fyear)), col = "grey", xlab = "year", ylab = "")
## Add the BLUP, along with the prediction limits
lines(fyear, forecast, lwd = 2)
lines(fyear, lower, lty = 2, lwd = 2)
lines(fyear, upper, lty = 2, lwd = 2)
## add a horizontal line at the mean
abline(h = mean(sqrt.rosslare.ds), lty = 3)
title("Forecasts for the deseasonalized square root wind speed")

## now add the seasonality estimate for the first 10 days in a year.
adj.forecast <- fitted(harm.model)[1:h] + sqrt.rosslare.forecast$pred
## adjust the lower and upper values of the interval
lower <- adj.forecast - me
upper <- adj.forecast + me

## Show the data for 1969 onwards
plot(year[year > 1969], sqrt.rosslare[year > 1969], type = "l",
     xlim = c(1969, max(fyear)), col = "grey", xlab = "year", ylab = "")
title("Forecasts for the square root wind speed")

## Add the BLUP, along with the prediction limits
lines(fyear, adj.forecast, lwd = 2)
lines(fyear, lower, lty = 2, lwd = 2)
lines(fyear, upper, lty = 2, lwd = 2)

## We square everything (forecast, lower limit, and upper limit)
## to get the forecast on the original wind speed (knots) scale.
## Show the data for 1969 onwards
plot(year[year > 1969], rosslare[year > 1969], type = "l",
```

```
xlim = c(1969, max(fyear)), col = "grey", xlab = "year", ylab = "")
title("Forecasts for the wind speed")
```

```
## Add the BLUP, along with the prediction limits
```

```
lines(fyear, adj.forecast^2, lwd = 2)
```

```
lines(fyear, lower^2, lty = 2, lwd = 2)
```

```
lines(fyear, upper^2, lty = 2, lwd = 2)
```

