

MATH 8090: State-Space Models II

Whitney Huang, Clemson University

11/17/2025

Contents

Local Level Model	2
Simulate data from local level models	2
Plot state vectors and observation vectors	2
Carrying out Kalman filter	3
Kalman filter: forecasting	5
Forecasting interval	5
Nile river flows missing values imputation	6
Kalman smoothing	9
Kalman smoothing: local level model example	10
Parameter Estimation	11
Generate data	11
Initial estimates	11
Function to evaluate the likelihood	12
Estimation	12
Global temperature example	13
EM algorithm example	15
Generate data (same as Example 6.6)	15
Initial Estimates	15
Standard Errors	16
Evaluate likelihood at estimates	16
Display summary of estimation	17
Bayesian Estimation Local Level Model	17
Generate data	17
Set up the Gibbs sampler	18
Progress bar	18
Pull out the results for easy plotting	19
Plot the results	19
References	20

Local Level Model

Simulate data from local level models

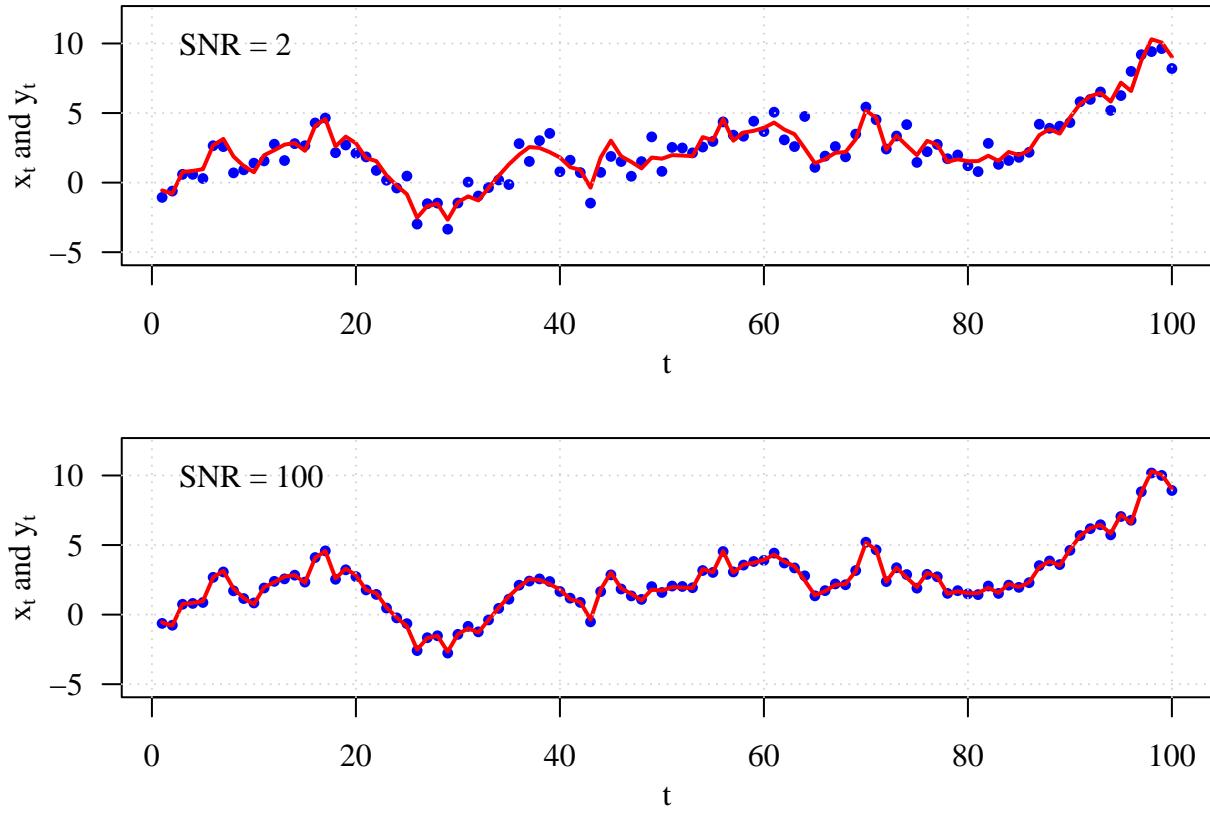
$$Y_t = X_t + W_t, \quad W_t \sim N(0, \sigma_W^2),$$
$$X_t = X_{t-1} + V_t, \quad V_t \sim N(0, \sigma_v^2).$$

Here $\mu_0 = 0$, $\sigma_0^2 = 1$, $\sigma_V^2 = 1$, $\sigma_W^2 = \sigma_V^2/\text{SNR}$.

```
set.seed(123)
mu0 <- 0; sig0 <- 1; sig2.V <- 1
X.0 <- rnorm(1, mean = mu0, sd = sqrt(sig0))
X <- cumsum(c(X.0, rnorm(99, sd = sqrt(sig2.V))))
W <- rnorm(100)
SNR <- 2; Y.2 <- X + W * sqrt(sig2.V / SNR)
SNR <- 100; Y.100 <- X + W * sqrt(sig2.V / SNR)
```

Plot state vectors and observation vectors

```
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), mfrow = c(2, 1),
     family = "serif")
plot(Y.2, col = "blue", pch = 16, cex = 0.75, xlab = "t",
      ylab = expression(paste(x[t], " and ", y[t])), main = "", ylim = c(-5.25, 12))
lines(X, col = "red", lwd = 2)
grid()
legend("topleft", legend = "SNR = 2", bty = "n")
plot(Y.100, col = "blue", pch = 16, cex = 0.75, xlab = "t",
      ylab = expression(paste(x[t], " and ", y[t])), main = "", ylim = c(-5.25, 12))
lines(X, col = "red", lwd = 2)
grid()
legend("topleft", legend = "SNR = 100", bty = "n")
```



Carrying out Kalman filter

The function below is adapted from Dr. Donald B. Percival's UW Stat 519 R codes.

1. Compute innovation:

$$U_t = Y_t - Y_t^{t-1} = Y_t - \mu_t^f.$$

2. Compute MSE for Y_t^{t-1} :

$$\Sigma_t^f + \sigma_W^2 = F_t.$$

3. Compute the new filtered value:

$$\mu_t^a = \mu_t^f + K_t U_t,$$

where $K_t = \Sigma_t^f / F_t$ is the **Kalman gain**.

4. Compute MSE for new filtered value:

$$\Sigma_t^a = \Sigma_t^f (1 - K_t).$$

5. Compute new forecast:

$$\mu_{t+1}^f = \mu_t^f + K_t U_t = \mu_t^a.$$

6. Compute MSE for new forecast:

$$\Sigma_{t+1}^f = \Sigma_t^f (1 - K_t) + \sigma_V^2 = \Sigma_t^a + \sigma_V^2.$$

```
KF.one.step.local.level <- function(X.t.tm1, P.t.tm1, Y.t, sig2.W, sig2.V){
  U.t <- if(is.na(Y.t)) NA else Y.t - X.t.tm1
  F.t <- P.t.tm1 + sig2.W
  K.t <- if(is.na(Y.t)) 0 else P.t.tm1 / F.t
  X.t.t <- X.t.tm1 + if(is.na(Y.t)) 0 else K.t * U.t
  P.t.t <- P.t.tm1 * (1 - K.t)
  X.tp1.t <- X.t.t
  P.tp1.t <- P.t.t + sig2.V
  structure(list(filter = X.t.t, forecast = X.tp1.t, filter.var = P.t.t,
                 forecast.var = P.tp1.t, innovation = U.t, innovation.var = F.t,
                 gain = K.t))
}

KF.n.steps.local.level <- function(ts, m.1 = 0, P.1 = 1, sig2.W = 1, sig2.V = 1){
  n <- length(ts)
  filter.ts <- forecast.ts <- filter.var.ts <- innovations.ts <- rep(0, n)
  forecast.var.ts <- innovations.var.ts <- gain.ts <- rep(0, n)
  X.forecast.in <- m.1; X.forecast.var.in <- P.1
  forecast.ts[1] <- X.forecast.in; forecast.var.ts[1] <- X.forecast.var.in
  Y.in <- ts[1]
  for(t in 1:n){
    temp <- KF.one.step.local.level(X.forecast.in, X.forecast.var.in, Y.in,
                                      sig2.W, sig2.V)
    filter.ts[t] <- temp$filter; filter.var.ts[t] <- temp$filter.var
    innovations.ts[t] <- temp$innovation; innovations.var.ts[t] <- temp$innovation.var
    gain.ts[t] <- temp$gain
    if(t < n){
      forecast.ts[t + 1] <- temp$forecast
      forecast.var.ts[t + 1] <- temp$forecast.var
      X.forecast.in <- temp$forecast
      X.forecast.var.in <- temp$forecast.var
      Y.in <- ts[t + 1]
    }
  }
  structure(list(filter.ts = filter.ts, forecast.ts = forecast.ts,
                 filter.var.ts = filter.var.ts, forecast.var.ts = forecast.var.ts,
                 innovations.ts = innovations.ts,
                 innovations.var.ts = innovations.var.ts,
                 gain.ts = gain.ts))
}
```

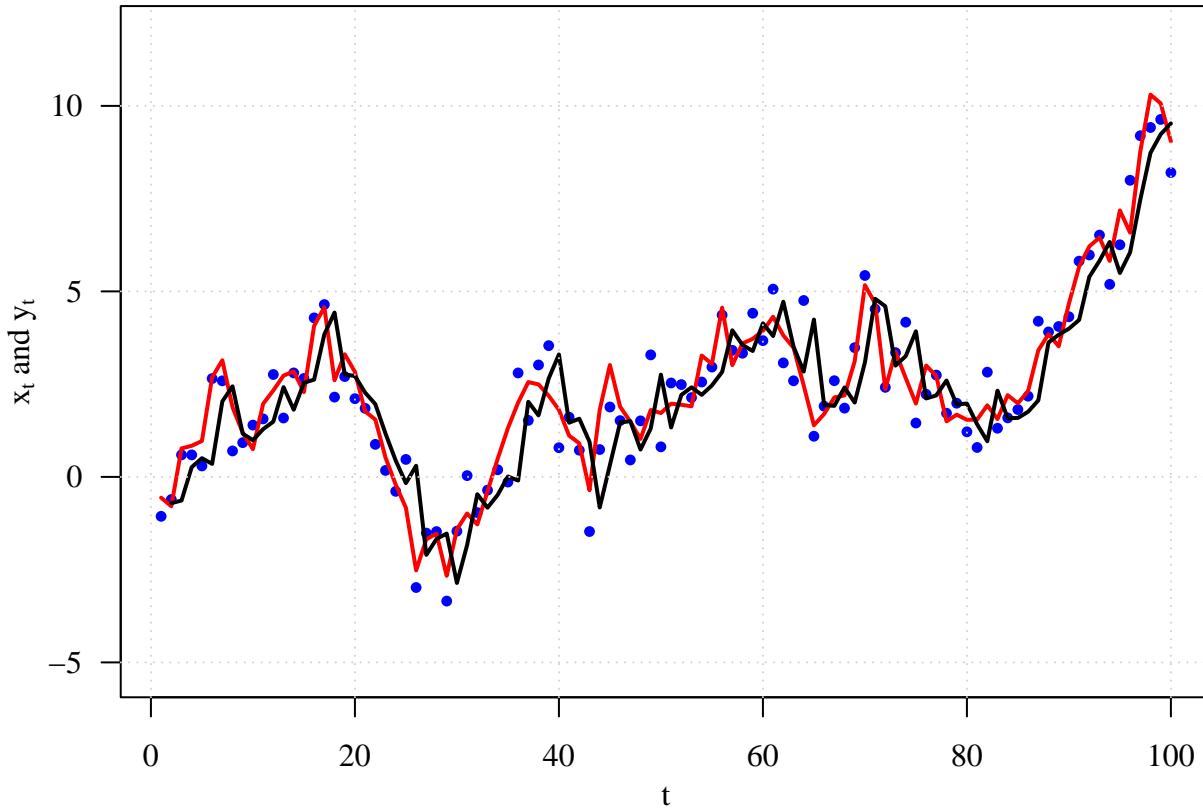
Kalman filter: forecasting

```

Y.2.KF <- KF.n.steps.local.level(Y.2, sig2.W = 0.5)

par(las = 1, mar = c(3.5, 3.5, 1, 0.6), family = "serif", mgp = c(2, 1, 0))
plot(Y.2, col = "blue", pch = 16, cex = 0.75, xlab = "t",
     ylab = expression(paste(x[t], " and ", y[t])), main = "", ylim = c(-5.25, 12))
lines(X, col = "red", lwd = 2)
lines(2:100, Y.2.KF$forecast.ts[-1], lwd = 2)
grid()

```

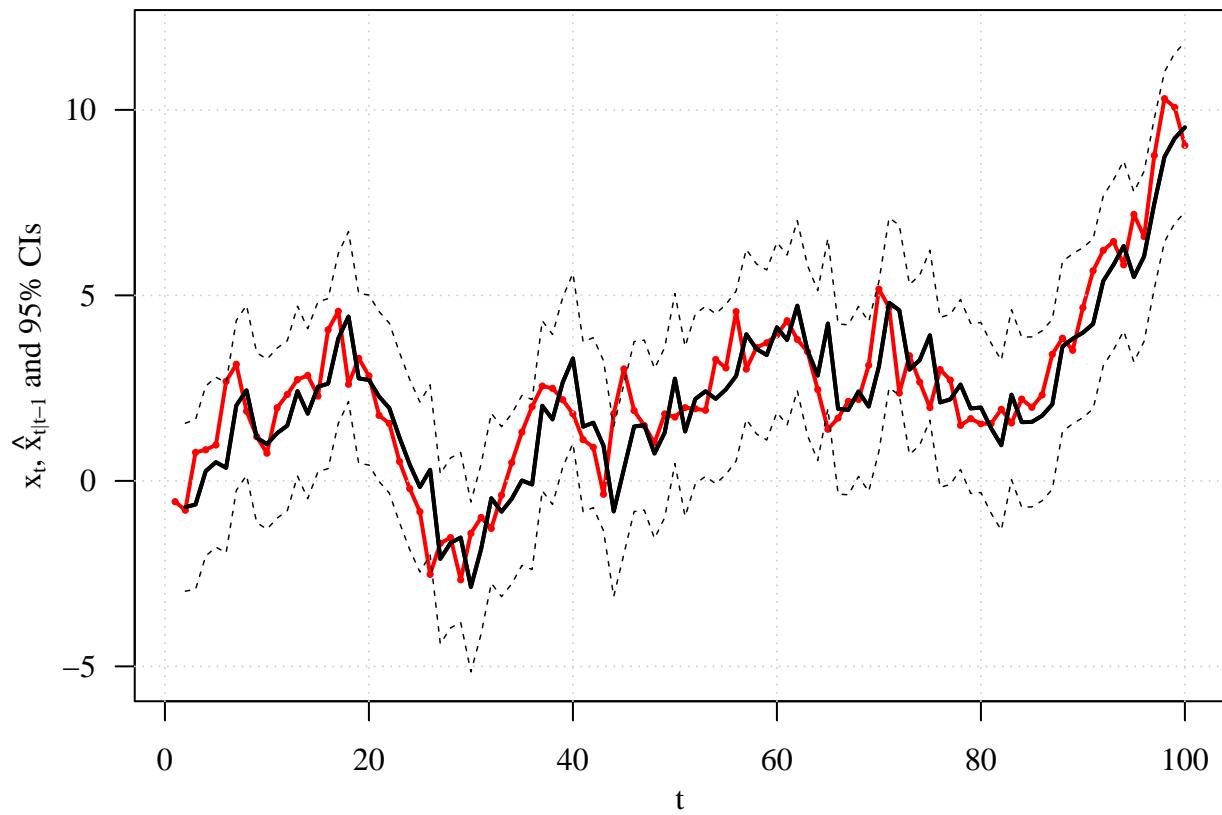


Forecasting interval

```

par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
plot(1:100, X, col = "red", pch = 16, cex = 0.5, xlab = "t",
     ylab = expression(paste(x[t], " ", hat(x)[paste(t, "|", t-1)], " and 95% CIs")),
     main = "", ylim = c(-5.25, 12))
lines(X, col = "red", lwd = 2)
lines(2:100, Y.2.KF$forecast.ts[-1], lwd = 2)
grid()
lines(2:100, Y.2.KF$forecast.ts[-1], lwd = 2)
ME <- qnorm(0.975) * sqrt(Y.2.KF$forecast.var[-1])
lines(2:100, Y.2.KF$forecast.ts[-1] - ME, lwd = .75, lty = 2)
lines(2:100, Y.2.KF$forecast.ts[-1] + ME, lwd = .75, lty = 2)

```



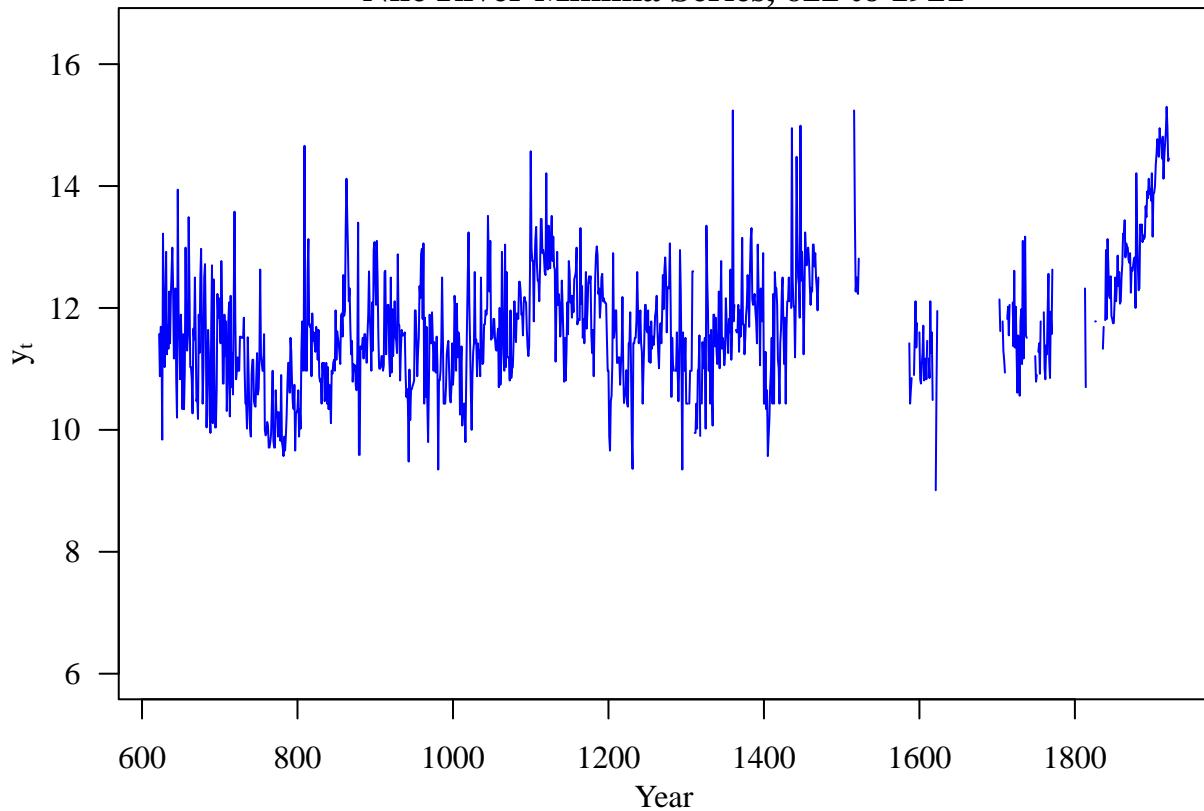
Nile river flows missing values imputation

The analysis below is adapted from Dr. Donald B. Percival's UW Stat 519 R codes.

```
nile <- scan("http://faculty.washington.edu/dbp/s519/Data/Nile-622-1921.txt")
nile.years <- 622:1921

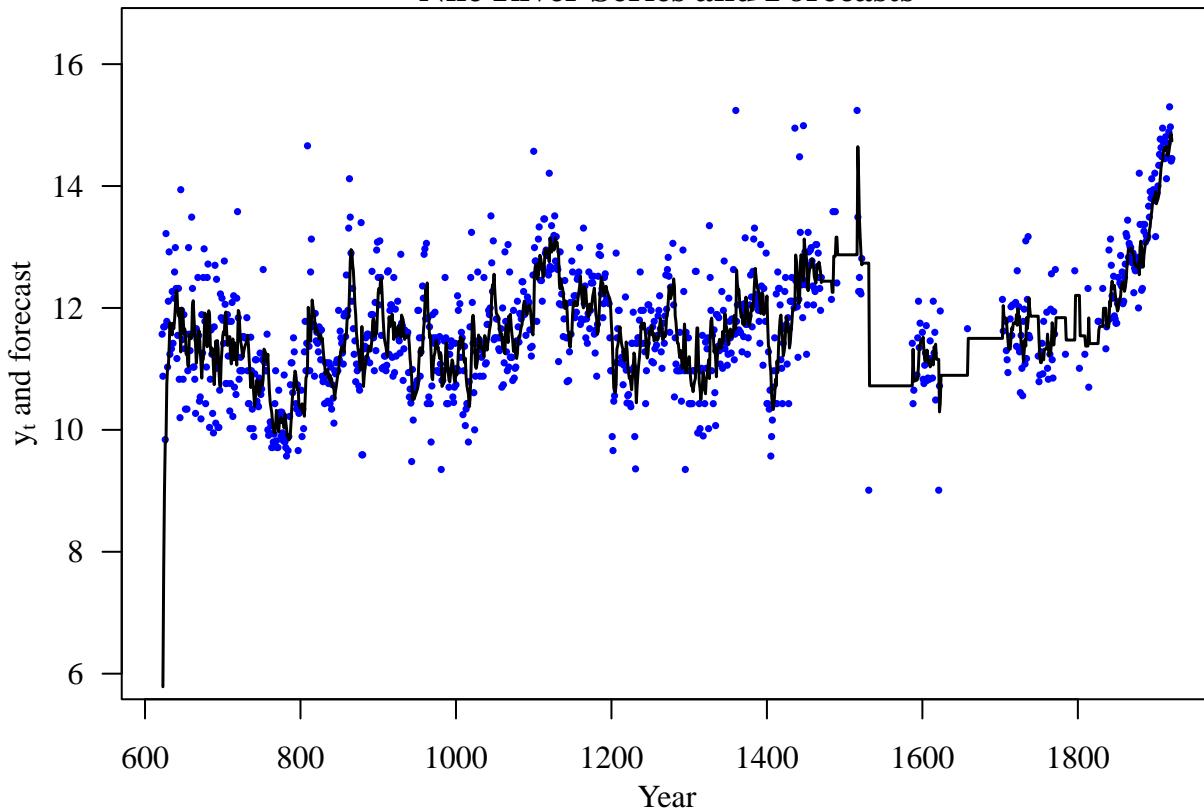
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
plot(nile.years, nile, ylim = c(6, 16.5), type = "l", col = "blue",
     xlab = "Year", ylab = expression(y[t]),
     main = "Nile River Minima Series, 622 to 1921")
```

Nile River Minima Series, 622 to 1921



```
#### Impute the missing values
nile.KF <- KF.n.steps.local.level(nile, P.1 = 1, sig2.W = 1, sig2.V = 0.1)
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
plot(nile.years, nile, pch = 16, cex = 0.5, col = "blue", ylim = c(6, 16.5),
     xlab = "Year", ylab = expression(paste(y[t], " and forecast")),
     main = "Nile River Series and Forecasts")
lines(nile.years[-1], nile.KF$forecast.ts[-1], lwd = 1.5)
```

Nile River Series and Forecasts

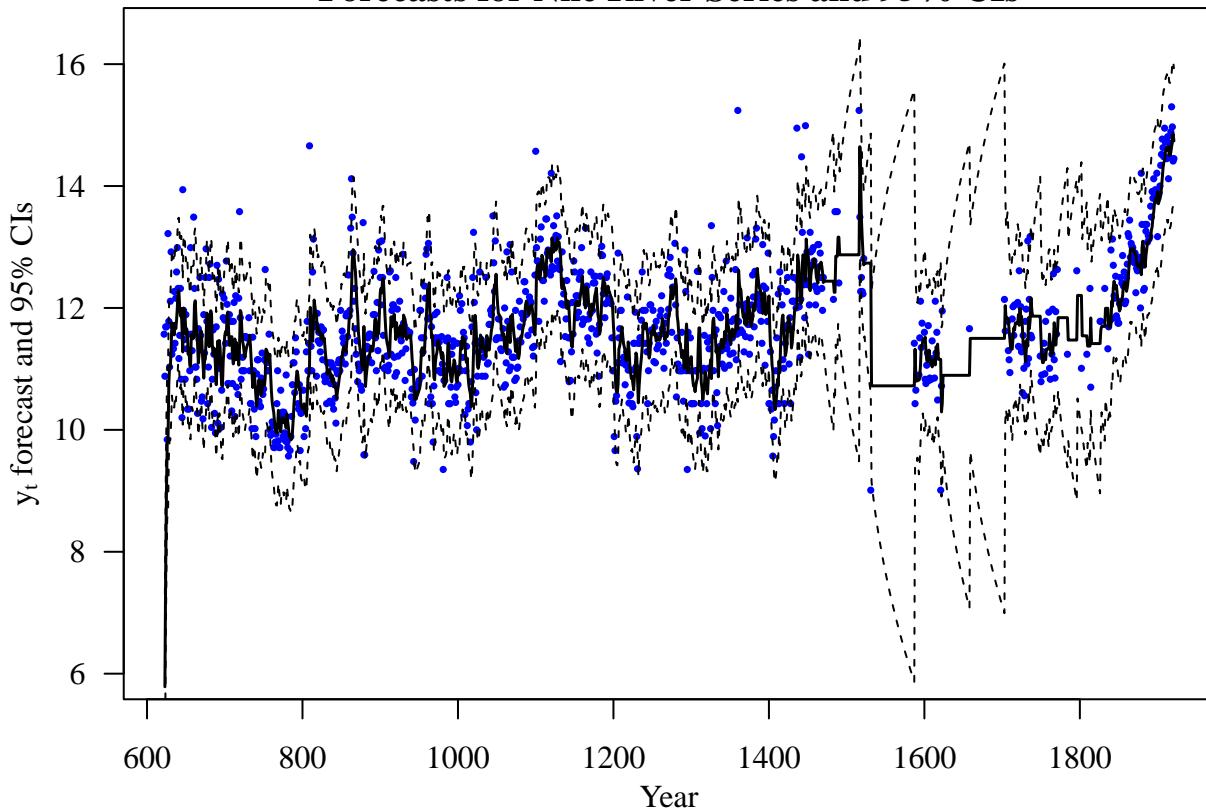


```

##### Construct confidence interval
par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
plot(nile.years, nile, pch = 16, cex = 0.5, col = "blue", ylim = c(6, 16.5),
     xlab = "Year", ylab = expression(paste(y[t], " forecast and 95% CIs")),
     main = "Forecasts for Nile River Series and 95% CIs")
ME <- qnorm(0.975) * sqrt(nile.KF$forecast.var.ts[-1])
lines(nile.years[-1], nile.KF$forecast.ts[-1], lwd = 1.5)
lines(nile.years[-1], nile.KF$forecast.ts[-1] - ME, lty = 2)
lines(nile.years[-1], nile.KF$forecast.ts[-1] + ME, lty = 2)

```

Forecasts for Nile River Series and 95% CIs



Kalman smoothing

The function below is adapted from Dr. Donald B. Percival's UW Stat 519 R codes.

```
KS.local.level <- function(KF.results){
  n <- length(KF.results$filter.ts)
  L.t.ts <- 1 - KF.results$gain.ts
  r.ts <- rep(0, n + 1)
  bg <- is.na(KF.results$innovations.ts)
  innov.0.for.NA <- KF.results$innovations.ts
  innov.0.for.NA[bg] <- 0
  for(t in n:1) r.ts[t] <- innov.0.for.NA[t] / KF.results$innovations.var.ts[t]
  + L.t.ts[t] * r.ts[t+1]
  smooth.ts <- KF.results$forecast.ts + KF.results$forecast.var.ts * r.ts[-(n + 1)]
  N.t.ts <- rep(0, n + 1)
  for(t in n:1) N.t.ts[t] <- 1 / KF.results$innovations.var.ts[t] +
    (L.t.ts[t])^2 * N.t.ts[t + 1]
  smooth.var.ts <- KF.results$forecast.var.ts -
    (KF.results$forecast.var.ts)^2 * N.t.ts[-(n + 1)]
  structure(list(L.t.ts = L.t.ts, r.ts = r.ts, smooth.ts = smooth.ts,
                 N.t.ts = N.t.ts, smooth.var.ts = smooth.var.ts))
}
```

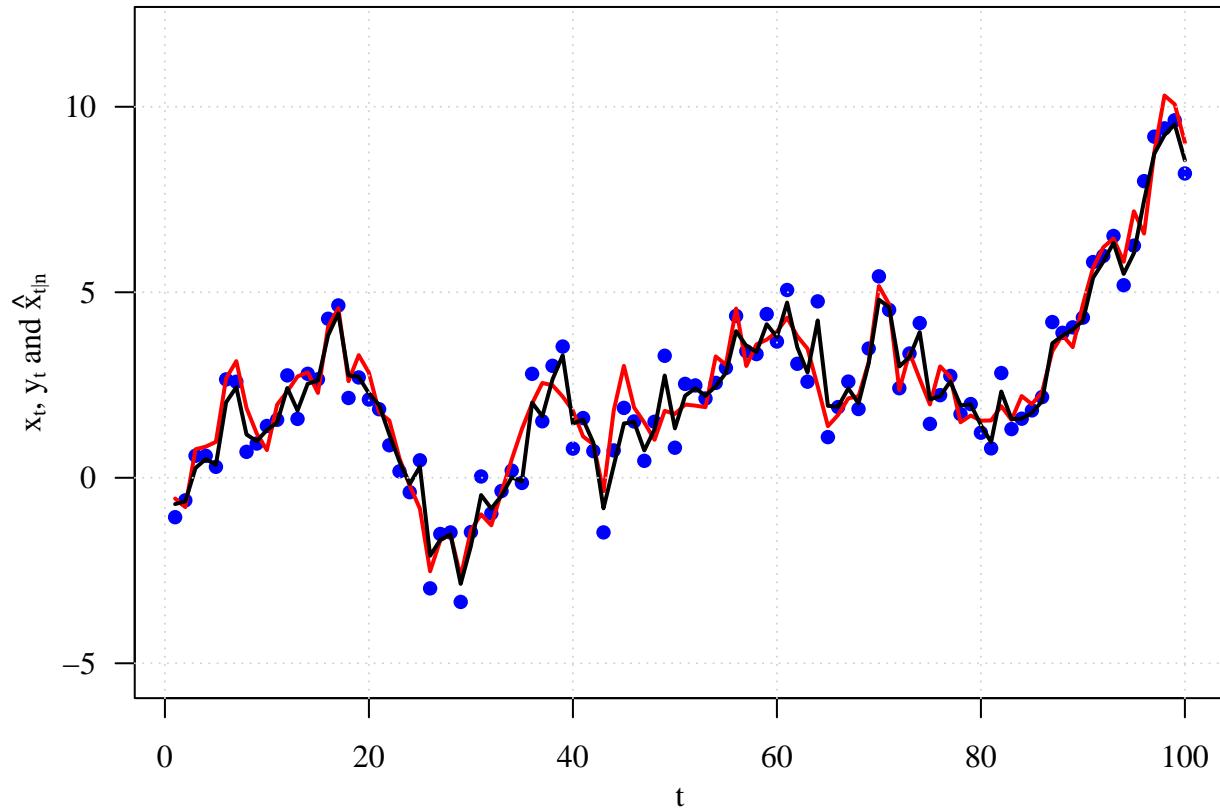
Kalman smoothing: local level model example

```

Y.2.KS <- KS.local.level(Y.2.KF)

par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
plot(Y.2, col = "blue", pch = 16, xlab = "t",
     ylab = expression(paste(x[t], ", ", y[t], " and ", hat(x)[paste(t, "|", n)])), main = "",
     ylim = c(-5.25, 12))
lines(X, col = "red", lwd = 2)
lines(Y.2.KS$smooth.ts, lwd = 2)
grid()

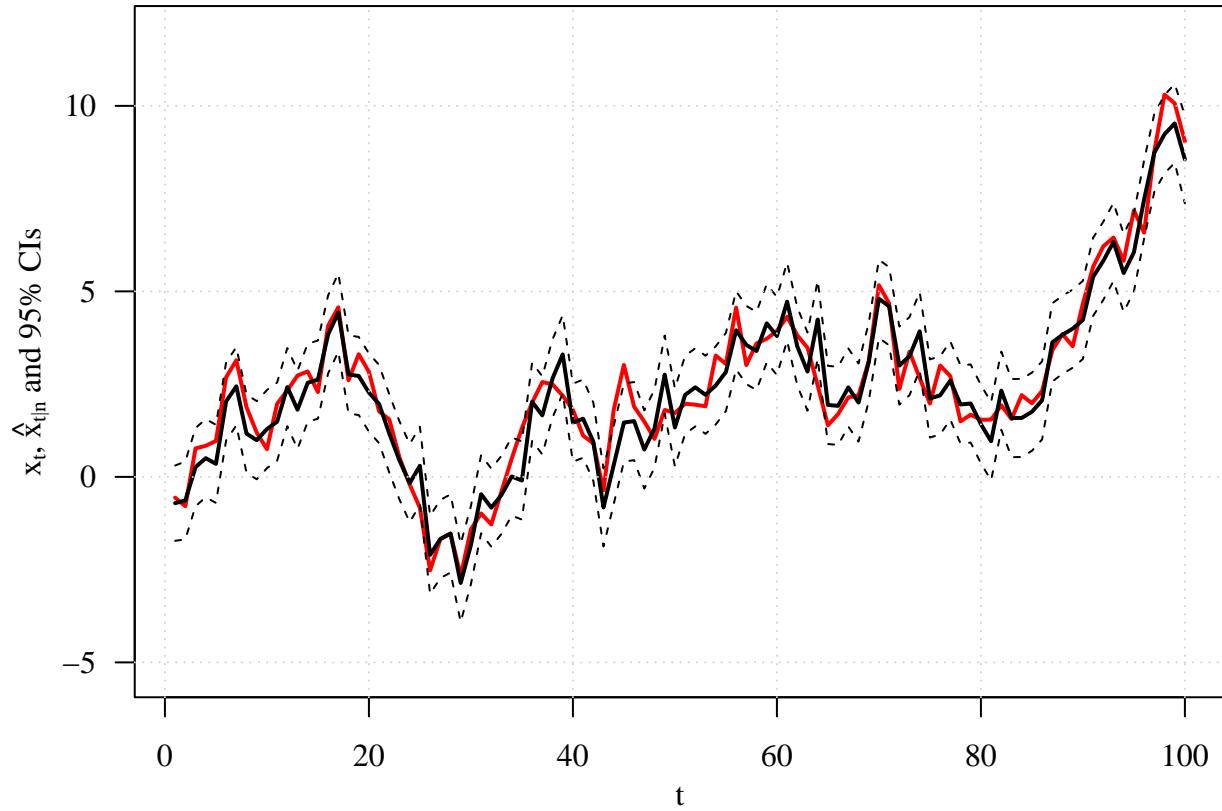
```



```

par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6))
plot(X, col = "red", typ = "l", lwd = 2, xlab = "t",
      ylab = expression(paste(x[t], ", ", hat(x)[paste(t, "|", n)], " and 95% CIs")),
      main = "", ylim = c(-5.25, 12))
lines(Y.2.KS$smooth.ts, lwd = 2)
lines(Y.2.KS$smooth.ts - 1.96 * sqrt(Y.2.KS$smooth.var.ts), lty = 2)
lines(Y.2.KS$smooth.ts + 1.96 * sqrt(Y.2.KS$smooth.var.ts), lty = 2)
grid()

```



Parameter Estimation

This example is taken from Shumway and Stoffer (2017) example 6.6

Generate data

```
library(astsa)
set.seed(123)
num = 100
N = num + 1
x <- sarima.sim(n = N, ar = .8)
y <- ts(x[-1] + rnorm(num, 0, 1))
```

Initial estimates

$$\begin{aligned}\phi^{(0)} &= \frac{\hat{\rho}_Y(2)}{\hat{\rho}_Y(1)}. \\ \sigma_W^{2(0)} &= (1 - \phi^{2(0)})\hat{\gamma}_Y(1)/\phi^{(0)}. \\ \sigma_V^{2(0)} &= \hat{\gamma}_Y(0) - \left[\frac{\sigma_W^{2(0)}}{1 - \phi^{2(0)}} \right]\end{aligned}$$

```
u = ts.intersect(y, lag(y, -1), lag(y, -2))
varu = var(u)
coru = cor(u)
```

```

phi = coru[1, 3] / coru[1, 2]
q = (1 - phi^2) * varu[1, 2] / phi
r = varu[1, 1] - q / (1 - phi^2)
(init.par = c(phi, sqrt(q), sqrt(r)))

```

```
## [1] 0.7614651 1.0020091 0.8744762
```

Function to evaluate the likelihood

```

Linn <- function(para){
  phi <- para[1]; sigw <- para[2]; sigv <- para[3]
  Sigma0 <- (sigw^2) / (1 - phi^2); Sigma0[Sigma0 < 0] = 0
  kf = Kfilter(y, 1, mu0 = 0, Sigma0, phi, sigw, sigv)
  return(kf$like)
}

```

Estimation

```
(est = optim(init.par, Linn, gr = NULL, method = "BFGS", hessian = TRUE,
            control = list(trace = 1, REPORT = 1)))
```

```

## initial value 84.170842
## iter 2 value 84.102702
## iter 3 value 83.916203
## iter 4 value 83.915653
## iter 5 value 83.889723
## iter 6 value 83.885783
## iter 7 value 83.885762
## iter 7 value 83.885762
## iter 7 value 83.885762
## final value 83.885762
## converged

## $par
## [1] 0.8213276 0.8308274 0.9691287
##
## $value
## [1] 83.88576
##
## $counts
## function gradient
##      29          7
##
## $convergence
## [1] 0
##
## $message
## NULL
##
```

```

## $hessian
##      [,1]      [,2]      [,3]
## [1,] 263.738652 74.14214 -9.936399
## [2,] 74.142142 69.77014 44.355806
## [3,] -9.936399 44.35581 85.616367

SE = sqrt(diag(solve(est$hessian)))
cbind(estimate = c(phi = est$par[1], sigw = est$par[2], sigv = est$par[3]), SE)

```

```

##           estimate        SE
## phi    0.8213276 0.08831157
## sigw  0.8308274 0.20920610
## sigv  0.9691287 0.15849779

```

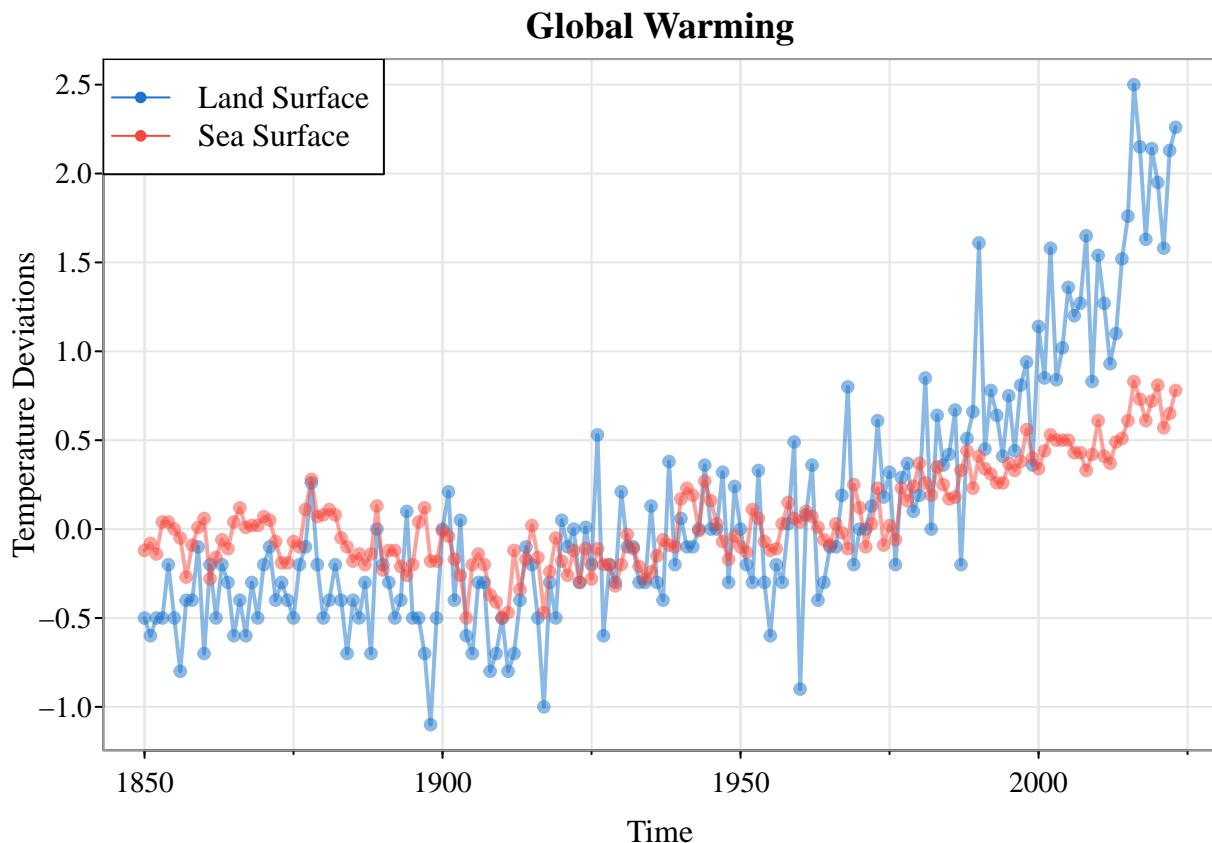
Global temperature example

This example is taken from Shumway and Stoffer (2017) example 6.7

```

par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
tsplot(cbind(gtemp_land, gtemp_ocean), spaghetti = TRUE,
       lwd = 2, pch = 20, type = "o", col=astsa.col(c(4,2),.5),
       ylab = "Temperature Deviations", main = "Global Warming")
legend("topleft", legend = c("Land Surface", "Sea Surface"), lty = 1,
       pch = 20, col = c(4, 2), bg = "white")

```



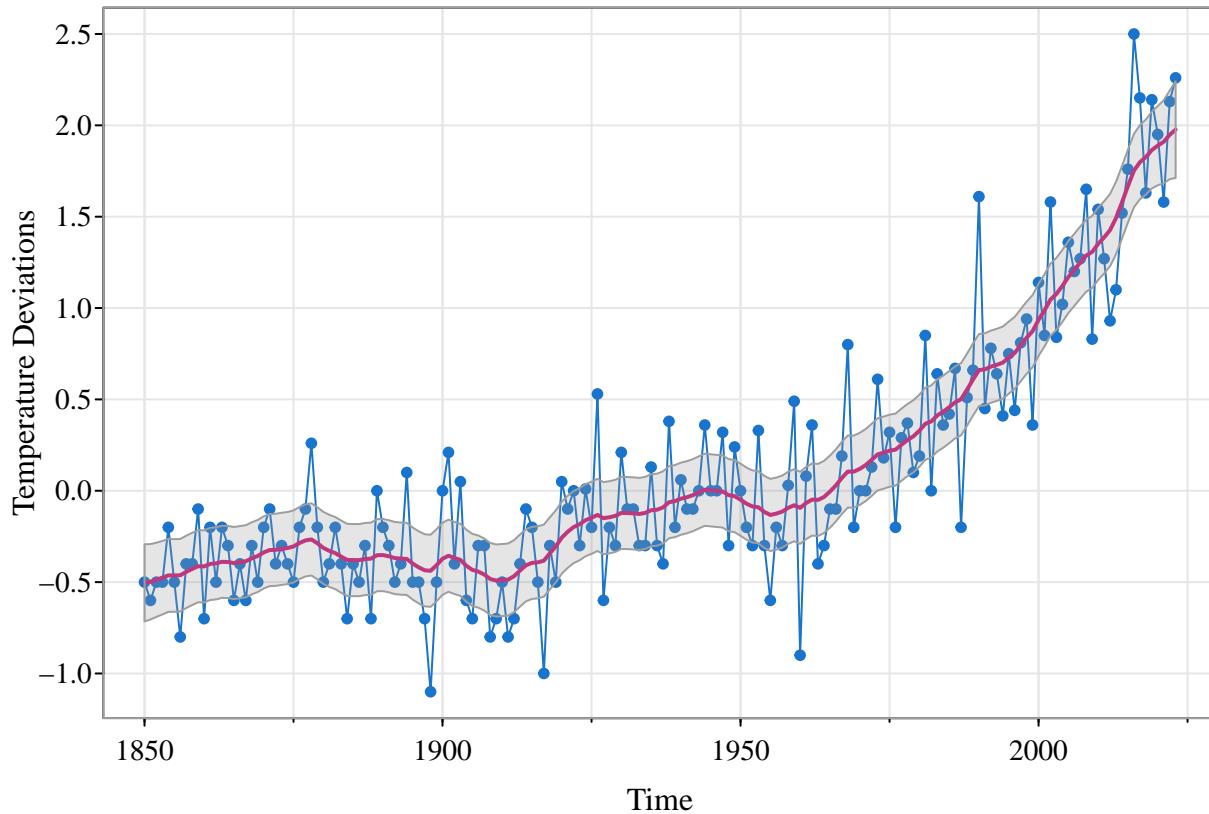
```

u = ssm(gtemp_land, A = 1, phi = 1, alpha = .01, sigw = .01,
         sigv = .1, fixphi = TRUE)

## initial value 548.927004
## iter 2 value 426.219829
## iter 3 value 425.534188
## iter 4 value 421.056500
## iter 5 value 420.929087
## iter 6 value 420.769864
## iter 7 value 419.824779
## iter 8 value 417.481820
## iter 9 value 308.066621
## iter 10 value 227.358420
## iter 11 value 91.340031
## iter 12 value -7.648163
## iter 13 value -52.574167
## iter 14 value -56.368335
## iter 15 value -59.350322
## iter 16 value -67.455015
## iter 17 value -70.682140
## iter 18 value -70.772274
## iter 19 value -79.846982
## iter 20 value -99.576830
## iter 21 value -100.354988
## iter 22 value -101.698806
## iter 23 value -104.630208
## iter 24 value -104.863636
## iter 25 value -105.139492
## iter 26 value -105.983246
## iter 27 value -106.008918
## iter 28 value -106.010596
## iter 29 value -106.010668
## iter 29 value -106.010668
## final value -106.010668
## converged
##           estimate      SE
## alpha 0.01429012 0.005140273
## sigw  0.06643629 0.013370144
## sigv  0.29494320 0.017369955

par(las = 1, mgp = c(2, 1, 0), mar = c(3.5, 3.5, 1, 0.6), family = "serif")
tsplot(gtemp_land, col = 4, type = "o", pch = 20, ylab = "Temperature Deviations")
lines(u$Xs, col = 6, lwd = 2)
xx = c(time(u$Xs), rev(time(u$Xs)))
yy = c(u$Xs - 2*sqrt(u$Ps), rev(u$Xs + 2*sqrt(u$Ps)))
polygon(xx, yy, border = 8, col = gray(.6, alpha = .25))

```



EM algorithm example

Generate data (same as Example 6.6)

```
library(nlme)
set.seed(123); num = 100; N = num + 1
x = sarima.sim(ar = .8, n = N)
y = ts(x[-1] + rnorm(num, 0, 1))
```

Initial Estimates

```
u = ts.intersect(y,lag(y,-1), lag(y,-2))
varu = var(u); coru = cor(u)
phi = coru[1,3] / coru[1,2]
q = (1 - phi^2) * varu[1, 2] / phi
r = varu[1, 1] - q/(1 - phi^2)
cr = sqrt(r); cq = sqrt(q); mu0 = 0; Sigma0 = 2.8
(em = EM(y, 1, mu0, Sigma0, phi, cq, cr, 75, .00001))
```

```
## iteration      -loglikelihood
##     1          84.36778
##     2          83.97942
##     3          83.82139
```

```

##      4      83.74255
##      5      83.69475
##      6      83.66085
##      7      83.63427
##      8      83.61222
##      9      83.59335
##     10      83.57691
##     11      83.56242
##     12      83.54955
##     13      83.53808
##     14      83.52781
##     15      83.51859
##     16      83.5103

## $Phi
## [1] 0.8106963
##
## $Q
## [1] 0.7752158
##
## $R
##      [,1]
## [1,] 0.8704274
##
## $mu0
##      [,1]
## [1,] 0.7842457
##
## $Sigma0
##      [,1]
## [1,] 0.1469216
##
## $like
## [1] 84.36778 83.97942 83.82139 83.74255 83.69475 83.66085 83.63427 83.61222
## [9] 83.59335 83.57691 83.56242 83.54955 83.53808 83.52781 83.51859 83.51030
##
## $niter
## [1] 16
##
## $cvg
## [1] 9.921766e-05

```

Standard Errors

```

phi = em$Phi; cq = chol(em$Q); cr = chol(em$R)
mu0 = em$mu0; Sigma0 = em$Sigma0
para = c(phi, cq, cr)

```

Evaluate likelihood at estimates

```

Linn = function(para){
  kf = Kfilter(y, 1, mu0, Sigma0, para[1], para[2], para[3])
  return(kf$like)
}
emhess = fdHess(para, function(para) Linn(para))
SE = sqrt(diag(solve(emhess$Hessian)))

```

Display summary of estimation

```

estimate = c(para, em$mu0, em$Sigma0); SE = c(SE, NA, NA)
u = cbind(estimate, SE)
rownames(u) = c("phi", "sigw", "sigv", "mu0", "Sigma0")
u

##           estimate          SE
## phi      0.8106963 0.09836856
## sigw    0.8804634 0.23235380
## sigv    0.9329670 0.17421057
## mu0     0.7842457      NA
## Sigma0   0.1469216      NA

```

Bayesian Estimation Local Level Model

This example is taken from Shumway and Stoffer (2017) example 6.7

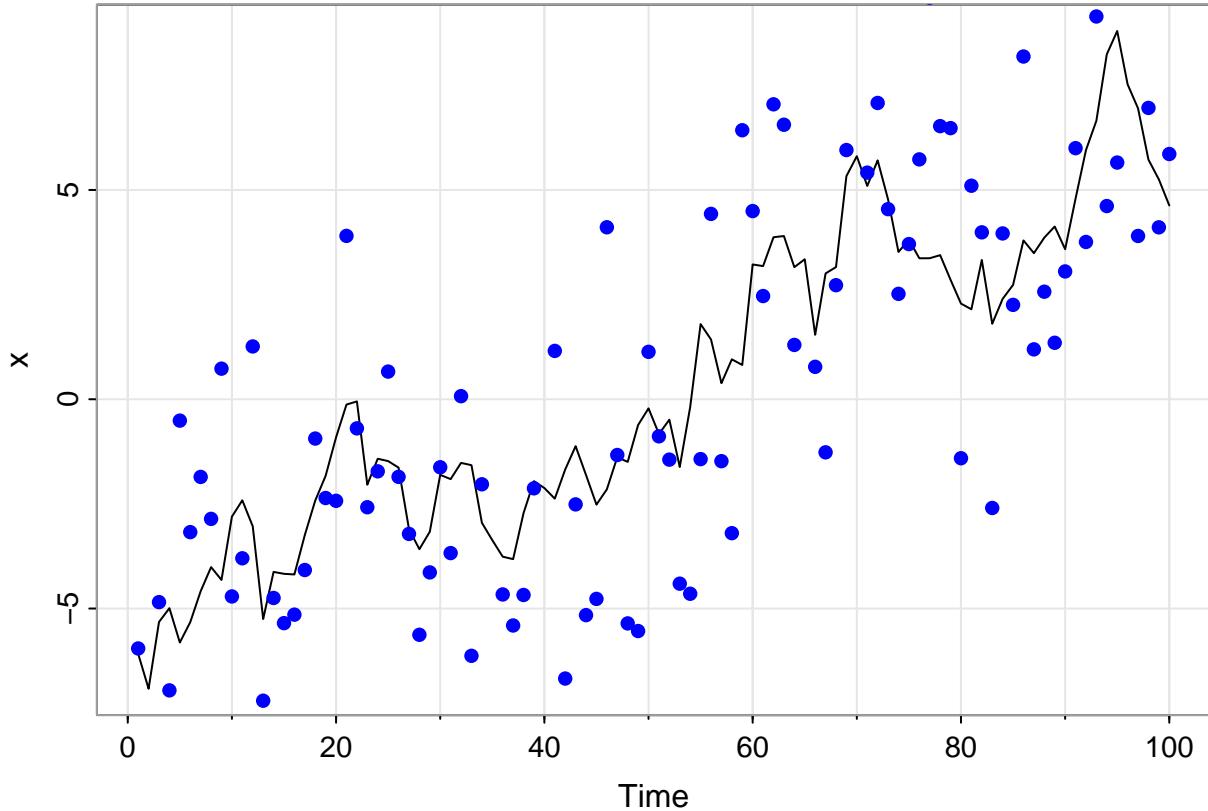
Generate data

```

set.seed(1)
sQ = 1; sR = 3; n = 100
mu0 = 0; Sigma0 = 10; x0 = rnorm(1, mu0, Sigma0)
w = rnorm(n); v = rnorm(n)
x = c(x0 + sQ * w[1]) # initialize states
y = c(x[1] + sR * v[1]) # initialize obs
for (t in 2:n){
  x[t] = x[t - 1] + sQ*w[t]
  y[t] = x[t] + sR * v[t]
}

tsplot(x, pch = 16)
points(1:100, y, pch = 16, col = "blue")

```



Set up the Gibbs sampler

```

burn = 50; n.iter = 1000
niter = burn + n.iter
draws = c()
# priors for R (a,b) and Q (c,d) IG distributions
a = 2; b = 2; c = 2; d = 1
# (1) initialize - sample sQ and sR
sR = sqrt(1/rgamma(1, a, b)); sQ = sqrt(1/rgamma(1, c, d))

```

Progress bar

```

pb = txtProgressBar(min = 0, max = niter, initial = 0, style = 3)

## | 

# run it
for (iter in 1:niter){
  ## (2) sample the states
  run = ffbs(y, 1, 0, 10, 1, sQ, sR) # ffbs(y,A,mu0,Sigma0,Phi,Ups,Gam,sQ,sR,input)
  ## (1) sample the parameters
  Xs = as.matrix(run$Xs)
  R = 1/rgamma(1, a + n/2, b + sum((y - Xs)^2)/2)
}

```

```

sR = sqrt(R)
Q = 1/rgamma(1,c + (n - 1)/2, d + sum(diff(Xs)^2)/2)
sQ = sqrt(Q)
## store everything
draws = rbind(draws, c(sQ, sR, Xs))
setTxtProgressBar(pb, iter)
}

##    |

close(pb)

```

Pull out the results for easy plotting

```

draws = draws[(burn + 1):(niter),]
q025 = function(x){quantile(x, 0.025)}
q975 = function(x){quantile(x, 0.975)}
xs = draws[, 3:(n + 2)]
lx = apply(xs, 2, q025)
mx = apply(xs, 2, mean)
ux = apply(xs, 2, q975)

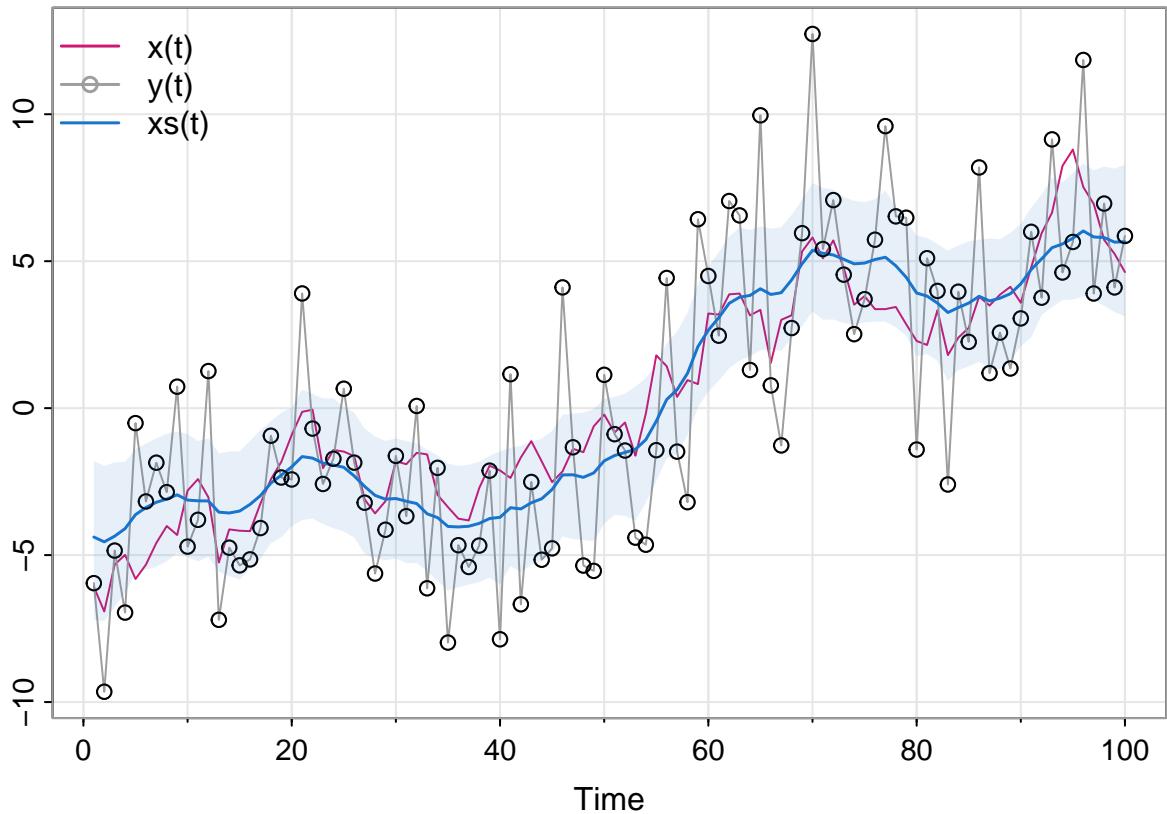
```

Plot the results

```

tsplot(cbind(x, y, mx), spag = TRUE, ylab = '', col = c(6, 8, 4),
       lwd = c(1, 1, 1.5), type = 'o', pch = c(NA, 1, NA))
legend('topleft', legend = c("x(t)", "y(t)", "xs(t)"),
       lty = 1, col = c(6, 8, 4), lwd = 1.5, bty = "n", pch = c(NA, 1, NA))
points(y)
xx = c(1:100, 100:1)
yy = c(lx, rev(ux))
polygon(xx, yy, border = NA, col = astsa.col(4, .1))

```



References

Shumway, Robert H, and David S Stoffer. 2017. *Time Series Analysis and Its Applications*. 4th ed. Springer.