

Lecture 21

Inference for Bivariate Categorical Data

Text: Chapter 10

STAT 8010 Statistical Methods I
April 7, 2020

Whitney Huang
Clemson University



Notes

Binomial Experiments and Inference for Binary Category Data

- Fixed number of n trials (sample size), each trial is an independent event (simple random sample)
- Binary outcomes ("success/failure"), where the probability of success, p , for each trial is constant
- The number of successes $X \sim \text{Bin}(n, p)$

We use a random sample from X to infer p , the population proportion



Notes

Multinomial Experiments and Inference for Multi-Category Data

- Fixed number of n trials, each trial is an independent event
- K possible outcomes, each with probability $p_k, k = 1, \dots, K$ where $\sum_{k=1}^K p_k = 1$
- $(X_1, X_2, \dots, X_K) \sim \text{Multi}(n, p_1, p_2, \dots, p_K)$

We use a random sample from (X_1, X_2, \dots, X_K) to infer $\{p_k\}_{k=1}^K$, the event probabilities

Question: How many parameters here? $\Rightarrow K - 1$




Notes

Example: Multinomial Probability

Suppose that in a three-way election for a large country, candidate 1 received 20% of the votes, candidate 2 received 35% of the votes, and candidate 3 received 45% of the votes. If ten voters are **selected randomly**, what is the probability that there will be exactly two supporter for candidate 1, three supporters for candidate 2 and five supporters for candidate 3 in the sample?

$$\mathbb{P}(X_1 = 2, X_2 = 3, X_3 = 5) = \frac{10!}{2!3!5!} (0.2)^2 (0.35)^3 (0.45)^5 \approx 0.08$$




Notes

Example: Estimating Multinomial Parameters

If we **randomly select** ten voters, two supporter for candidate 1, three supporters for candidate 2 and five supporters for candidate 3 in the sample. What would our best guess for the population proportion each candidate would received?

$$\begin{aligned}\hat{p}_1 &= \frac{x_1}{n} = \frac{2}{10} = .2 \\ \hat{p}_2 &= \frac{x_2}{n} = \frac{3}{10} = .3 \\ \hat{p}_3 &= 1 - \hat{p}_1 - \hat{p}_2 = 1 - .2 - .3 = .5\end{aligned}$$



Notes

Pearson's χ^2 Test (Multinomial Test)

Goal: To test if the parameters of a multinomial distribution equal specified values

- The Hypotheses:

$$H_0 : p_1 = p_{1,0}; p_2 = p_{2,0}; \dots, p_K = p_{K,0}$$


H_a : At least one is different

- The Test Statistic:

$$\chi^2_* = \sum_{k=1}^K \frac{(O_k - E_k)^2}{E_k},$$

where O_k is the observed frequency for the k_{th} event and E_k is the expected frequency under H_0


- The Null Distribution: $\chi^2_* \sim \chi^2_{df=K-1}$
- Assumption: $np_k > 5, k = 1, \dots, K$



Notes

χ^2 -Distribution

chi-square



Notes

Example: Testing Mendel's Theories (pp 22–23, “Categorical Data Analysis” 2nd Ed by Alan Agresti)

“Among its many applications, Pearson’s test was used in genetics to test Mendel’s theories of natural inheritance. Mendel crossed pea plants of pure yellow strain (dominant strain) plants of pure green strain. He predicted that second generation hybrid seeds would be 75% yellow and 25% green. One experiment produced $n = 8023$ seeds, of which $X_1 = 6022$ were yellow and $X_2 = 2001$ were green.”

Use Pearson’s χ^2 test to assess Mendel’s hypothesis.

Notes

Testing Mendel's Theory

- $H_0 : p_1 = .75; p_2 = .25$ vs. $H_a : \text{at least one } p_i \neq p_{i,0}, i = 1, 2$
- $\chi_*^2 = \sum_{k=1}^K \frac{(O_k - E_k)^2}{E_k}$, where
 $K = 2, O_1 = 6022, O_2 = 2001, E_1 = 8023 \times .75 = 6017.25, E_2 = 8023 \times .25 = 2005.75$. We have
 $\chi_{obs}^2 = \frac{(6022 - 6017.25)^2}{6017.25} + \frac{(2001 - 2005.75)^2}{2005.75} = 0.0150$.
- Since $\chi_{obs}^2 = 0.0150 < \chi_{0.05,1}^2 = 3.8414$, we fail to reject the H_0 .
- We do not have enough evidence to reject Mendel's Theory.

Notes

Color Preference Example

In Child Psychology, color preference by young children is used as an indicator of emotional state. In a study of 112 children, each was asked to choose “favorite” color from the 7 colors indicated below. Test if there is evidence of a preference at the 5% level.

Color	Blue	Red	Green	White	Purple	Black	Other
Frequency	13	14	8	17	25	15	20

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.10

Notes

Color Preference Example Cont'd

- $H_0 : p_1 = p_2 = \cdots = p_7 = \frac{1}{7}$ vs. H_a : At least one $p_k \neq \frac{1}{7}, k = 1, \cdots, 7$.
- $E_k = 112 \times \frac{1}{7} = 16, k = 1, \cdots, 7$.
 $\chi^2_{obs} = \frac{(13-16)^2 + (14-16)^2 + \cdots + (20-16)^2}{16} = 11$.
- Since $\chi^2_{obs} = 11 < 12.59 = \chi^2_{0.05,6}$, we fail to reject H_0 at 0.05 level.
- We do not have enough evidence that there exists color preference at 0.05 level.

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.11

Notes

An Example of Bivariate Categorical Data

A psychologist is interested in whether or not handedness is related to gender. She collected data on handedness for 100 individuals and the data set is summarized in the table below

	Right-handed	Left-handed	Total
Males	43	9	52
Females	44	4	48
Total	87	13	100

- Grand total: 100
- Marginal total for males: 52
- Marginal total for females: 48
- Marginal total for right-handed: 87
- Marginal total for left-handed: 13

This is an example of a contingency table

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.12

Notes

Contingency Tables

- Bivariate categorical data is typically displayed in a contingency table
- The number in each cell is the frequency for each category level combination
- Contingency table for the previous example:

	Right-handed	Left-handed	Total
Males	43	9	52
Females	44	4	48
Total	87	13	100

For a given contingency table, we want to test **if two variables have a relationship or not?** $\Rightarrow \chi^2$ -Test

Notes

χ^2 -Test for Independence

- Define the null and alternative hypotheses:

H_0 : there is no relationship between the 2 variables

H_a : there is a relationship between the 2 variables

- (If necessary) Calculate the marginal totals, and the grand total

- Calculate the expected cell frequencies:

$$\text{Expected cell frequency} = \frac{\text{Row Total} \times \text{Column Total}}{\text{Grand Total}}$$

- Calculate the partial χ^2 values (χ^2 value for each cell of the table):

$$\text{Partial } \chi^2 \text{ value} = \frac{(\text{observed} - \text{expected})^2}{\text{expected}}$$

Notes

χ^2 -Test for Independence Cont'd

- Calculate the χ^2 statistic:

$$\chi^2_{obs} = \sum \text{partial } \chi^2 \text{ value}$$

- Calculate the degrees of freedom (df)

$$df = (\text{\#of rows} - 1) \times (\text{\#of columns} - 1)$$

- Find the χ^2 critical value with respect to α

- Draw the conclusion:

Reject H_0 if χ^2_{obs} is bigger than the χ^2 critical value \Rightarrow
There is an statistical evidence that there is a relationship between the two variables at α level

Notes

Handedness/Gender Example Revisited

	Right-handed	Left-handed	Total
Males	43	9	52
Females	44	4	48
Total	87	13	100

Is the percentage left-handed men in the population different from the percentage of left-handed women?

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.16

Notes

Example

A 2011 study was conducted in Kalamazoo, Michigan. The objective was to determine if parents' marital status affects children's marital status later in their life. In total, 2,000 children were interviewed. The columns refer to the parents' marital status. Use the contingency table below to conduct a χ^2 test from beginning to end. Use $\alpha = .10$

(Observed)	Married	Divorced	Total
Married	581	487	
Divorced	455	477	
Total			

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.17

Notes

Example Cont'd

- 1 Define the Null and Alternative hypotheses:
- H_0 : there is no relationship between parents' marital status and childrens' marital status.
- H_a : there is a relationship between parents' marital status and childrens' marital status
- 2 Calculate the marginal totals, and the grand total

(Observed)	Married	Divorced	Total
Married	581	487	1068
Divorced	455	477	932
Total	1036	964	2000

Inference for Bivariate Categorical Data

CLEMSON UNIVERSITY

21.18

Notes

Example Cont'd

- Calculate the expected cell counts

(Expected)	Married	Divorced
Married	$\frac{1068 \times 1036}{2000} = 553.224$	$\frac{1068 \times 964}{2000} = 514.776$
Divorced	$\frac{932 \times 1036}{2000} = 482.776$	$\frac{932 \times 964}{2000} = 449.224$

- Calculate the partial χ^2 values

partial χ^2	Married	Divorced
Married	$\frac{(581 - 553.224)^2}{553.224} = 1.39$	$\frac{(487 - 514.776)^2}{514.776} = 1.50$
Divorced	$\frac{(455 - 482.776)^2}{482.776} = 1.60$	$\frac{(477 - 449.224)^2}{449.224} = 1.72$

- Calculate the χ^2 statistic

$$\chi^2 = 1.39 + 1.50 + 1.60 + 1.72 = 6.21$$

- Calculate the degrees of freedom (df)

$$\text{The } df \text{ is } (2 - 1) \times (2 - 1) = 1$$

- Find the χ^2 critical value with respect to α from the χ^2 table

$$\text{The } \chi^2_{\alpha=0.1, df=1} = 2.71$$

- Draw your conclusion:

We reject H_0 and conclude that there is a relationship between parents' marital status and childrens' marital status.

Notes

Example

The following contingency table contains enrollment data for a random sample of students from several colleges at Purdue University during the 2006-2007 academic year. The table lists the number of male and female students enrolled in each college. Use the two-way table to conduct a χ^2 test from beginning to end. Use $\alpha = .01$

(Observed)	Female	Male	Total
Liberal Arts	378	262	640
Science	99	175	274
Engineering	104	510	614
Total	581	947	1528

Notes

Example Cont'd

(Expected)	Female	Male
Liberal Arts	$\frac{640 \times 581}{1528} = 243.35$	$\frac{640 \times 947}{1528} = 396.65$
Science	$\frac{274 \times 581}{1528} = 104.18$	$\frac{274 \times 947}{1528} = 169.82$
Engineering	$\frac{614 \times 581}{1528} = 233.46$	$\frac{614 \times 947}{1528} = 380.54$

partial χ^2	Female	Male
Lib Arts	$\frac{(378 - 243.35)^2}{243.35} = 74.50$	$\frac{(262 - 396.65)^2}{396.65} = 45.71$
Sci	$\frac{(99 - 104.18)^2}{104.18} = 0.26$	$\frac{(175 - 169.82)^2}{169.82} = 0.16$
Eng	$\frac{(104 - 233.46)^2}{233.46} = 71.79$	$\frac{(510 - 380.54)^2}{380.54} = 44.05$

$$\chi^2 = 74.50 + 45.71 + 0.26 + 0.16 + 71.79 + 44.05 = 236.47$$

$$\text{The } df = (3 - 1) \times (2 - 1) = 2 \Rightarrow \text{Critical value}$$

$$\chi^2_{\alpha=0.01, df=2} = 9.21$$

Therefore we reject H_0 (at .01 level) and conclude that there is a relationship between gender and major.

Notes

R Code & Output

```
table <- matrix(c(378, 99, 104,
                  262, 175, 510), 3, 2)
colnames(table) <- c("Female", "Male")
rownames(table) <- c("Liberal Arts", "Science",
"Engineering")
table
```

	Female	Male
Liberal Arts	378	262
Science	99	175
Engineering	104	510

```
chisq.test(table)
```

Pearson's Chi-squared test

data: table
X-squared = 236.47, df = 2, p-value < 2.2e-16



Notes

Take Another Look at the Example

(Proportion)	Female	Male	Total
Liberal Arts	.59 (.65)	.41 (.28)	(.42)
Science	.36 (.17)	.64 (.18)	(.18)
Engineering	.17 (.18)	.83 (.54)	(.40)
Total	.38	.62	1

Rejecting $H_0 \Rightarrow$ conditional probabilities are not consistent with marginal probabilities



Notes

Example: Comparing Two Population Proportions

Let $p_1 = \mathbb{P}(\text{Female}|\text{Liberal Arts})$ and $p_2 = \mathbb{P}(\text{Female}|\text{Science})$.

$n_1 = 640, X_1 = 378, n_2 = 274, X_2 = 99$

- $H_0 : p_1 - p_2 = 0$ vs. $H_a : p_1 - p_2 \neq 0$
- $z_{obs} = \frac{.59 - .36}{\sqrt{\frac{.52 \times .48}{640} + \frac{.52 \times .48}{274}}} = 6.36 > z_{0.025} = 1.96$
- We do have enough statistical evidence to conclude that $p_1 \neq p_2$ at .05% significant level.



Notes

R Code & Output

```
prop.test(x = c(378, 99), n = c(640, 274),
          correct = F)

2-sample test for equality of
proportions without continuity
correction

data:  c(378, 99) out of c(640, 274)
X-squared = 40.432, df = 1, p-value =
2.036e-10
alternative hypothesis: two.sided
95 percent confidence interval:
 0.1608524 0.2977699
sample estimates:
 prop 1    prop 2 
0.5906250 0.3613139
```

Inference for
Bivariate
Categorical Data

CLEMSON
UNIVERSITY

21.25

Notes

Example: Test for Homogeneity

Let $p_1 = \mathbb{P}(\text{Liberal Arts})$, $p_2 = \mathbb{P}(\text{Science})$,
 $p_3 = \mathbb{P}(\text{Engineering})$

- The Hypotheses:

$H_0 : p_1 = p_2 = p_3 = \frac{1}{3}$

$H_a : \text{At least one is different}$

- The Test Statistic:

$$\chi^2_{obs} = \frac{(640 - 509.33)^2}{509.33} + \frac{(274 - 509.33)^2}{509.33} + \frac{(614 - 509.33)^2}{509.33}$$

$$= 33.52 + 108.73 + 21.51 = 163.76 > \chi^2_{.05, df=2} = 5.99$$

- Rejecting H_0 at .05 level

Inference for
Bivariate
Categorical Data

CLEMSON
UNIVERSITY

21.26

Notes

R Code & Output

```
chisq.test(x = c(640, 274, 614), p = rep(1/3, 3))

Chi-squared test for given
probabilities

data:  c(640, 274, 614)
X-squared = 163.76, df = 2, p-value
< 2.2e-16
```

Inference for
Bivariate
Categorical Data

CLEMSON
UNIVERSITY

21.27

Notes
